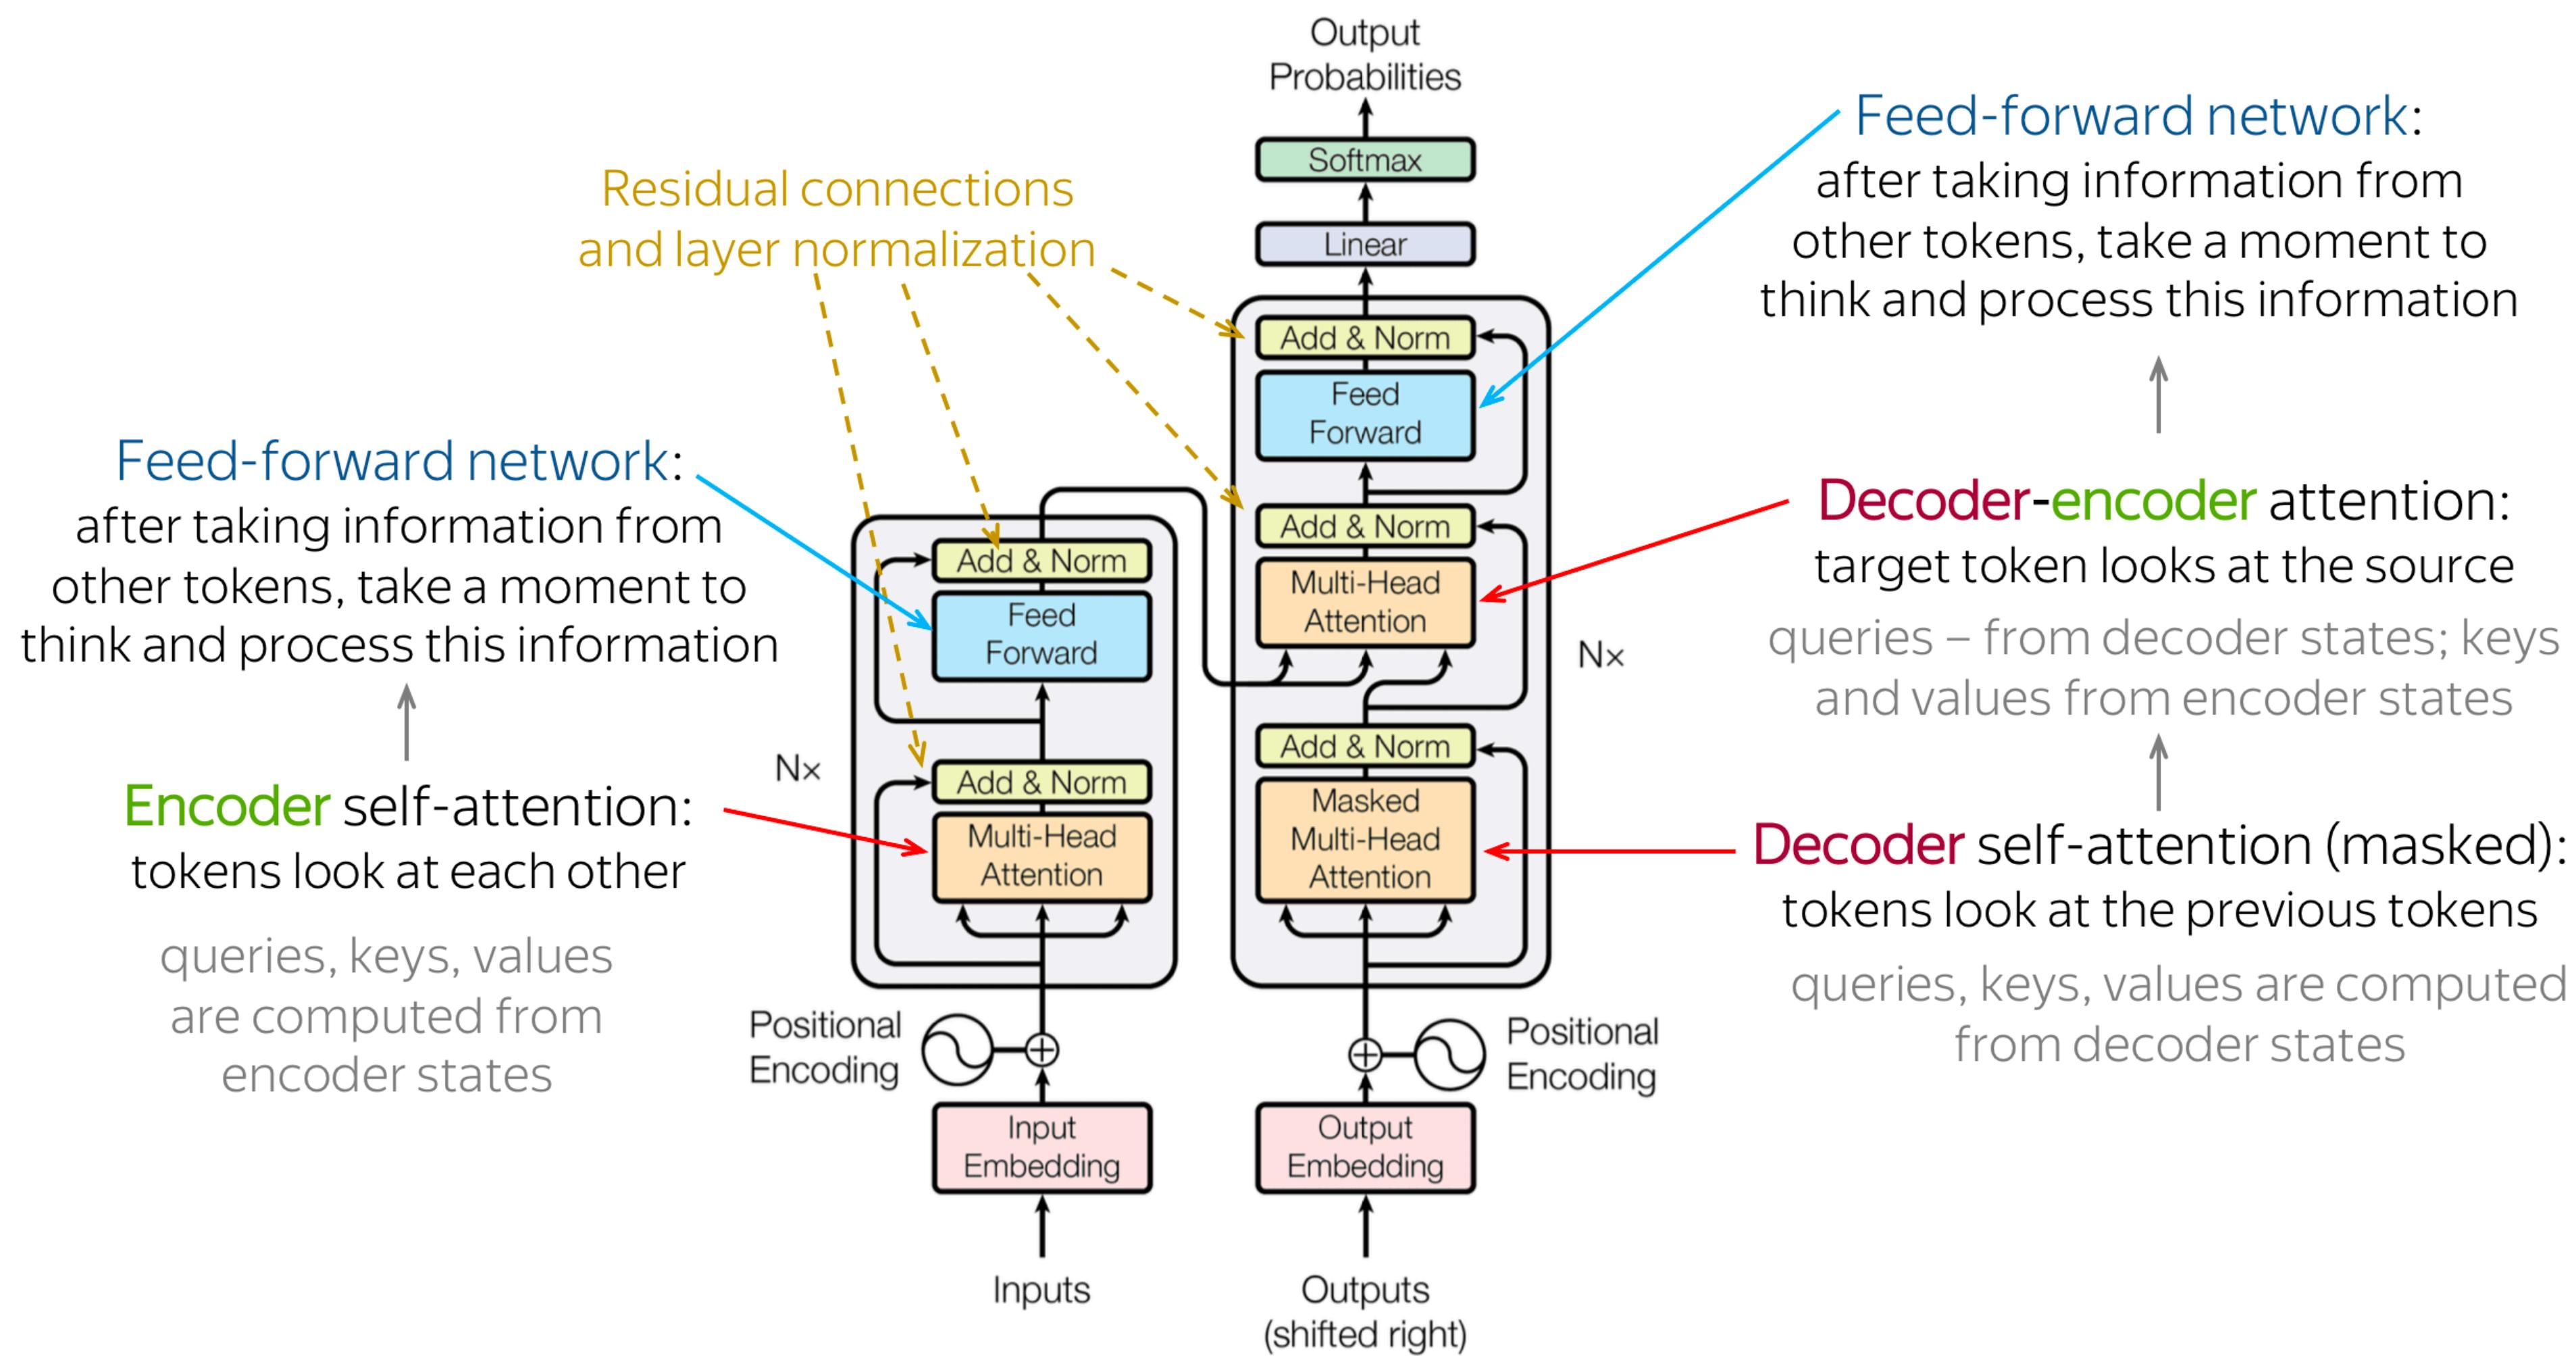


(Deep) Generative Models

Phillip Rust, May 16, 2022

Previous Lecture – Transformers



Today's Lecture

1. Supervised vs Unsupervised Learning
2. Generative Modeling
3. Applications of generative models
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
5. (Variational) Autoencoders
6. Generative Adversarial Networks (GANs)

Today's Lecture

- 1. Supervised vs Unsupervised Learning**
2. Generative Modeling
3. Applications of generative models
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
5. (Variational) Autoencoders
6. Generative Adversarial Networks (GANs)

Supervised vs Unsupervised Learning

Supervised Learning

- Dataset: (x, y)
 $x = \text{input (features)}, y = \text{label}$
- **Goal:** Learn a *mapping function* $f: x \rightarrow y$
- **Examples:** Classification, Regression, Object Detection, Image Captioning, ...



→ **Cat**

“Horrible service! I ordered a cat but they delivered raw chicken instead. 0/10”

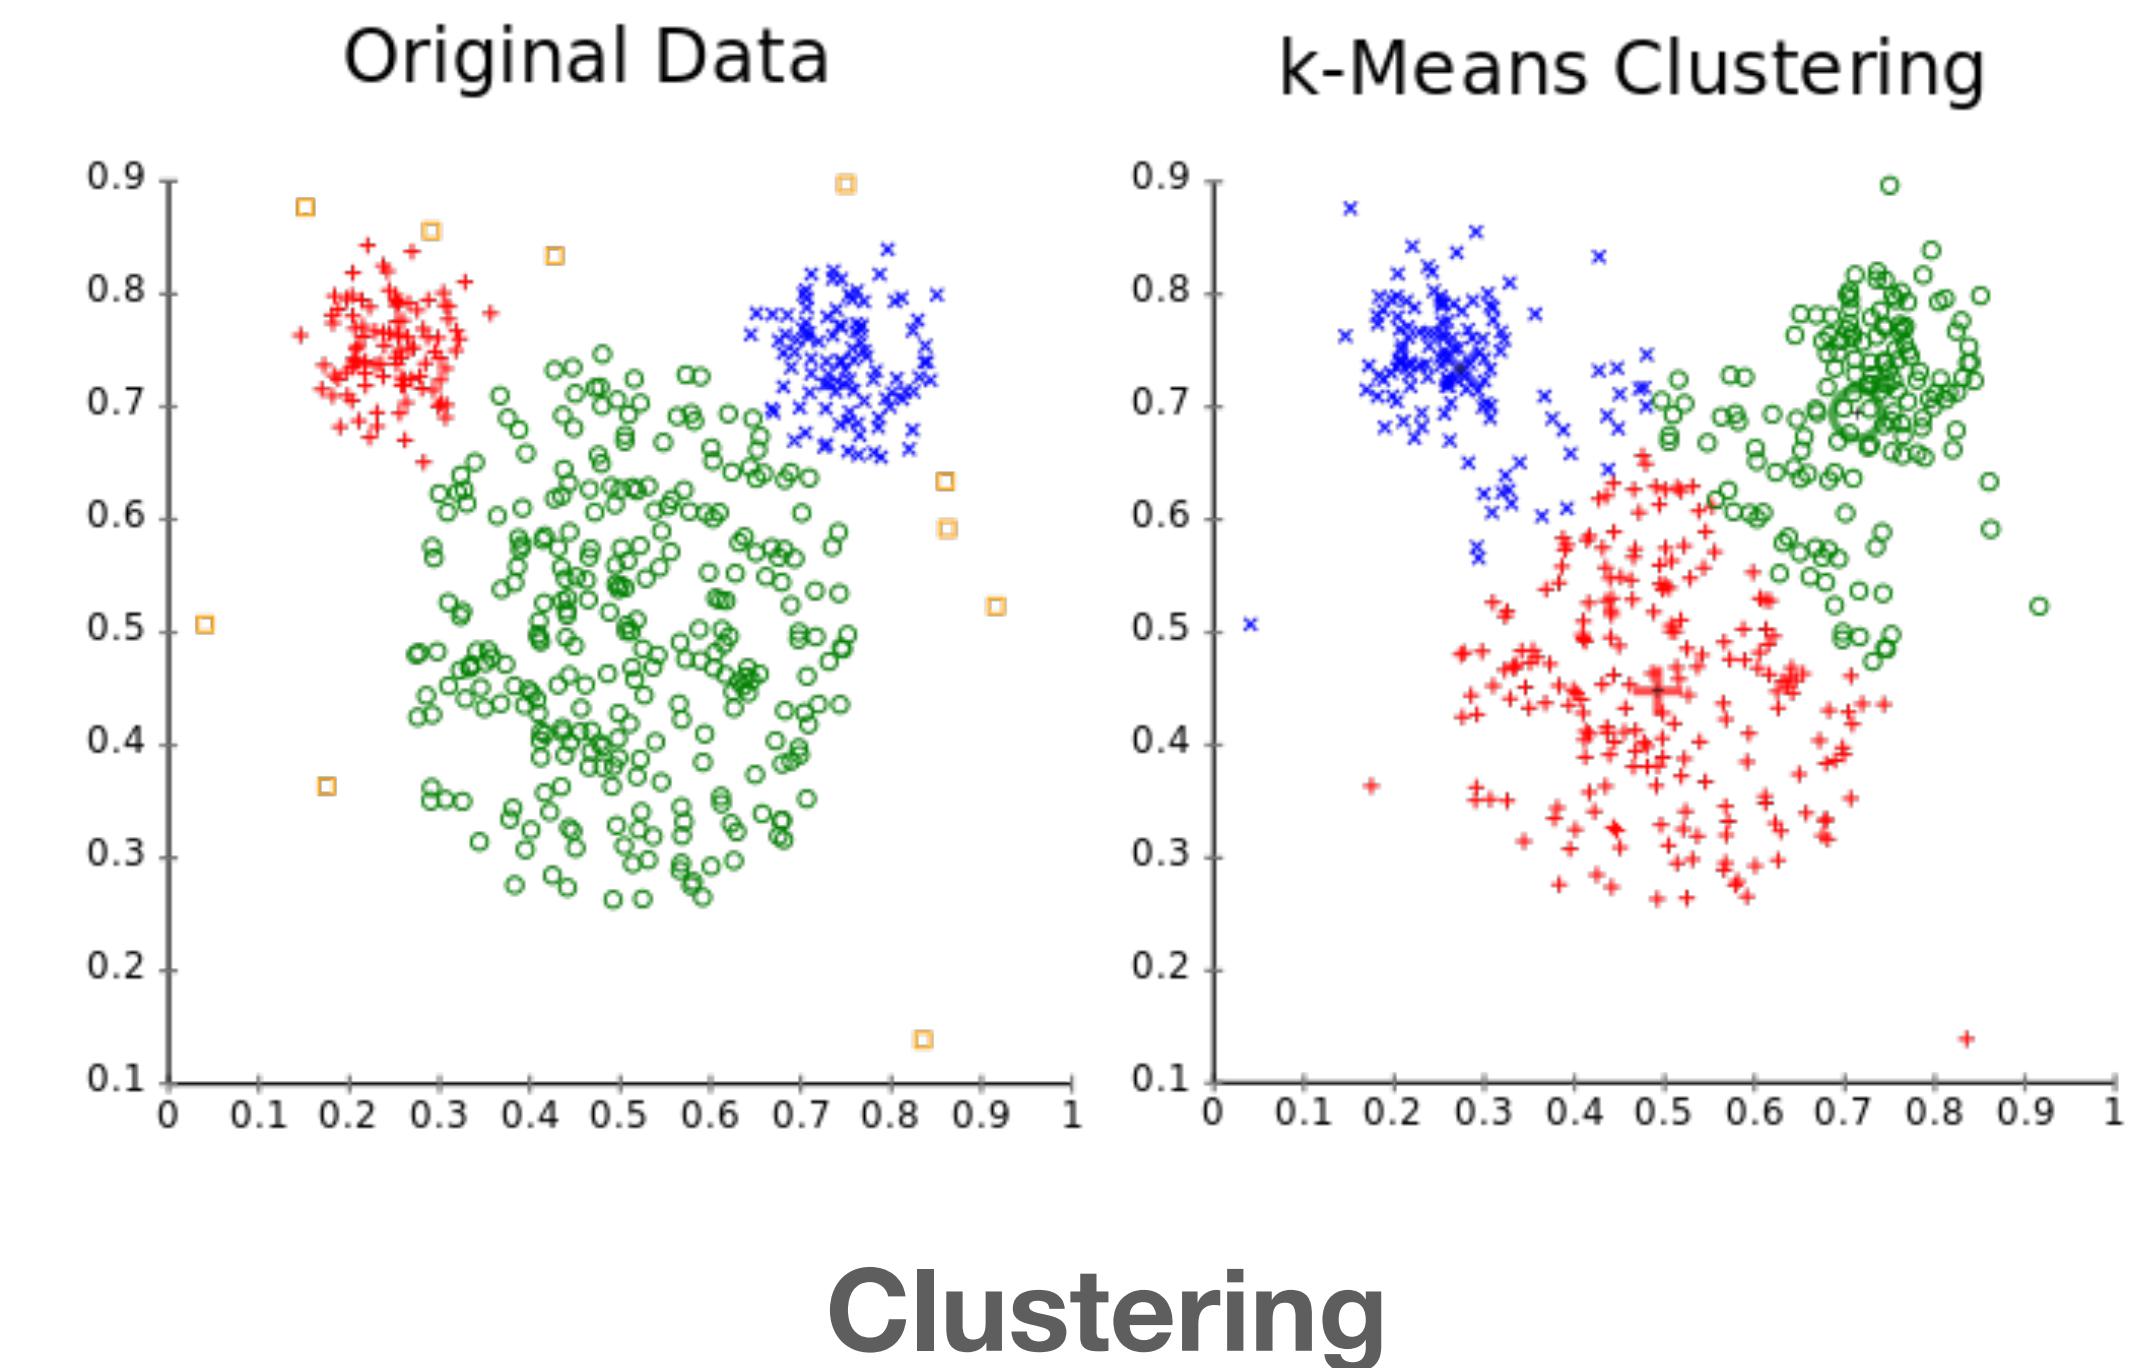
→ **Negative**

Supervised vs Unsupervised Learning

Unsupervised Learning

- Dataset: x (only inputs, no labels)
- **Goal:** Learn some underlying hidden structure of the data
- **Examples:** Clustering, dimensionality reduction, density estimation, ...

Figure from https://upload.wikimedia.org/wikipedia/commons/thumb/0/09/ClusterAnalysis_Mouse.svg/1024px-ClusterAnalysis_Mouse.svg.png

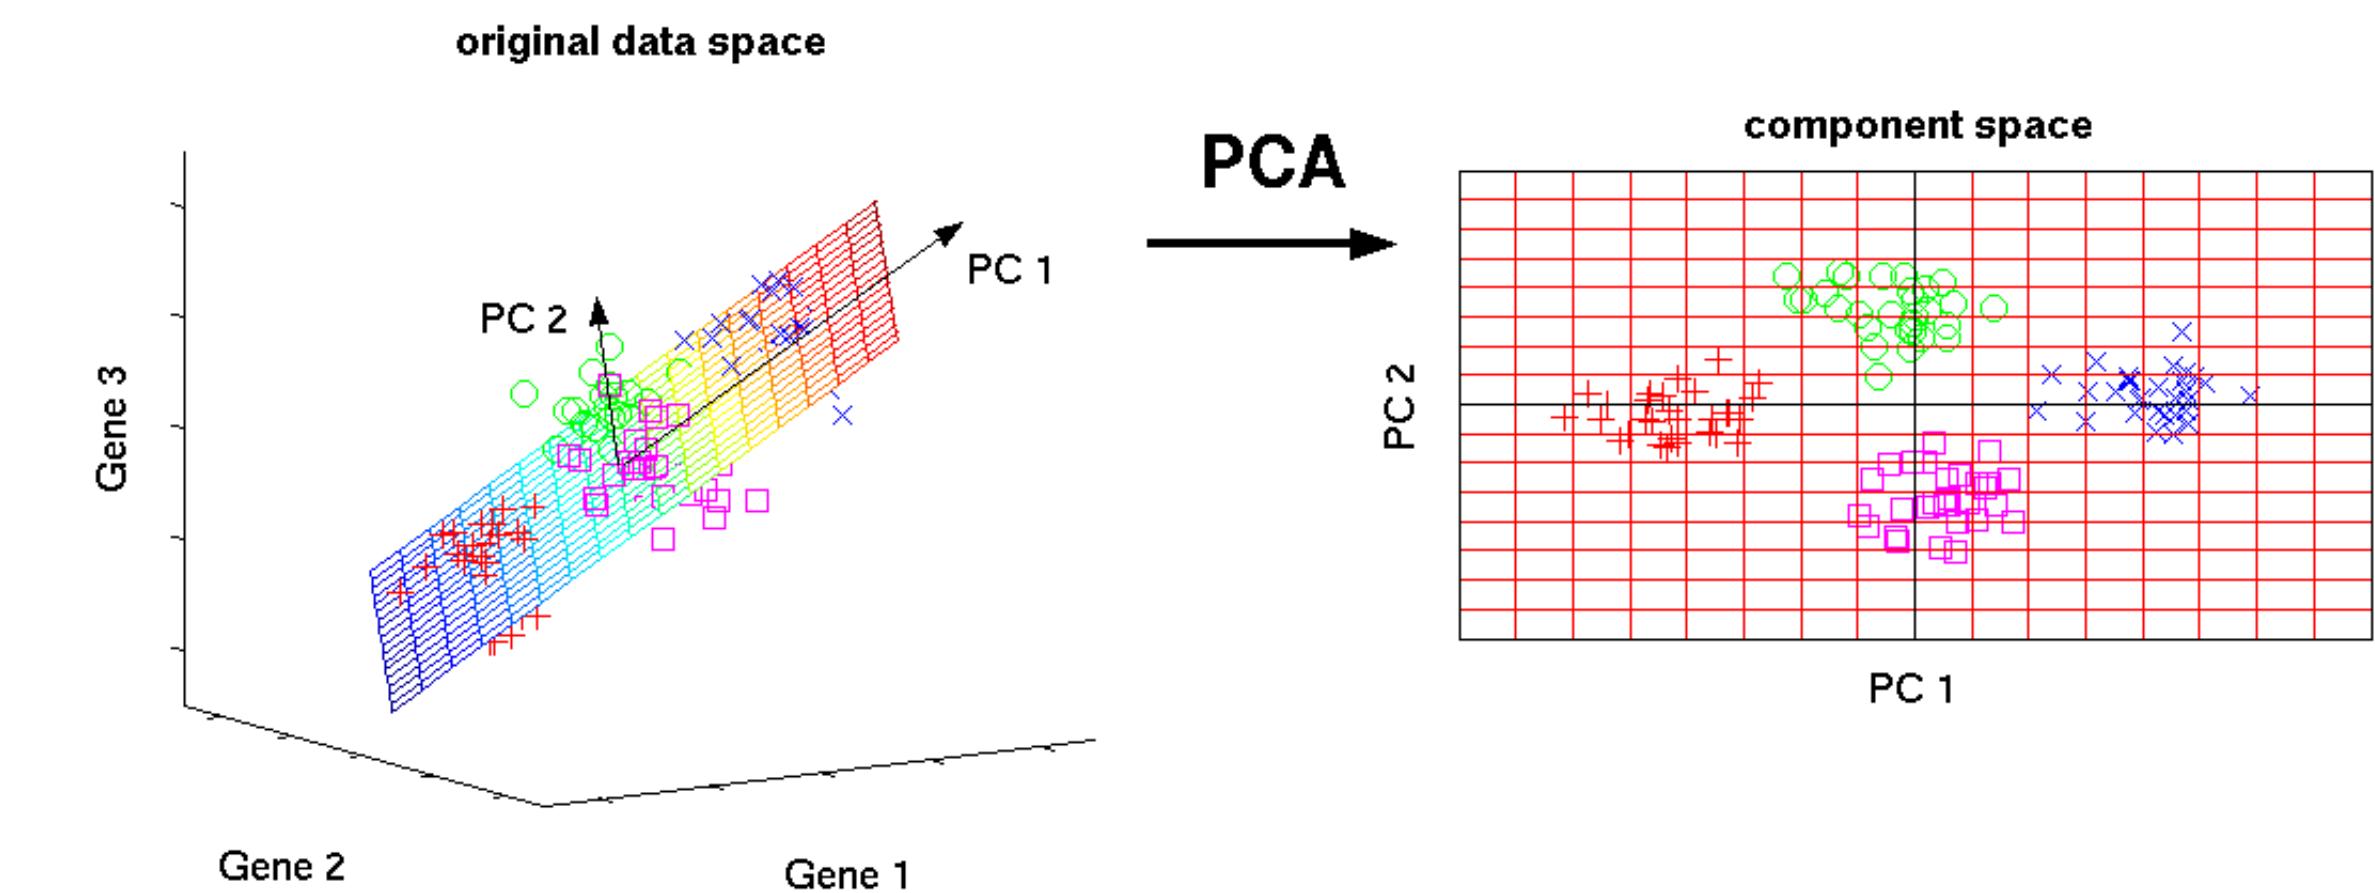


Supervised vs Unsupervised Learning

Unsupervised Learning

- Dataset: x (only inputs, no labels)
- **Goal:** Learn some underlying hidden structure of the data
- **Examples:** Clustering, dimensionality reduction, density estimation, ...

Figure from http://phdthesis-bioinformatics-maxplanckinstitute-molecularplantphys.matthias-scholz.de/fig_pca_illu3d.png



Dimensionality Reduction (3D \rightarrow 2D)

Supervised vs Unsupervised Learning

Unsupervised Learning

- Dataset: x (only inputs, no labels)
- **Goal:** Learn some underlying hidden structure of the data
- **Examples:** Clustering, dimensionality reduction, density estimation, ...

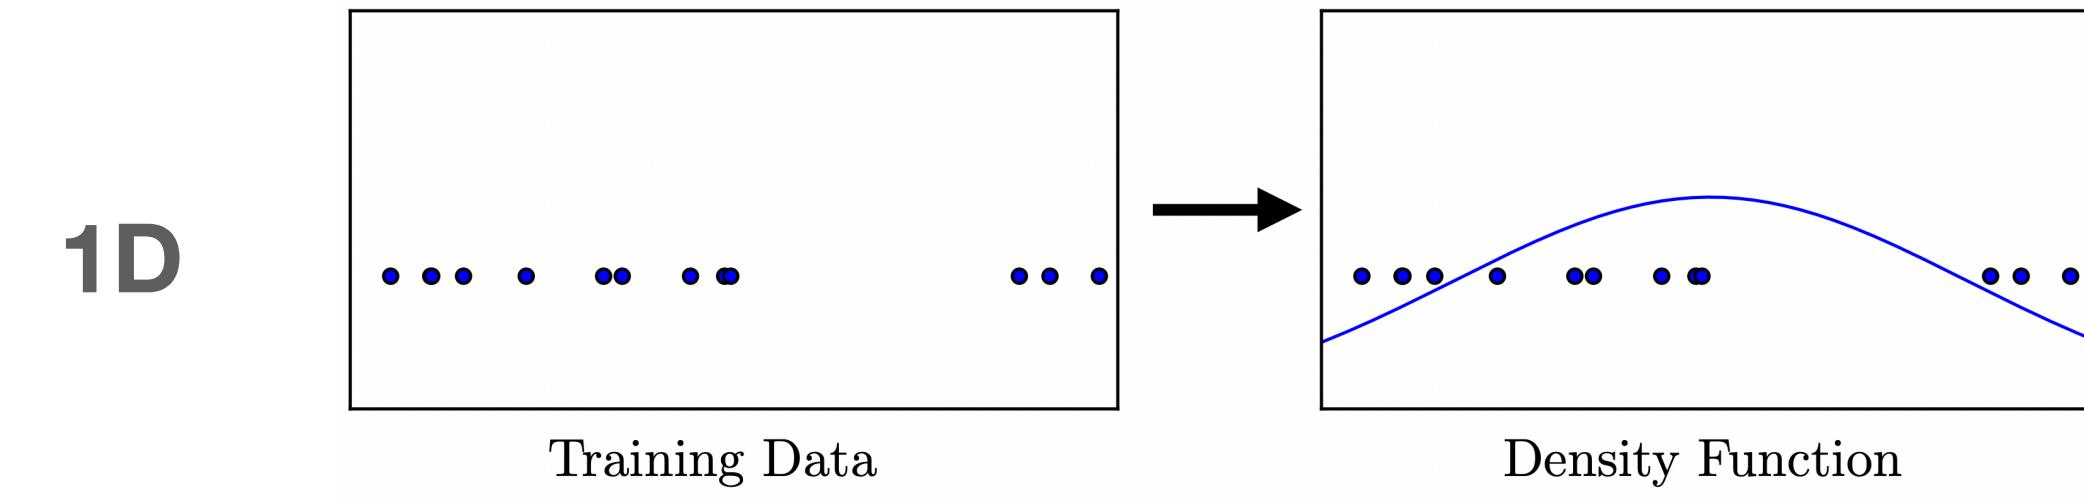


Figure from https://www.iangoodfellow.com/slides/2018-06-22-gan_tutorial.pdf

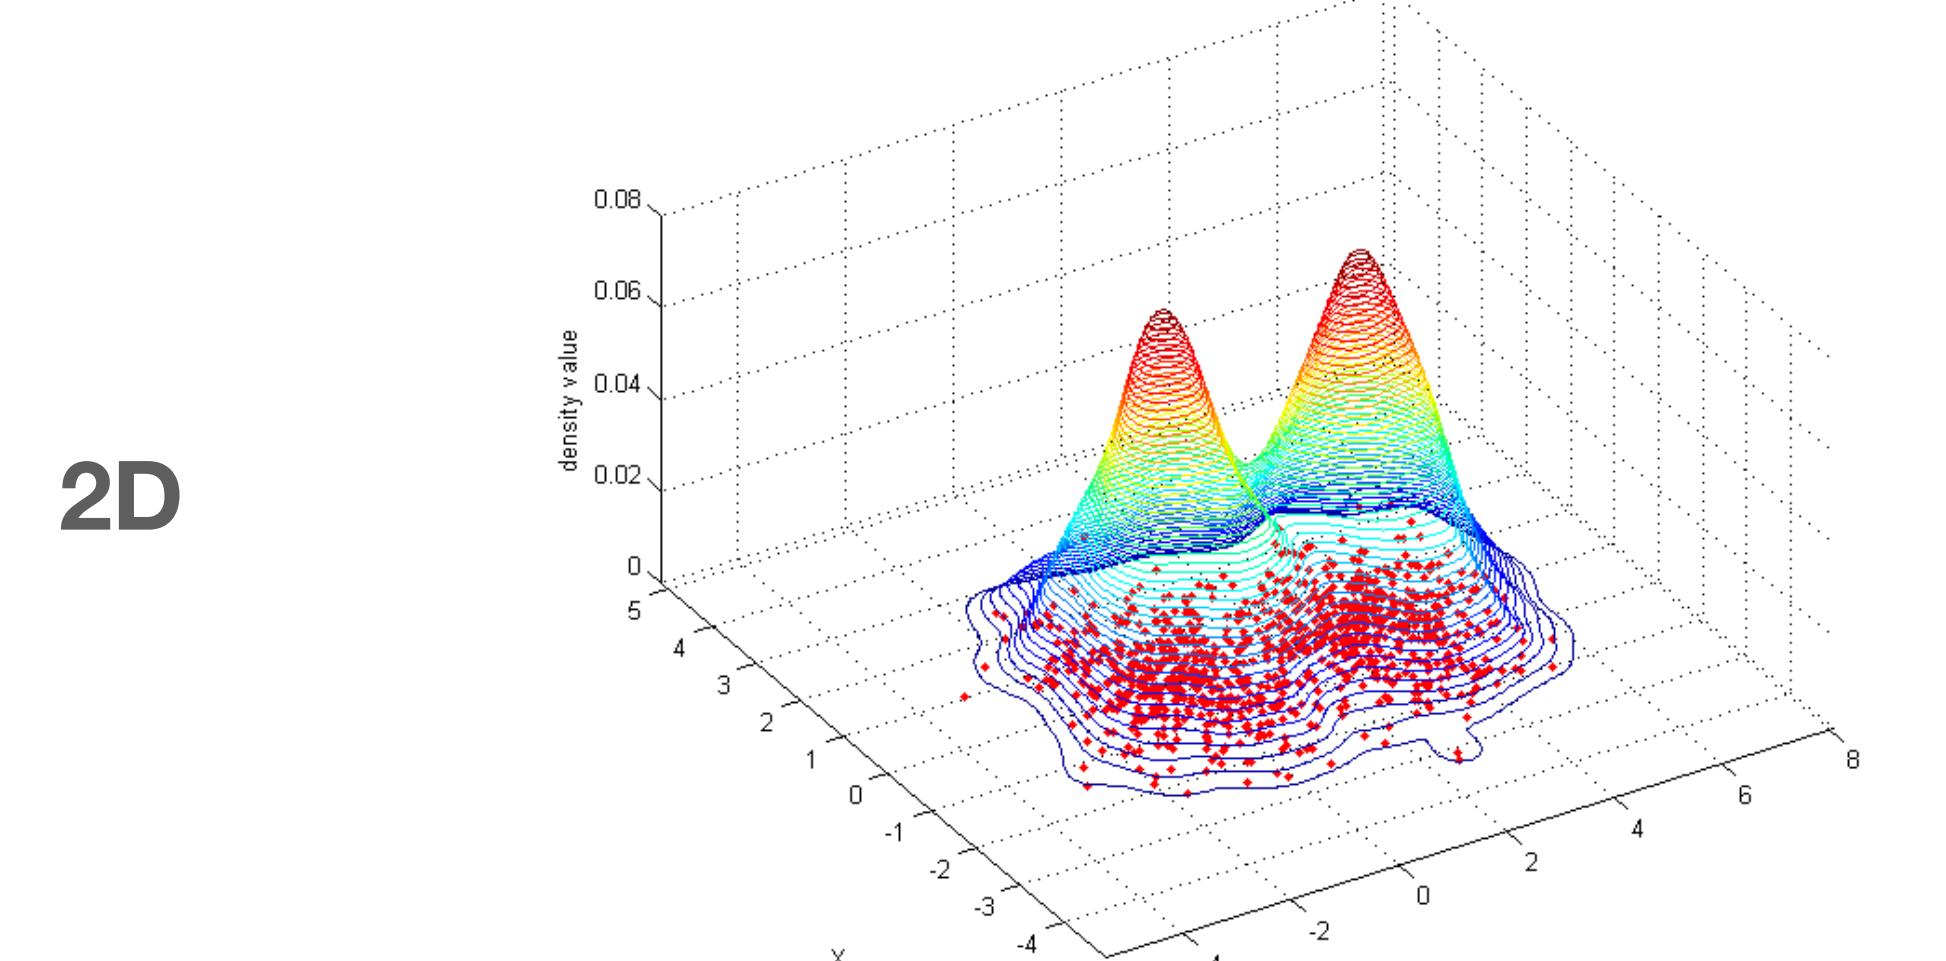


Figure from https://upload.wikimedia.org/wikipedia/commons/1/1b/Bivariate_example.png

Density Estimation

Supervised vs Unsupervised Learning

Supervised Learning

- Dataset: (x, y)
 x = input (features), y = label
- **Goal:** Learn a *mapping function* $f: x \rightarrow$
- **Examples:** Classification, Regression, Object Detection, Image Captioning, ...

Unsupervised Learning

- Dataset: x (only inputs, no labels)
- **Goal:** Learn some underlying hidden structure of the data
- **Examples:** Clustering, dimensionality reduction, density estimation, ...

Today's Lecture

1. Supervised vs Unsupervised Learning
- 2. Generative Modeling**
3. Applications of generative models
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
5. (Variational) Autoencoders
6. Generative Adversarial Networks (GANs)

Generative Modeling

Learning the data-generating distribution through density estimation

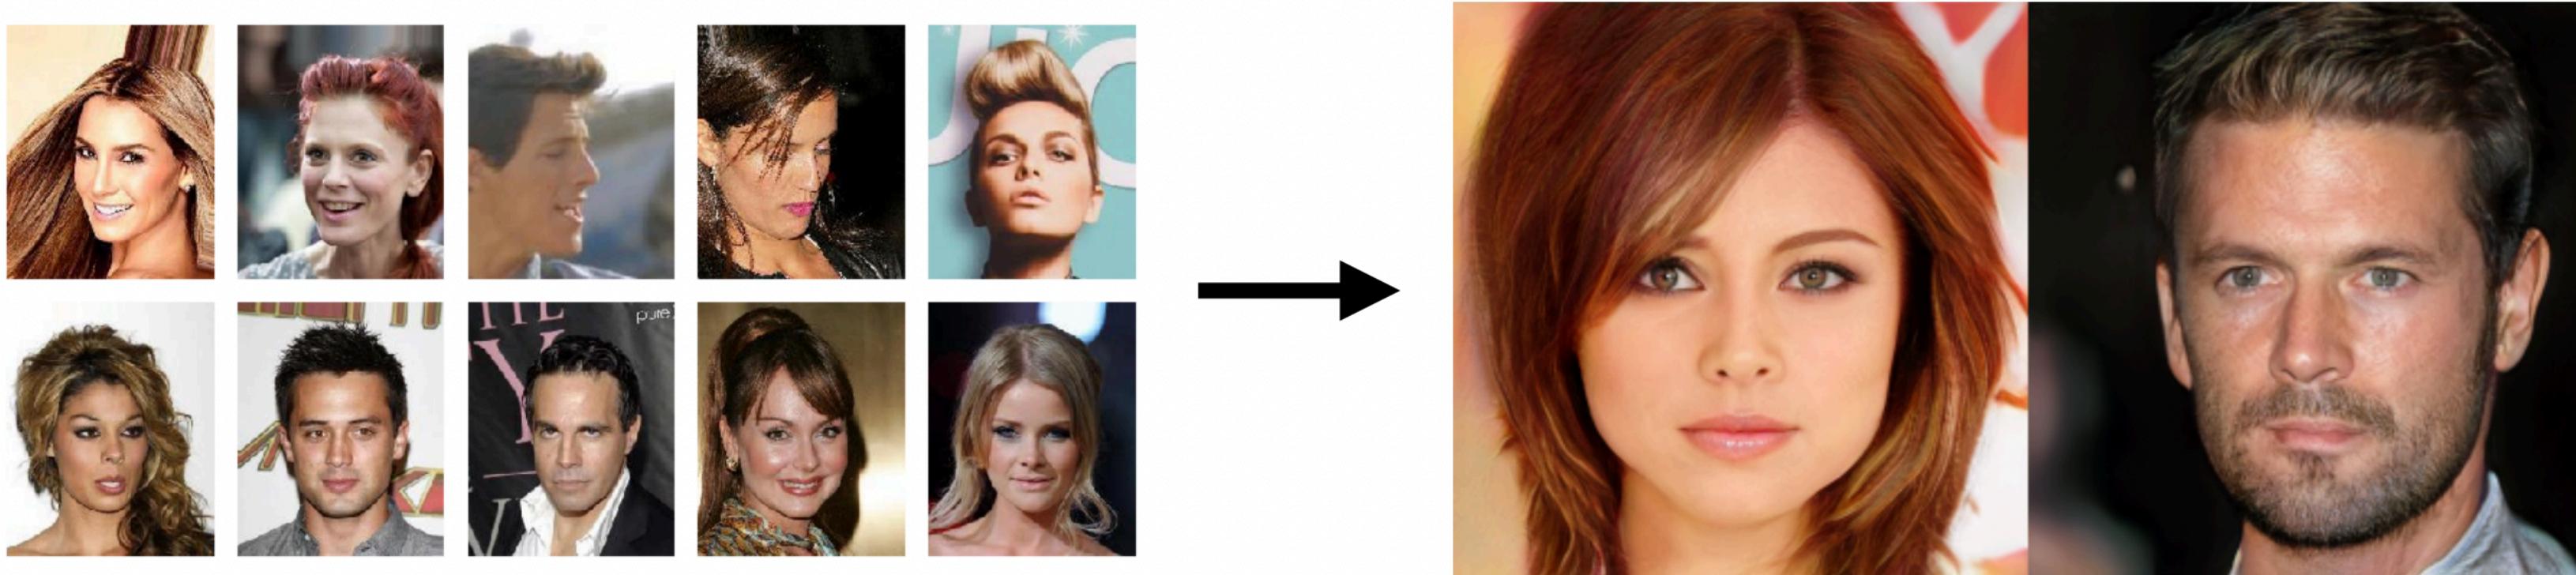


Figure taken from
https://www.iangoodfellow.com/slides/2018-06-22-gan_tutorial.pdf

Training Data
(CelebA)

Sample Generator
(Karras et al, 2017)

Premise: There is a distribution $p(x)$ that underlies our training data x

Plan: We want to learn the distribution $p(x)$ as closely as possible and sample from it to generate new data

Method: Density estimation

Explicit : Explicitly define $p_{model}(x)$ and solve for it

Implicit : Learn to sample from $p_{model}(x)$ without having to define it

Today's Lecture

1. Supervised vs Unsupervised Learning
2. Generative Modeling
- 3. Applications of generative models**
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
5. (Variational) Autoencoders
6. Generative Adversarial Networks (GANs)

Applications of generative models

Creative purposes

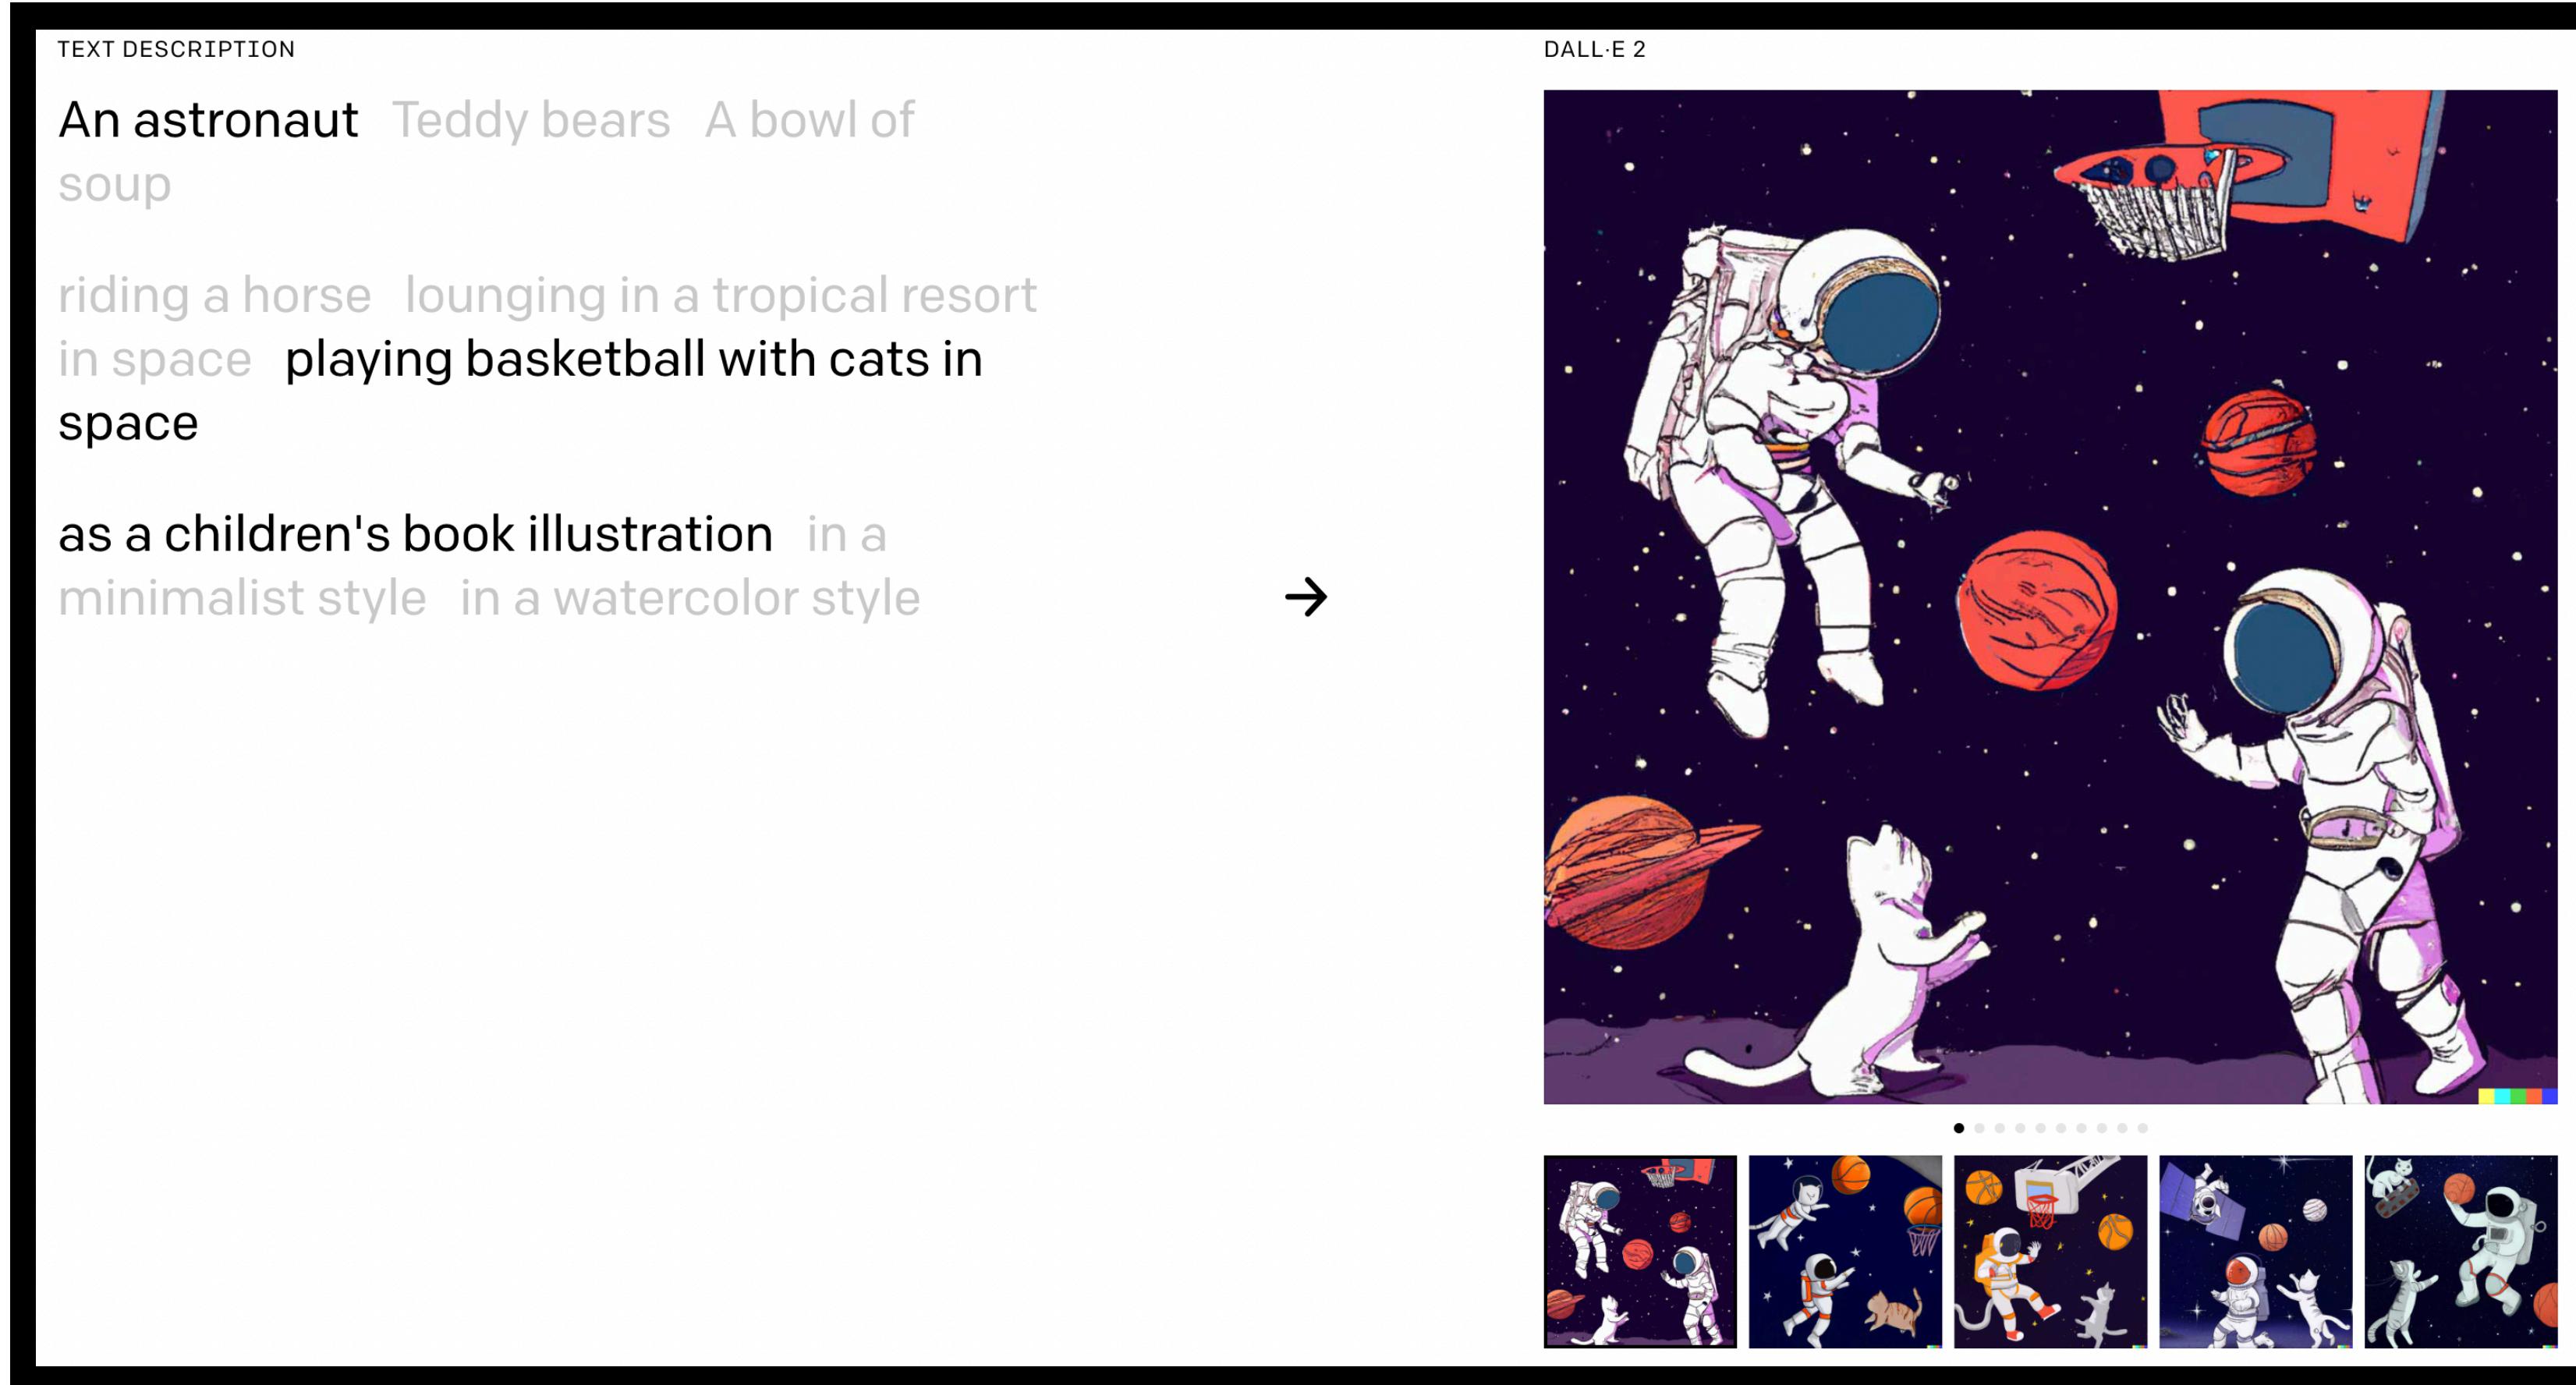


Figure from <https://openai.com/dall-e-2/>



Copyright Sofia Crespo
Figure from <https://blogs.nvidia.com/blog/2022/05/12/ai-art-times-square-endangered-species/>

Applications of generative models

Creative purposes



Copyright Michael Hull/Times Square Arts

Figure from <https://blogs.nvidia.com/blog/2022/05/12/ai-art-times-square-endangered-species/>

Applications of generative models

Compression & Super-resolution

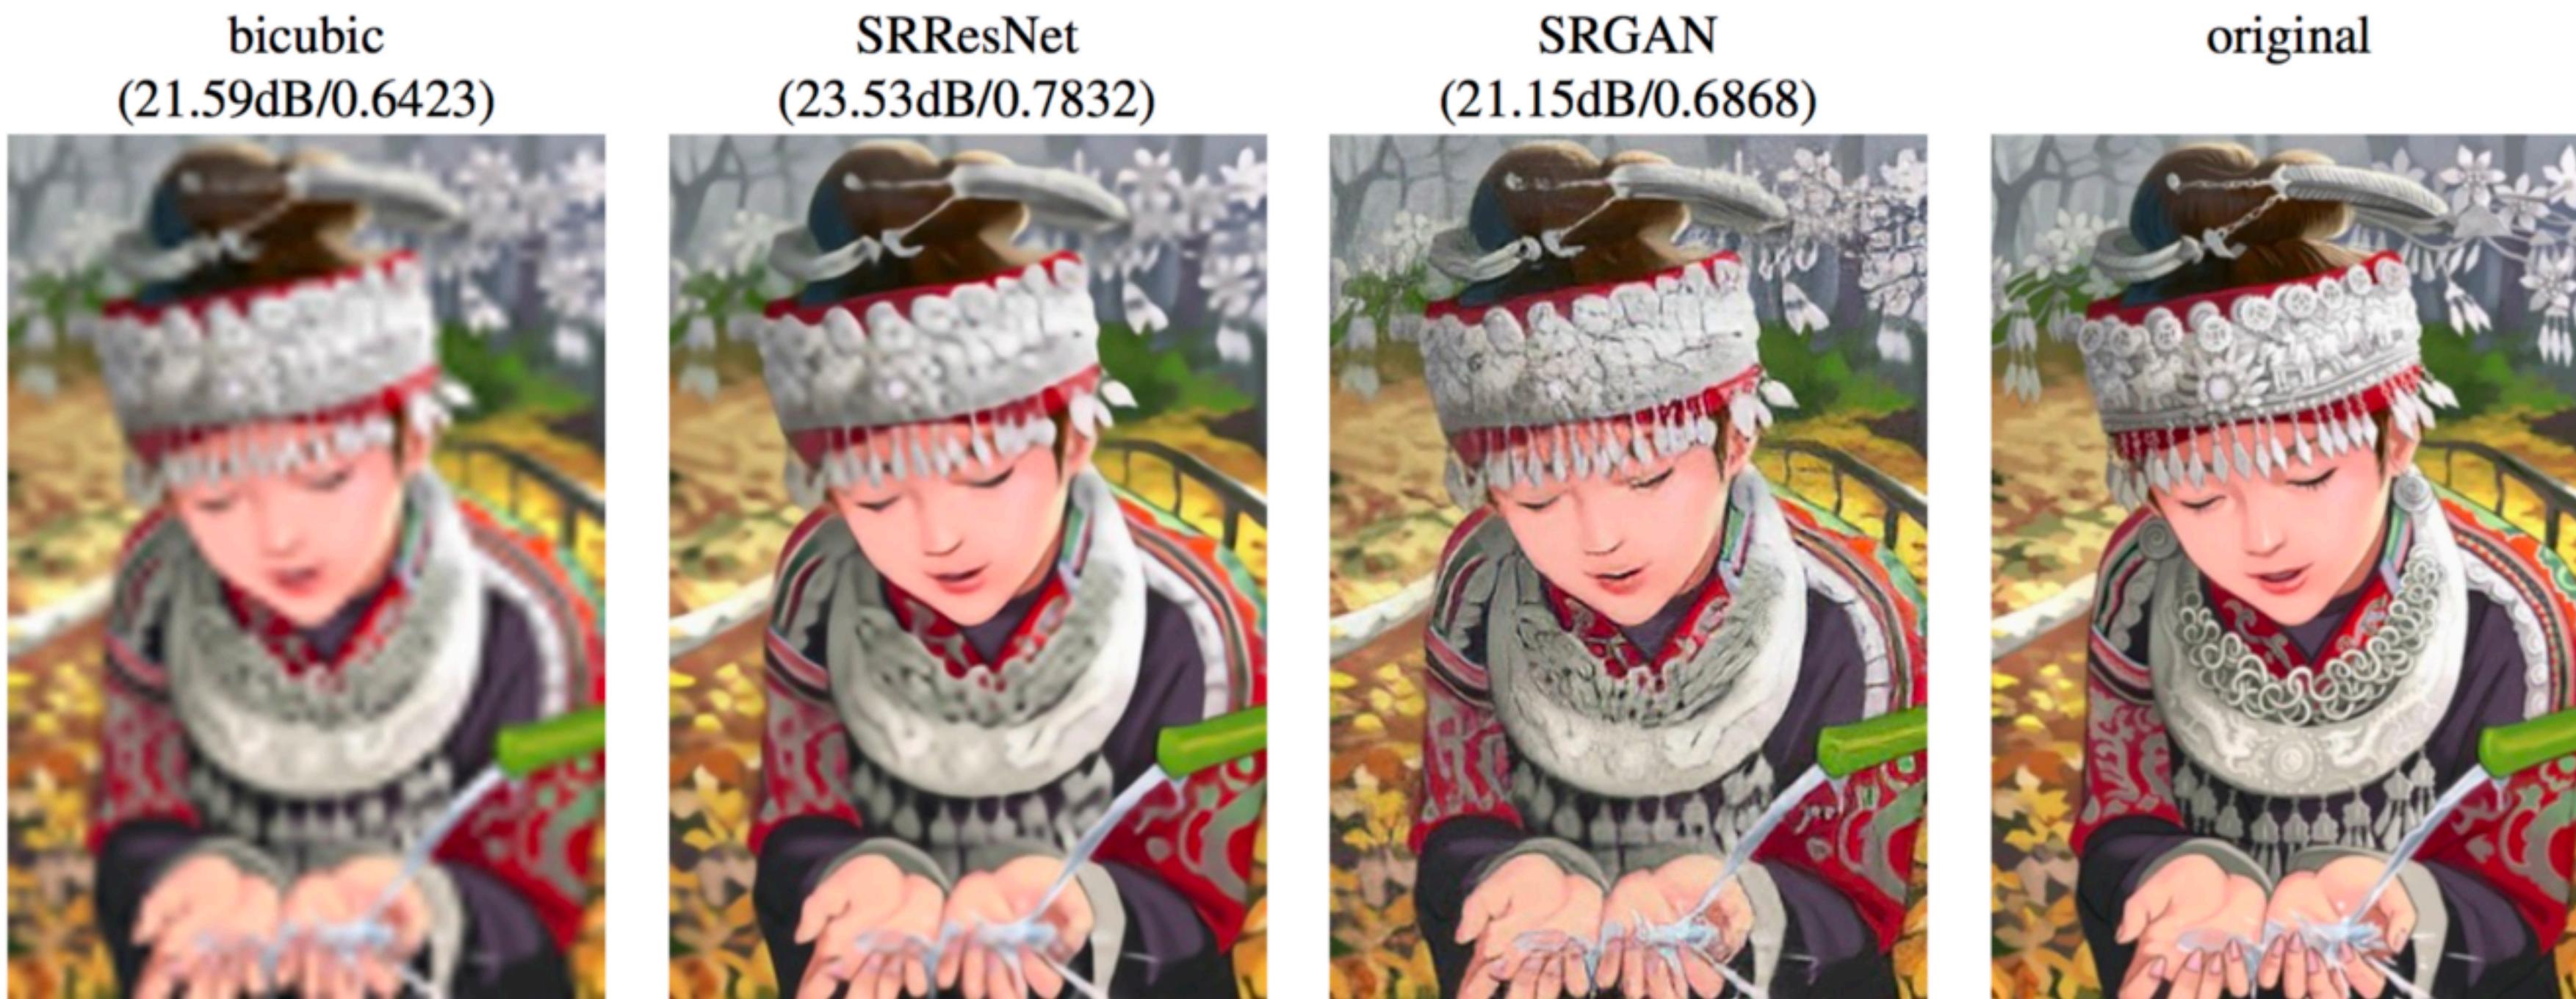


Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]

Applications of generative models

Drug Design & Other medical applications

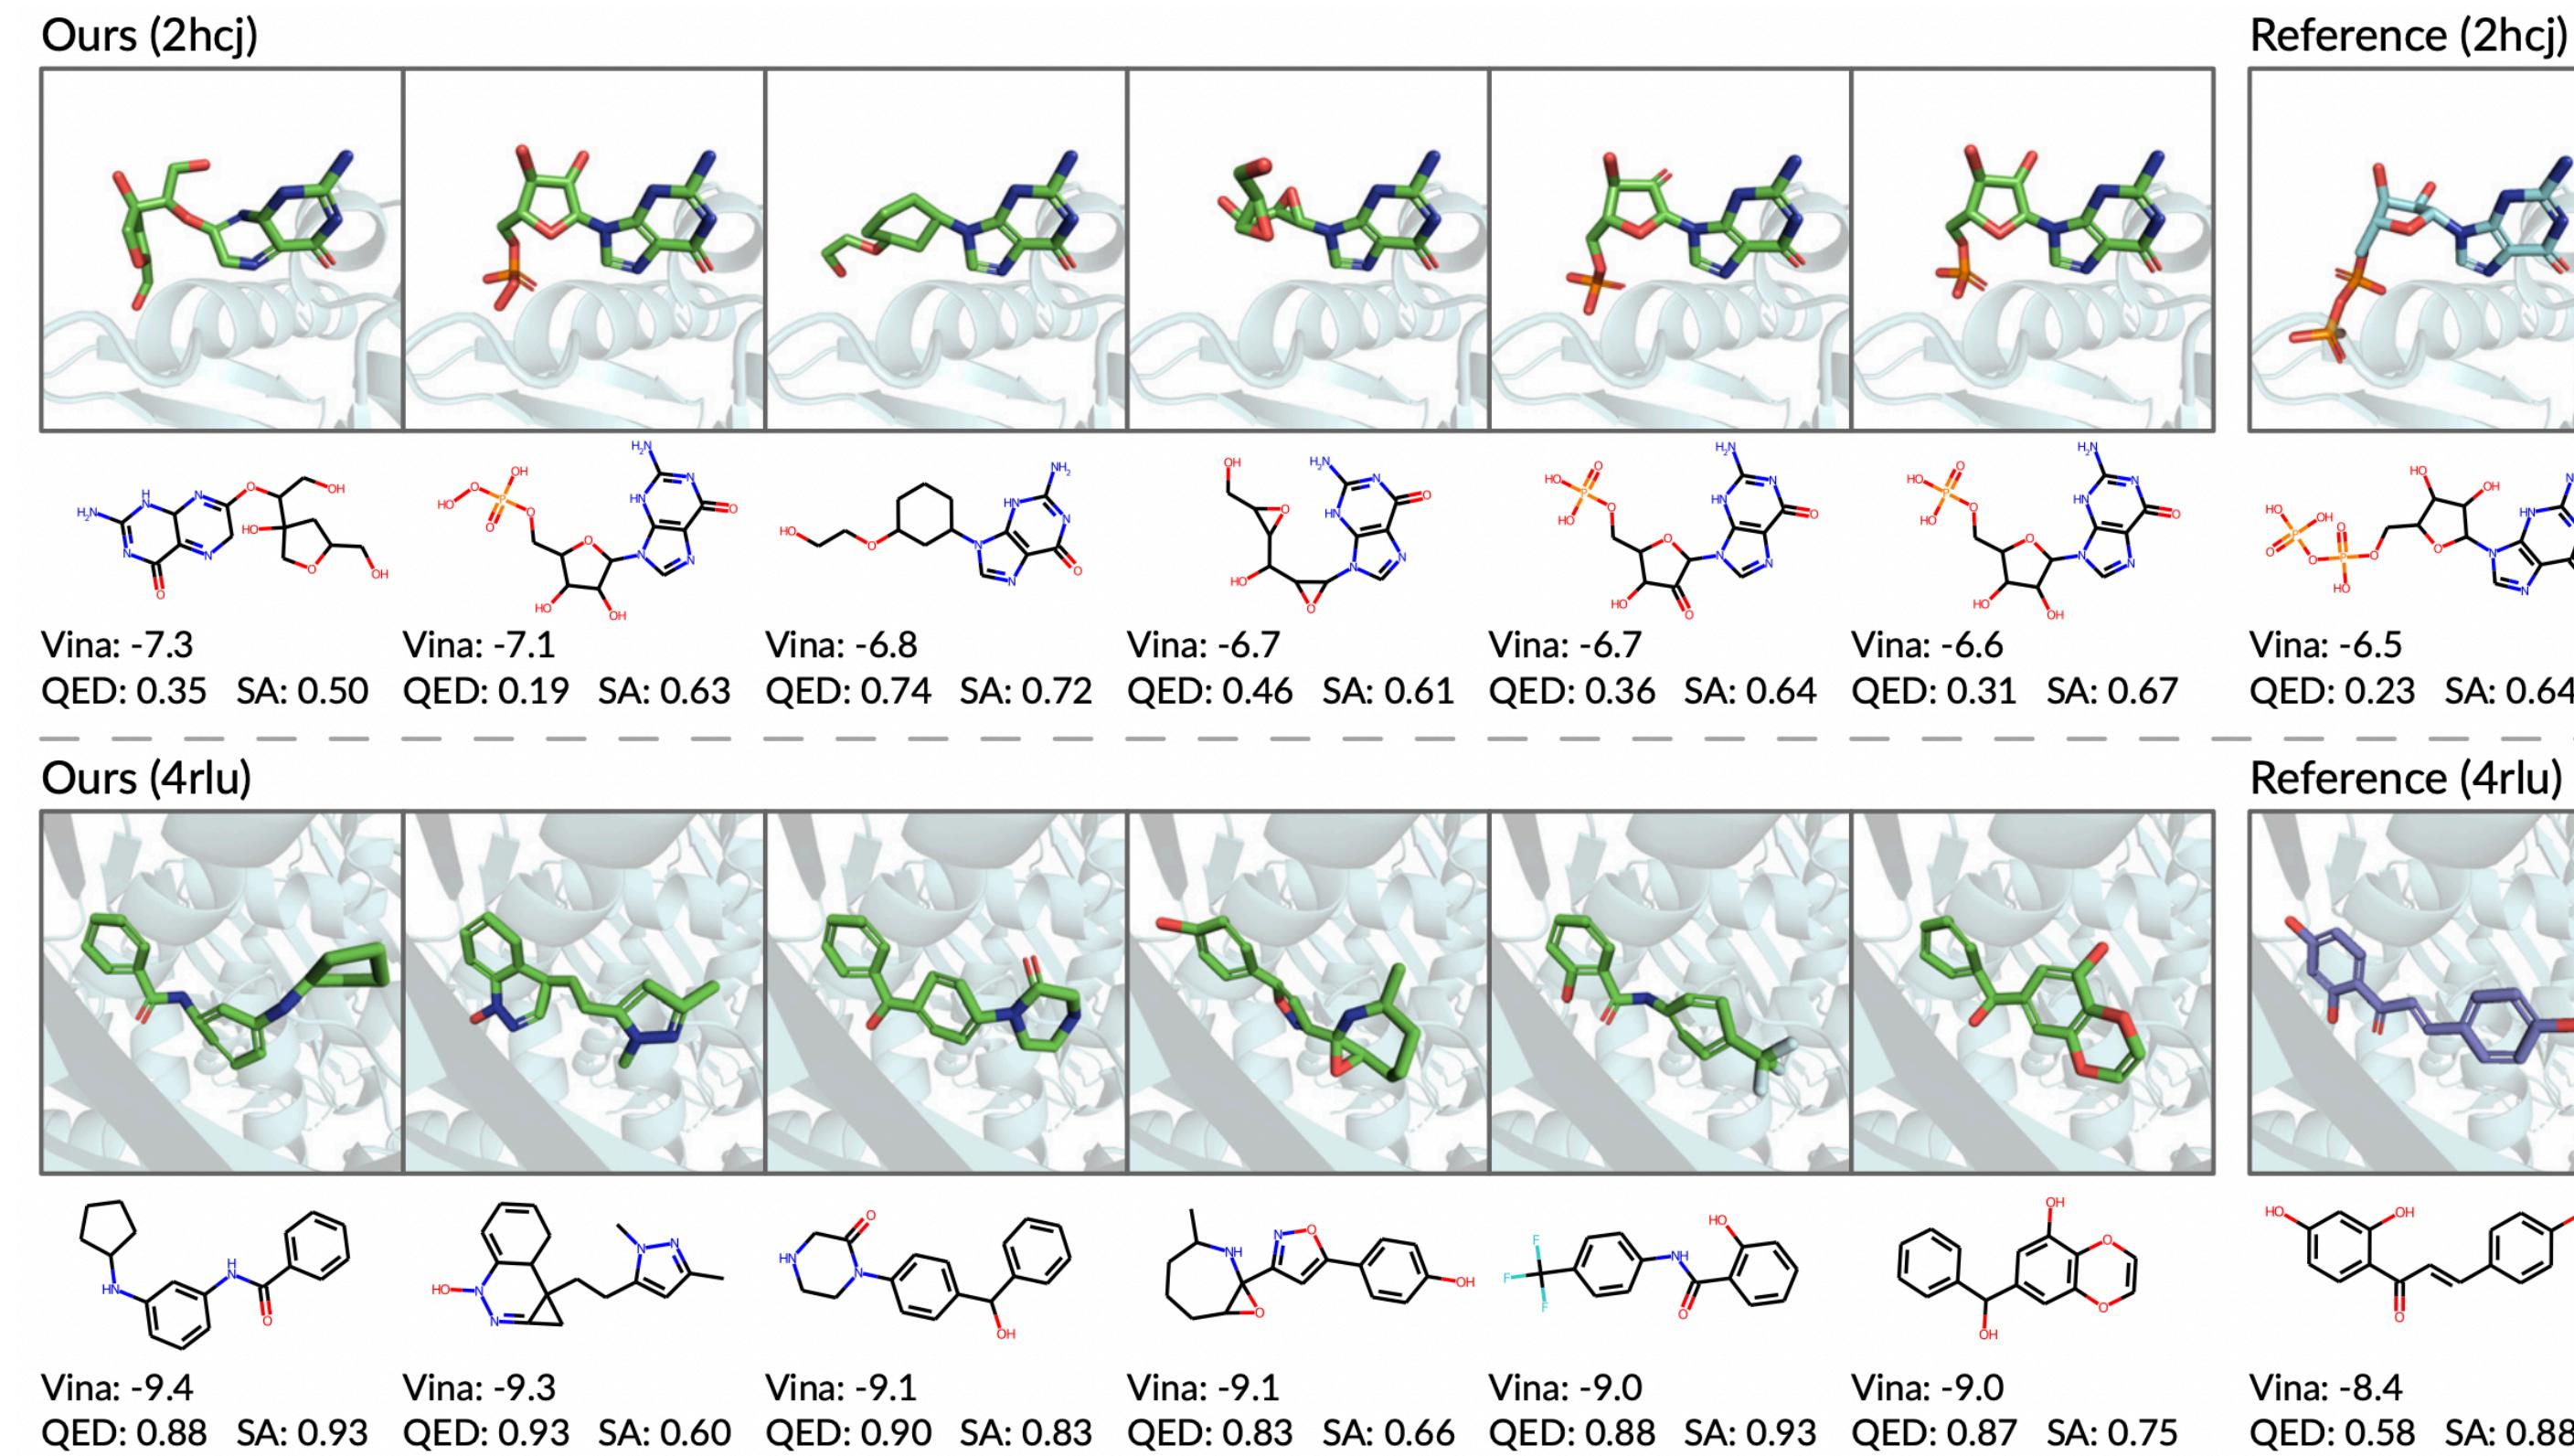


Figure 4: Generated molecules with top binding affinity and the reference molecule for two representative binding sites. Lower Vina score indicates higher binding affinity.

Applications of generative models

- Learning general purpose features for downstream tasks (e.g. classification)
- Gaining insights from high-dimensional data (physics, medical imaging, etc.)
- Modeling the physical world for simulation and planning (robotics, reinforcement learning)
- ...

Generative models – Taxonomy

Maximum likelihood estimation: Find the set of parameters θ^* that maximizes the (log) likelihood of our data \mathbf{x}

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\text{model}}(\mathbf{x} | \theta)$$

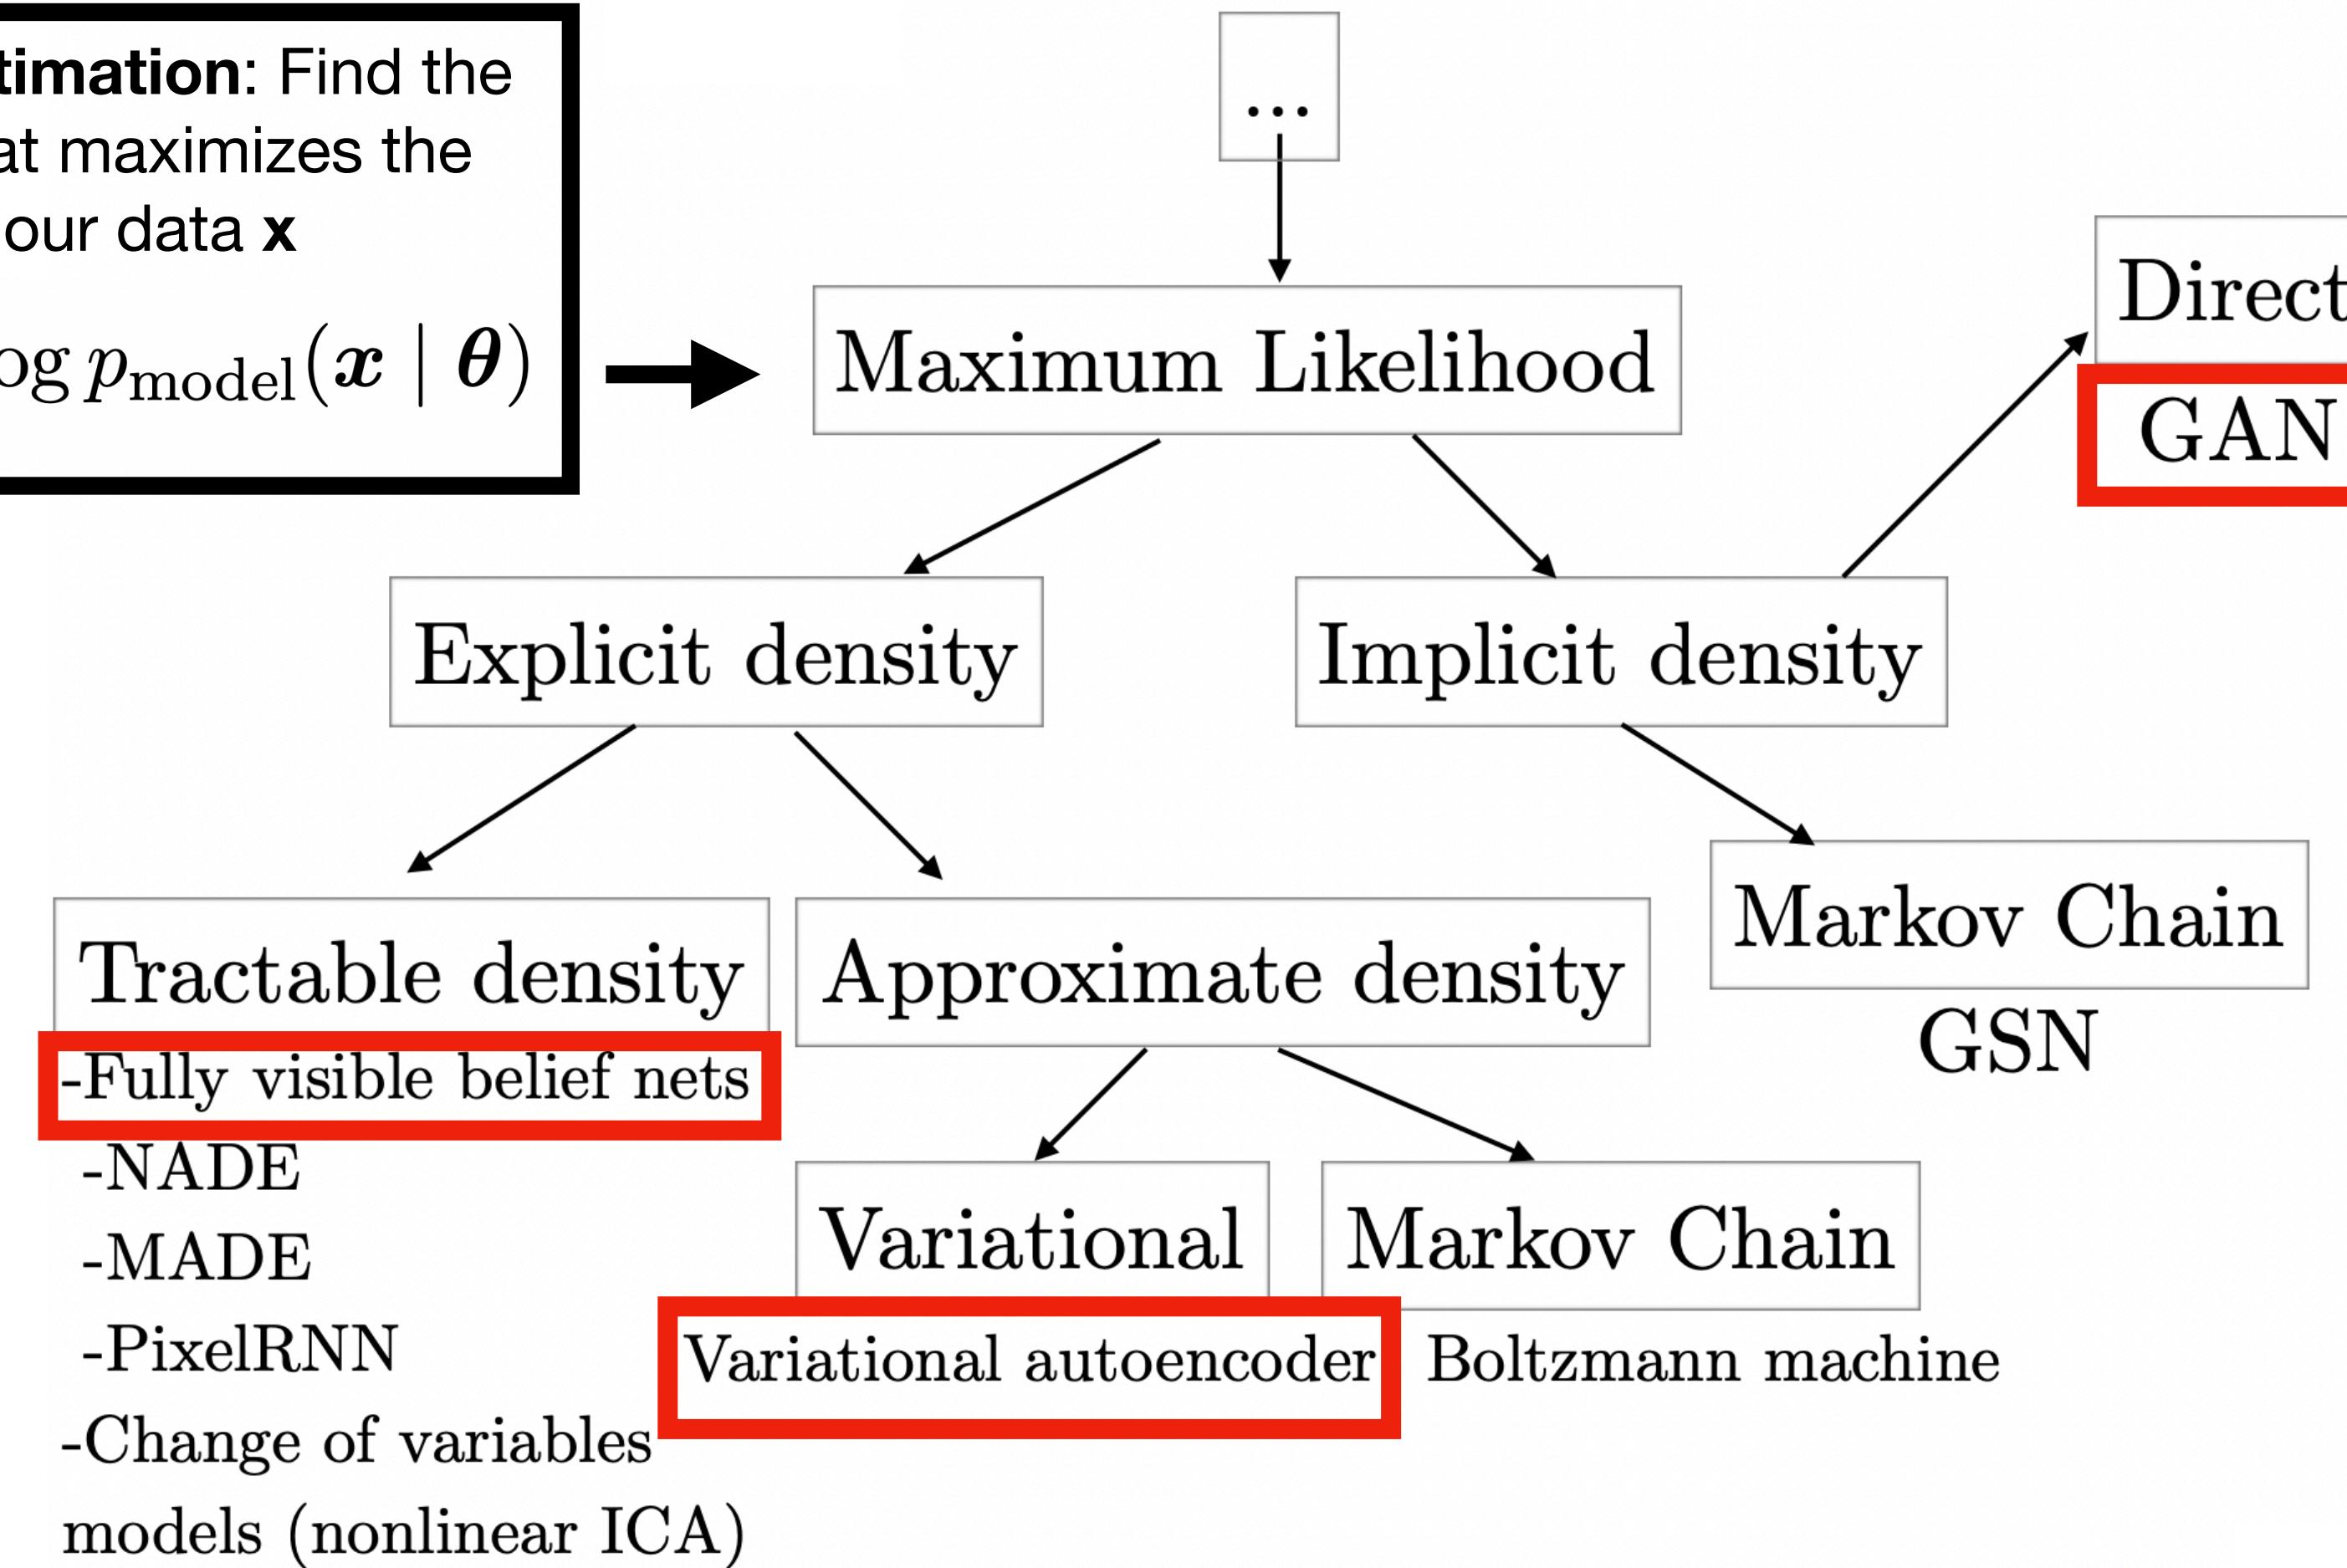


Figure taken from <https://www.iangoodfellow.com/slides/2016-12-04-NIPS.pdf>

Today's Lecture

1. Supervised vs Unsupervised Learning
2. Generative Modeling
3. Applications of generative models
- 4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders**
5. (Variational) Autoencoders
6. Generative Adversarial Networks (GANs)

Fully-visible Belief Networks (FVBNs)

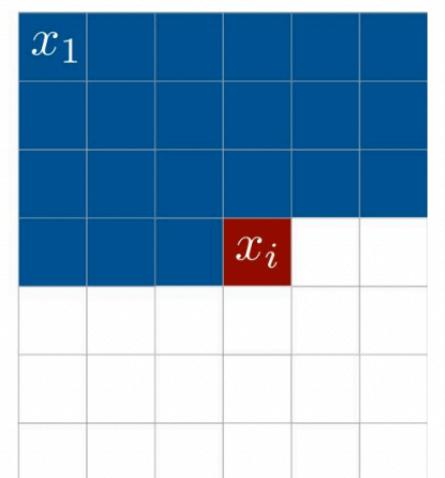
Autoregressive Decoders

- Explicit density model
- Decompose joint likelihood into a product of tractable conditional likelihoods using the chain rule

$$p(x) = p(x_1, x_2, \dots, x_n) \rightarrow p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Likelihood of image x Joint likelihood of each pixel in the image

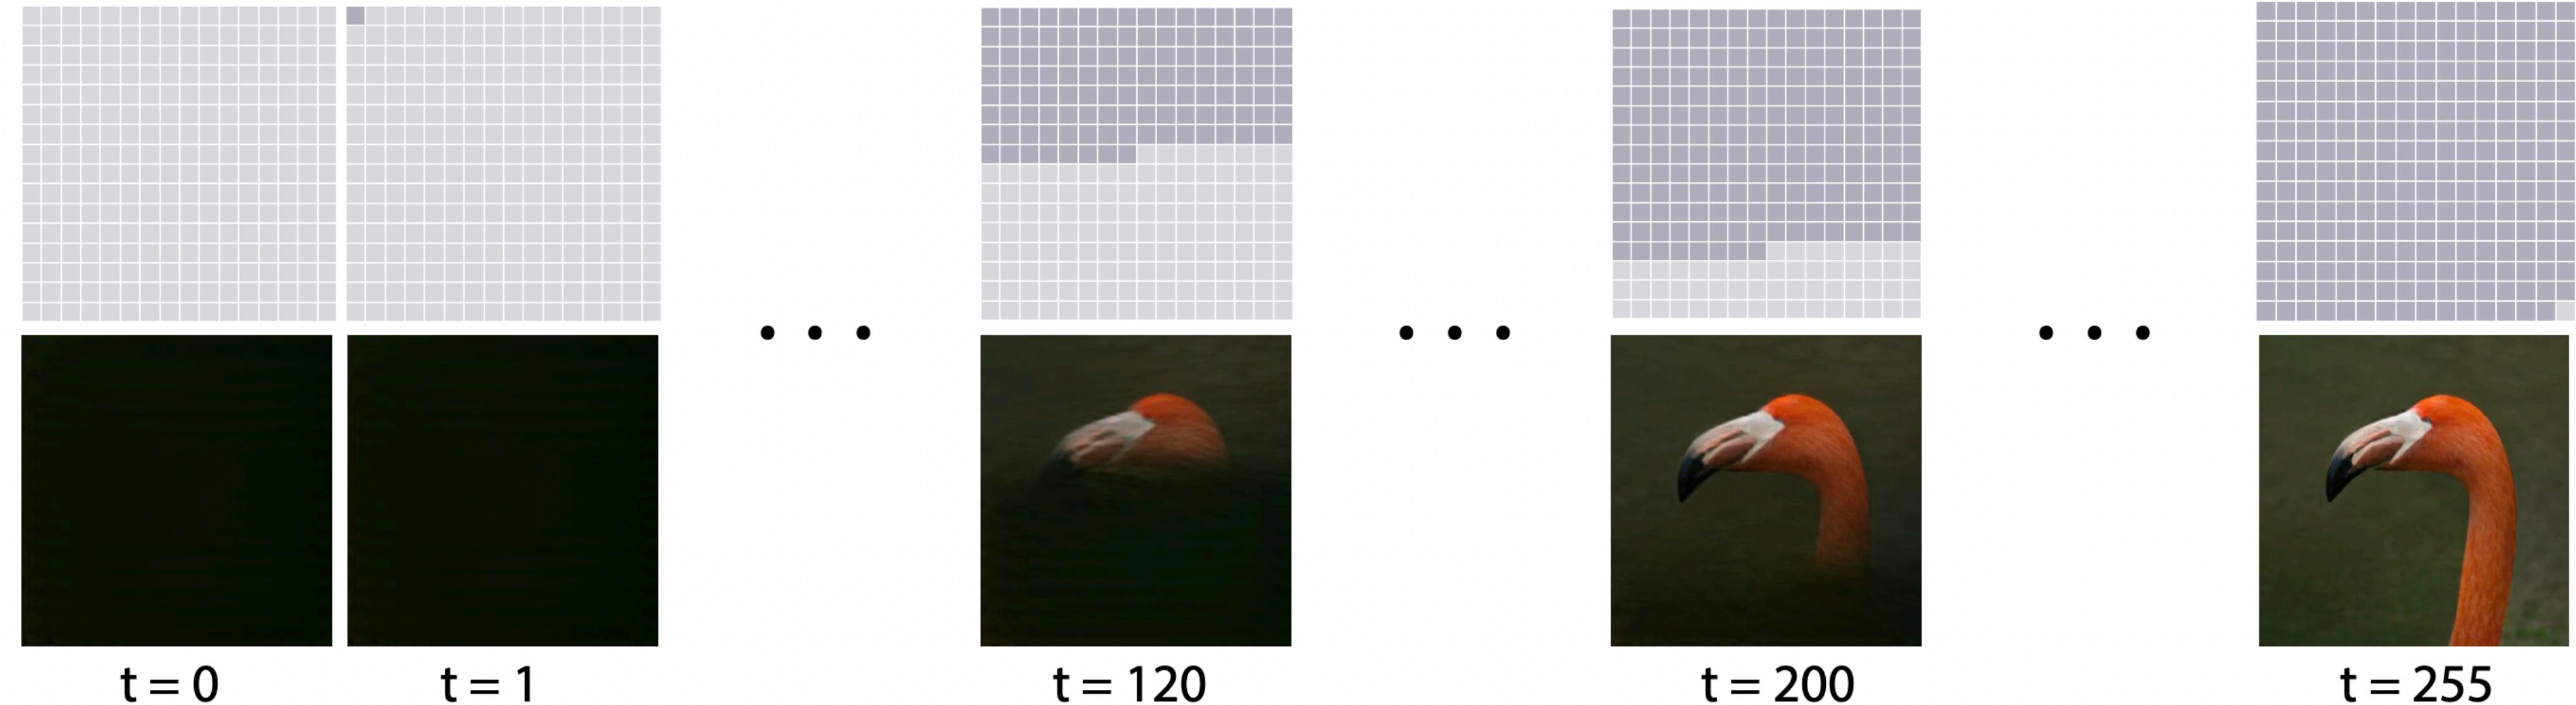
Likelihood of image x Probability of i 'th pixel value given all previous pixels



Fully-visible Belief Networks (FVBNs)

Autoregressive Decoders

Sequential
Decoding
with Autoregressive
Transformers



Chang et al. (2022), [MaskGIT: Masked Generative Image Transformer](#)

Fully-visible Belief Networks (FVBNs)

Autoregressive Decoders

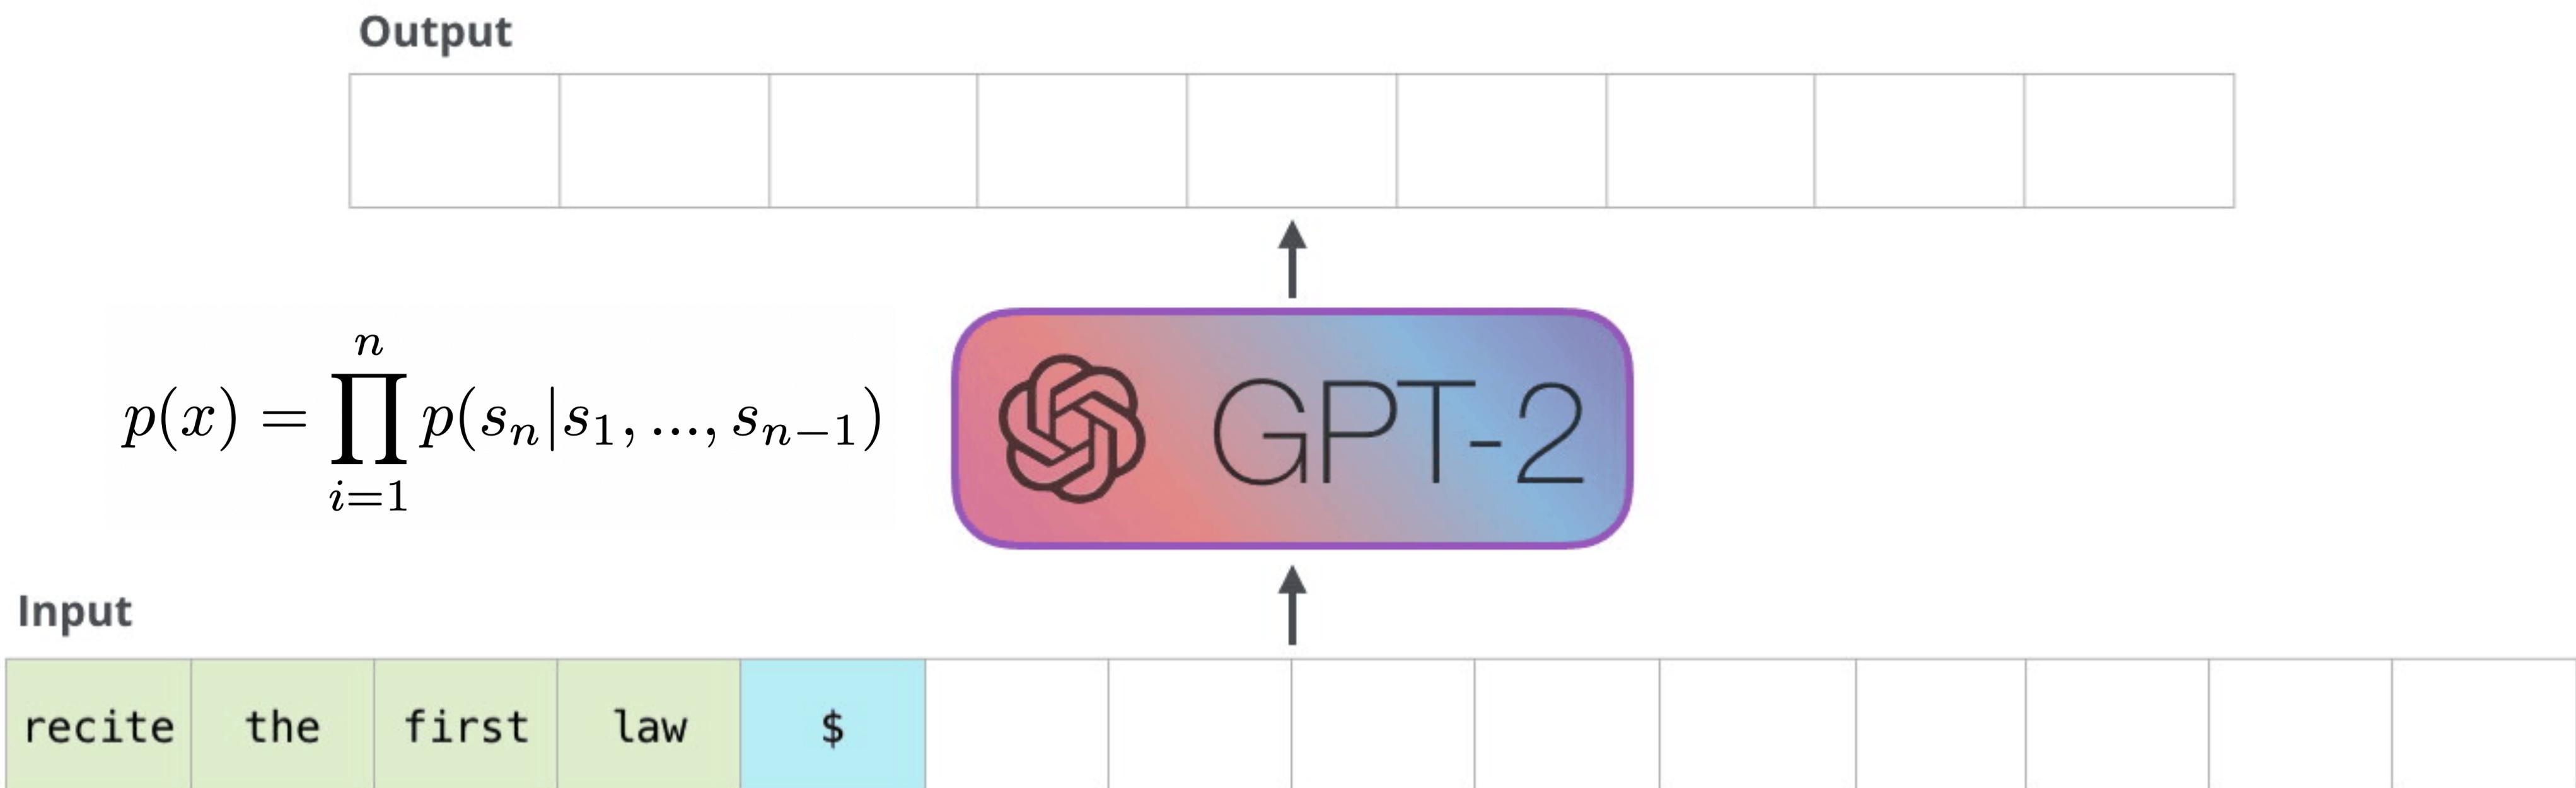


Figure taken from <https://jalammar.github.io/illustrated-gpt2/>

Pros & Cons of Autoregressive Decoders

- Pros:
 - $p(x)$ can be computed explicitly and exactly
 - Easy optimization via single objective function
 - High output quality
- Cons:
 - Slow $O(n)$ decoding, n = output sequence length
 - Not controlled by latent variable (z)

Today's Lecture

1. Supervised vs Unsupervised Learning
2. Generative Modeling
3. Applications of generative models
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
- 5. (Variational) Autoencoders**
6. Generative Adversarial Networks (GANs)

Variational Autoencoders (VAEs)

- Autoregressive decoders define a tractable density function, directly optimised through MLE in the form of gradient descent

$$p_{\theta}(x) = \prod_{i=1}^n p_{\theta}(x_i|x_1, \dots, x_{i-1})$$

- VAEs define an intractable density function with latent variable z

$$p_{\theta}(x) = \int p_{\theta}(z)p_{\theta}(x|z)dz$$

**Yay, no chain dependencies,
so we can sample the whole
thing at once :))**

**Oh no, how are we supposed
to optimize this though :((**

**And what is a latent variable
 z ??**

Autoencoders

- Unsupervised model
- Learns lower-dimensional feature representation \mathbf{z} from **unlabelled** training data \mathbf{x}
- Bottleneck architecture, i.e., $\dim(\mathbf{z}) < \dim(\mathbf{x})$

Why?

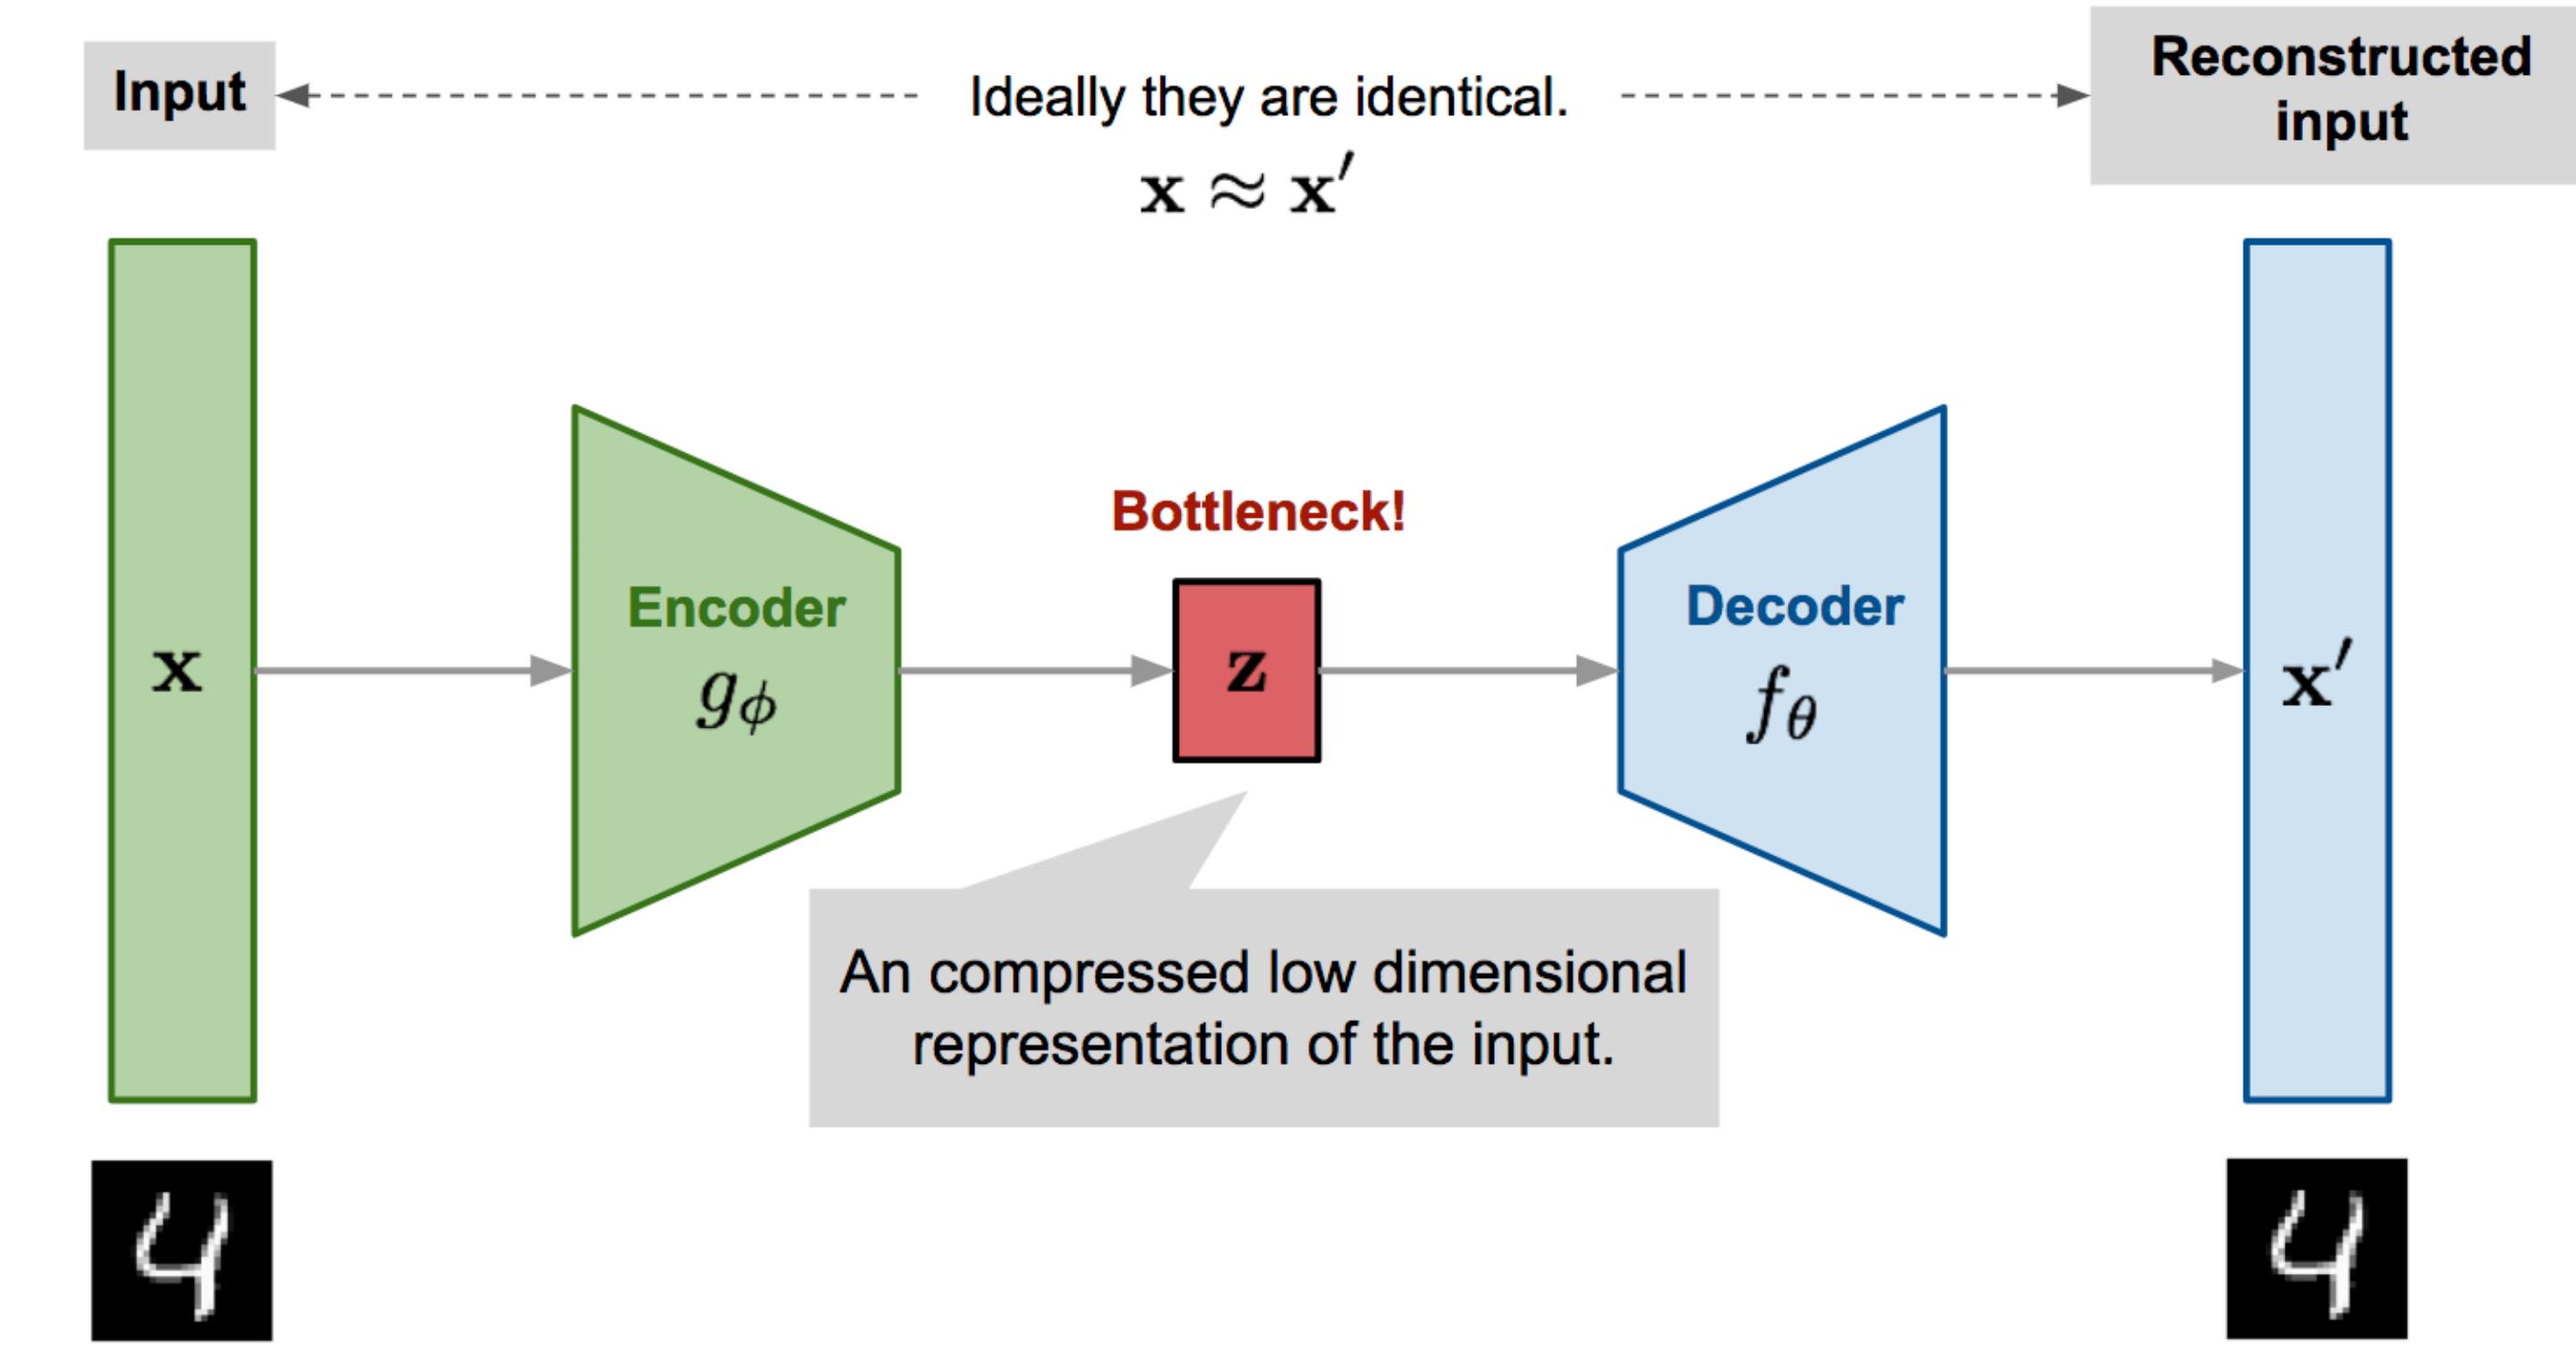
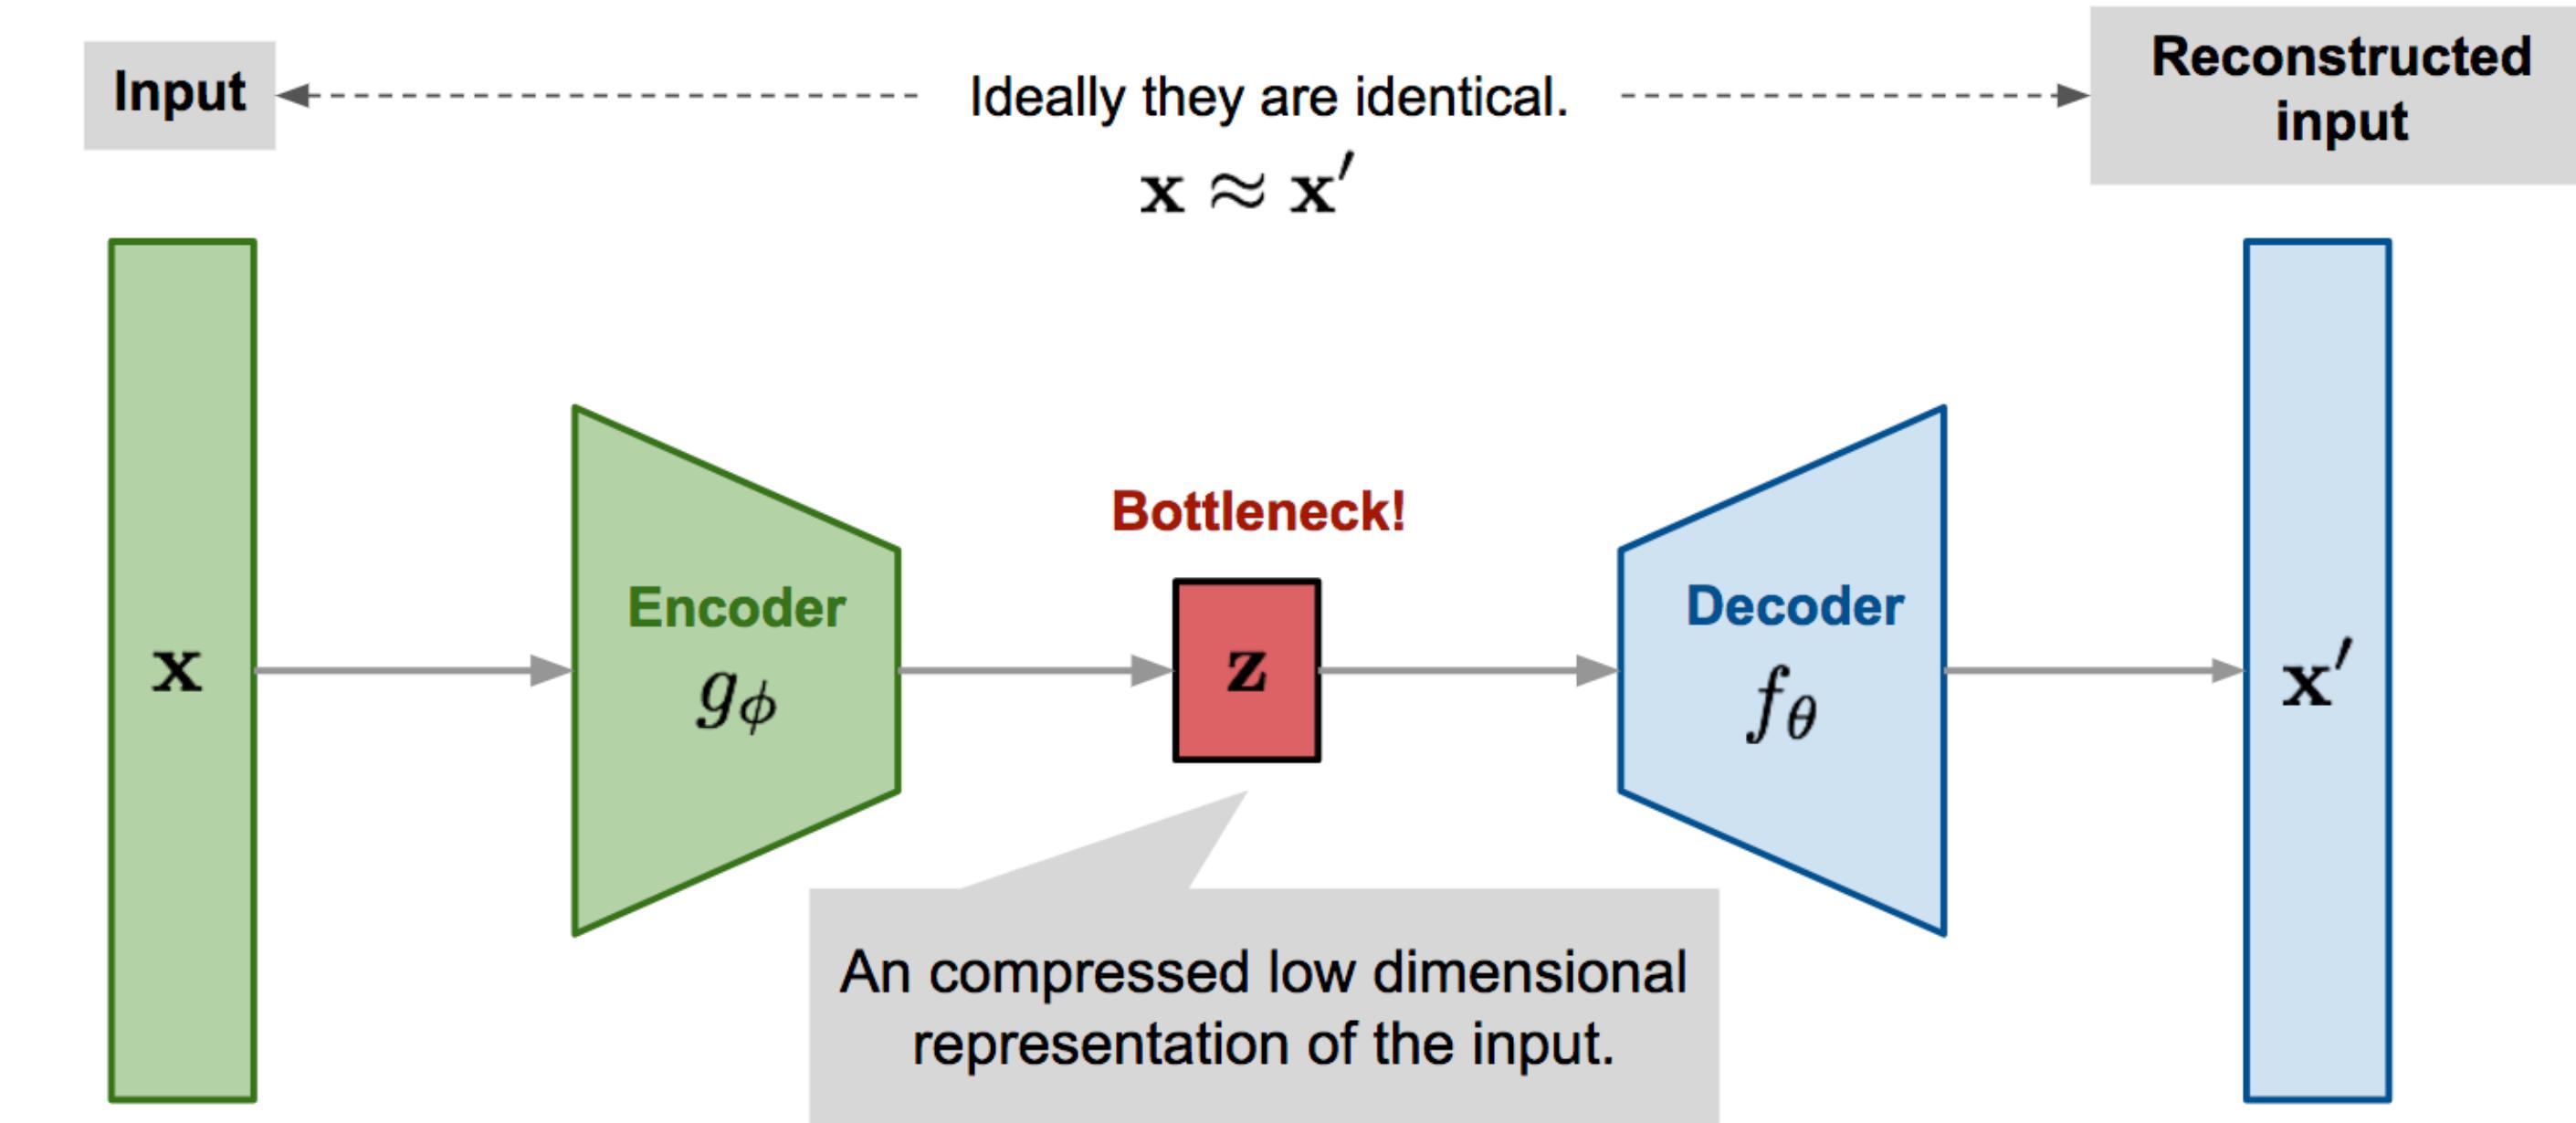


Figure taken from <https://lilianweng.github.io/posts/2018-08-12-vae/>

Autoencoders

- Unsupervised model
- Learns lower-dimensional feature representation \mathbf{z} from **unlabelled** training data \mathbf{x}
- Bottleneck architecture, i.e.,
 $\dim(\mathbf{z}) < \dim(\mathbf{x})$



Why?

To learn meaningful factors of variation rather than just memorizing

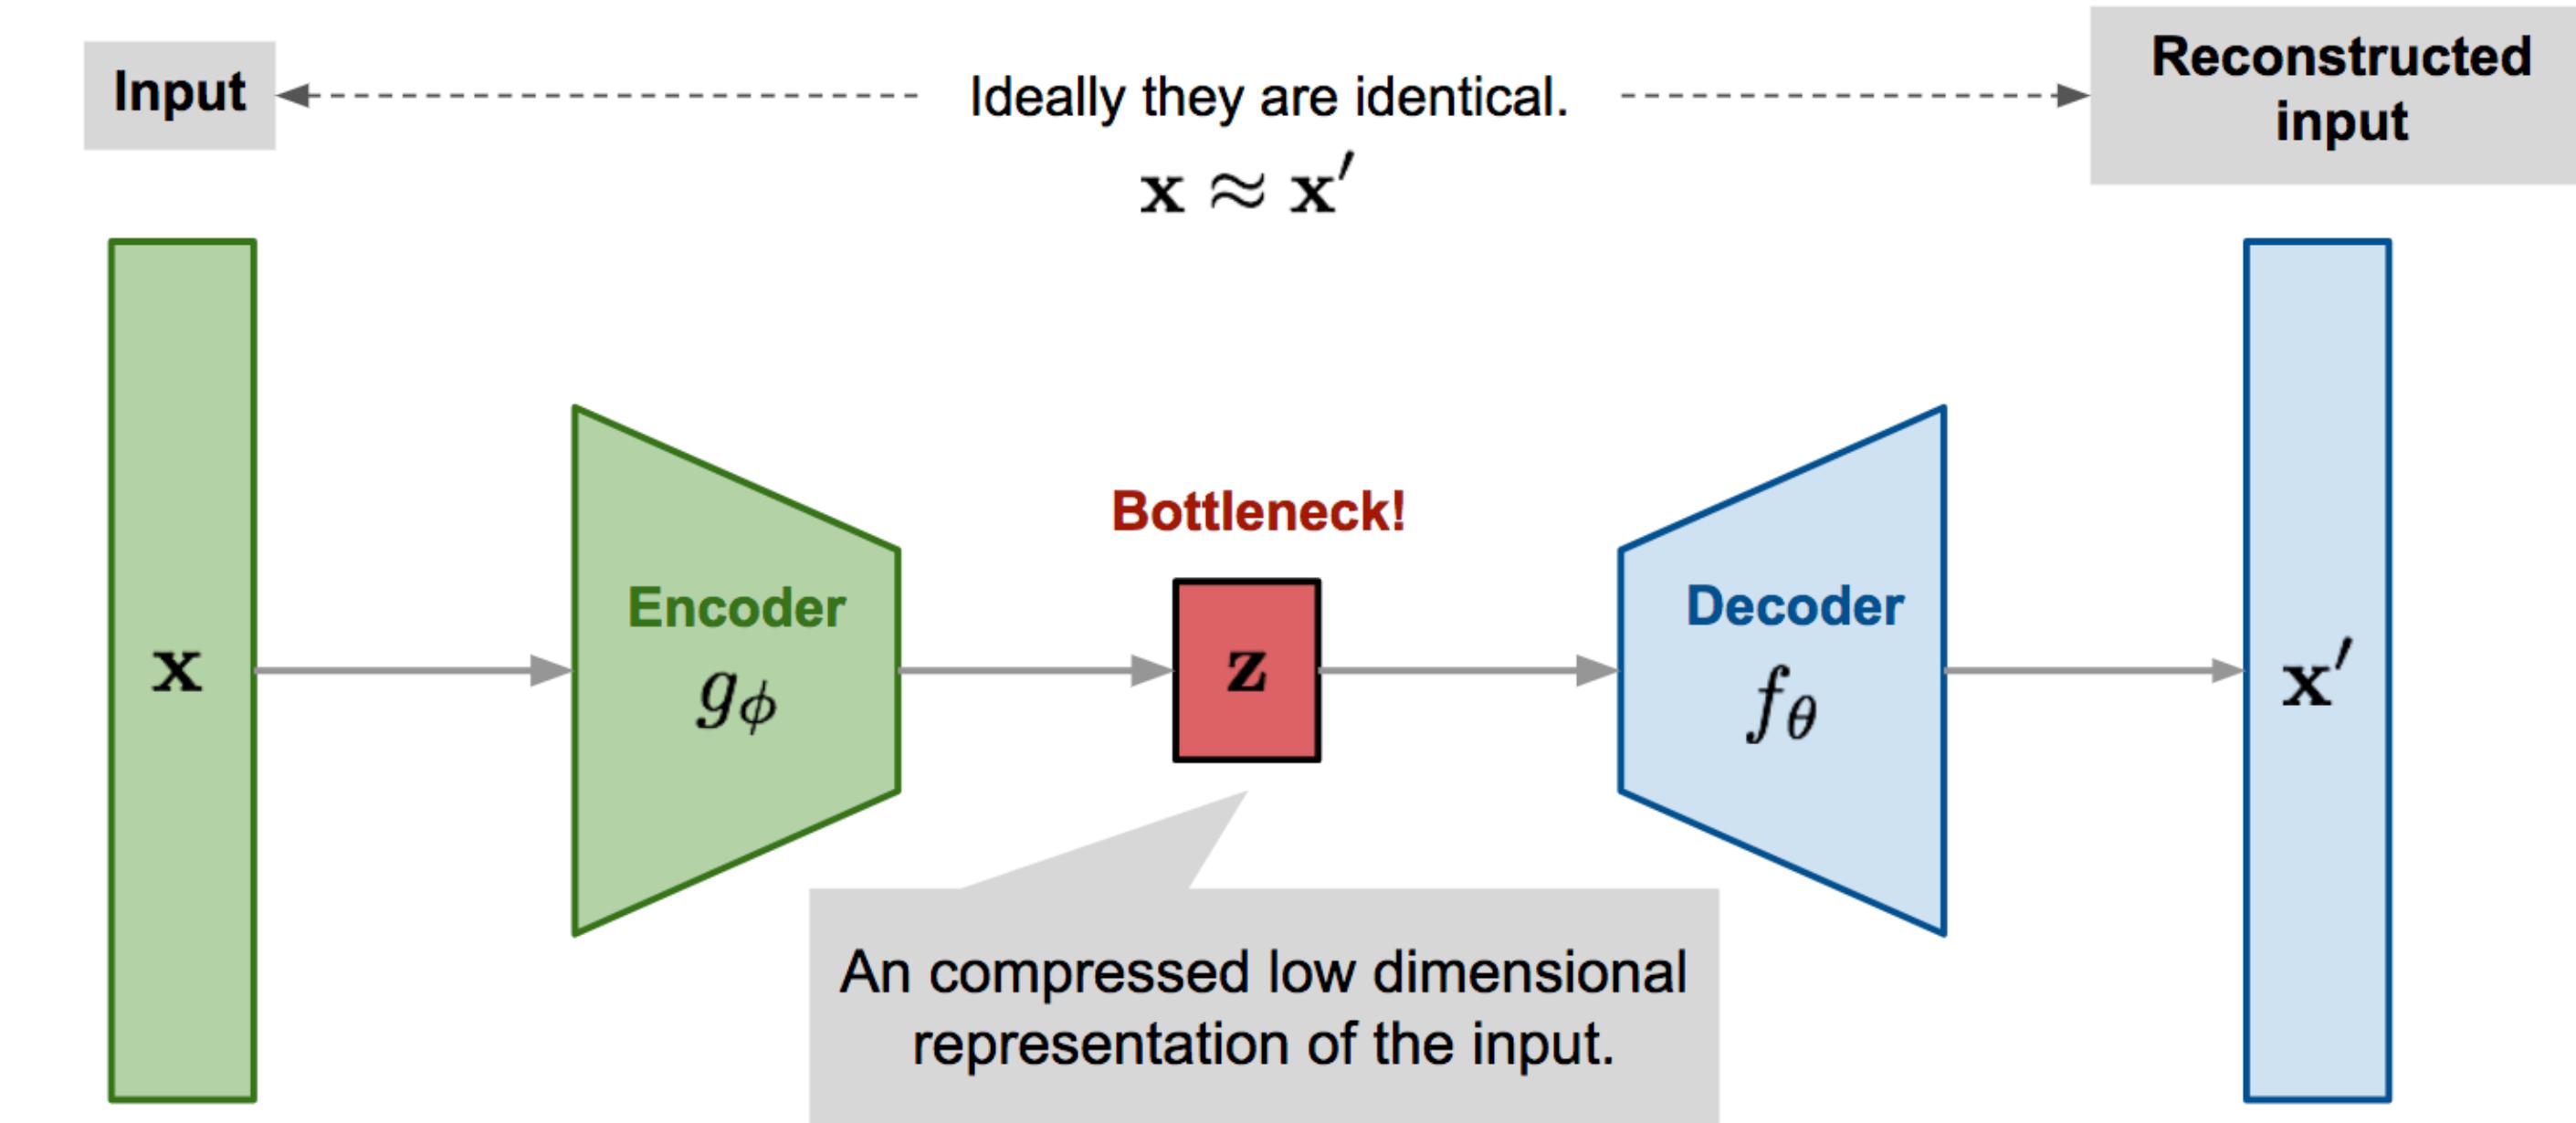
4

4

Figure taken from <https://lilianweng.github.io/posts/2018-08-12-vae/>

Autoencoders

- Unsupervised model
- Learns lower-dimensional feature representation \mathbf{z} from **unlabelled** training data \mathbf{x}
- Bottleneck architecture, i.e., $\text{dim}(\mathbf{z}) < \text{dim}(\mathbf{x})$



Why?

To learn meaningful factors of variation rather than just memorizing

4

4

Figure taken from <https://lilianweng.github.io/posts/2018-08-12-vae/>

$$L_{\text{AE}}(\theta, \phi) = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}^{(i)} - f_\theta(g_\phi(\mathbf{x}^{(i)})))^2$$

Autoencoders

- Unsupervised model
- Learns lower-dimensional feature representation \mathbf{z} from **unlabelled** training data \mathbf{x}
- Bottleneck architecture, i.e.,
 $\dim(\mathbf{z}) < \dim(\mathbf{x})$

Why?

To learn meaningful factors of variation rather than just memorizing

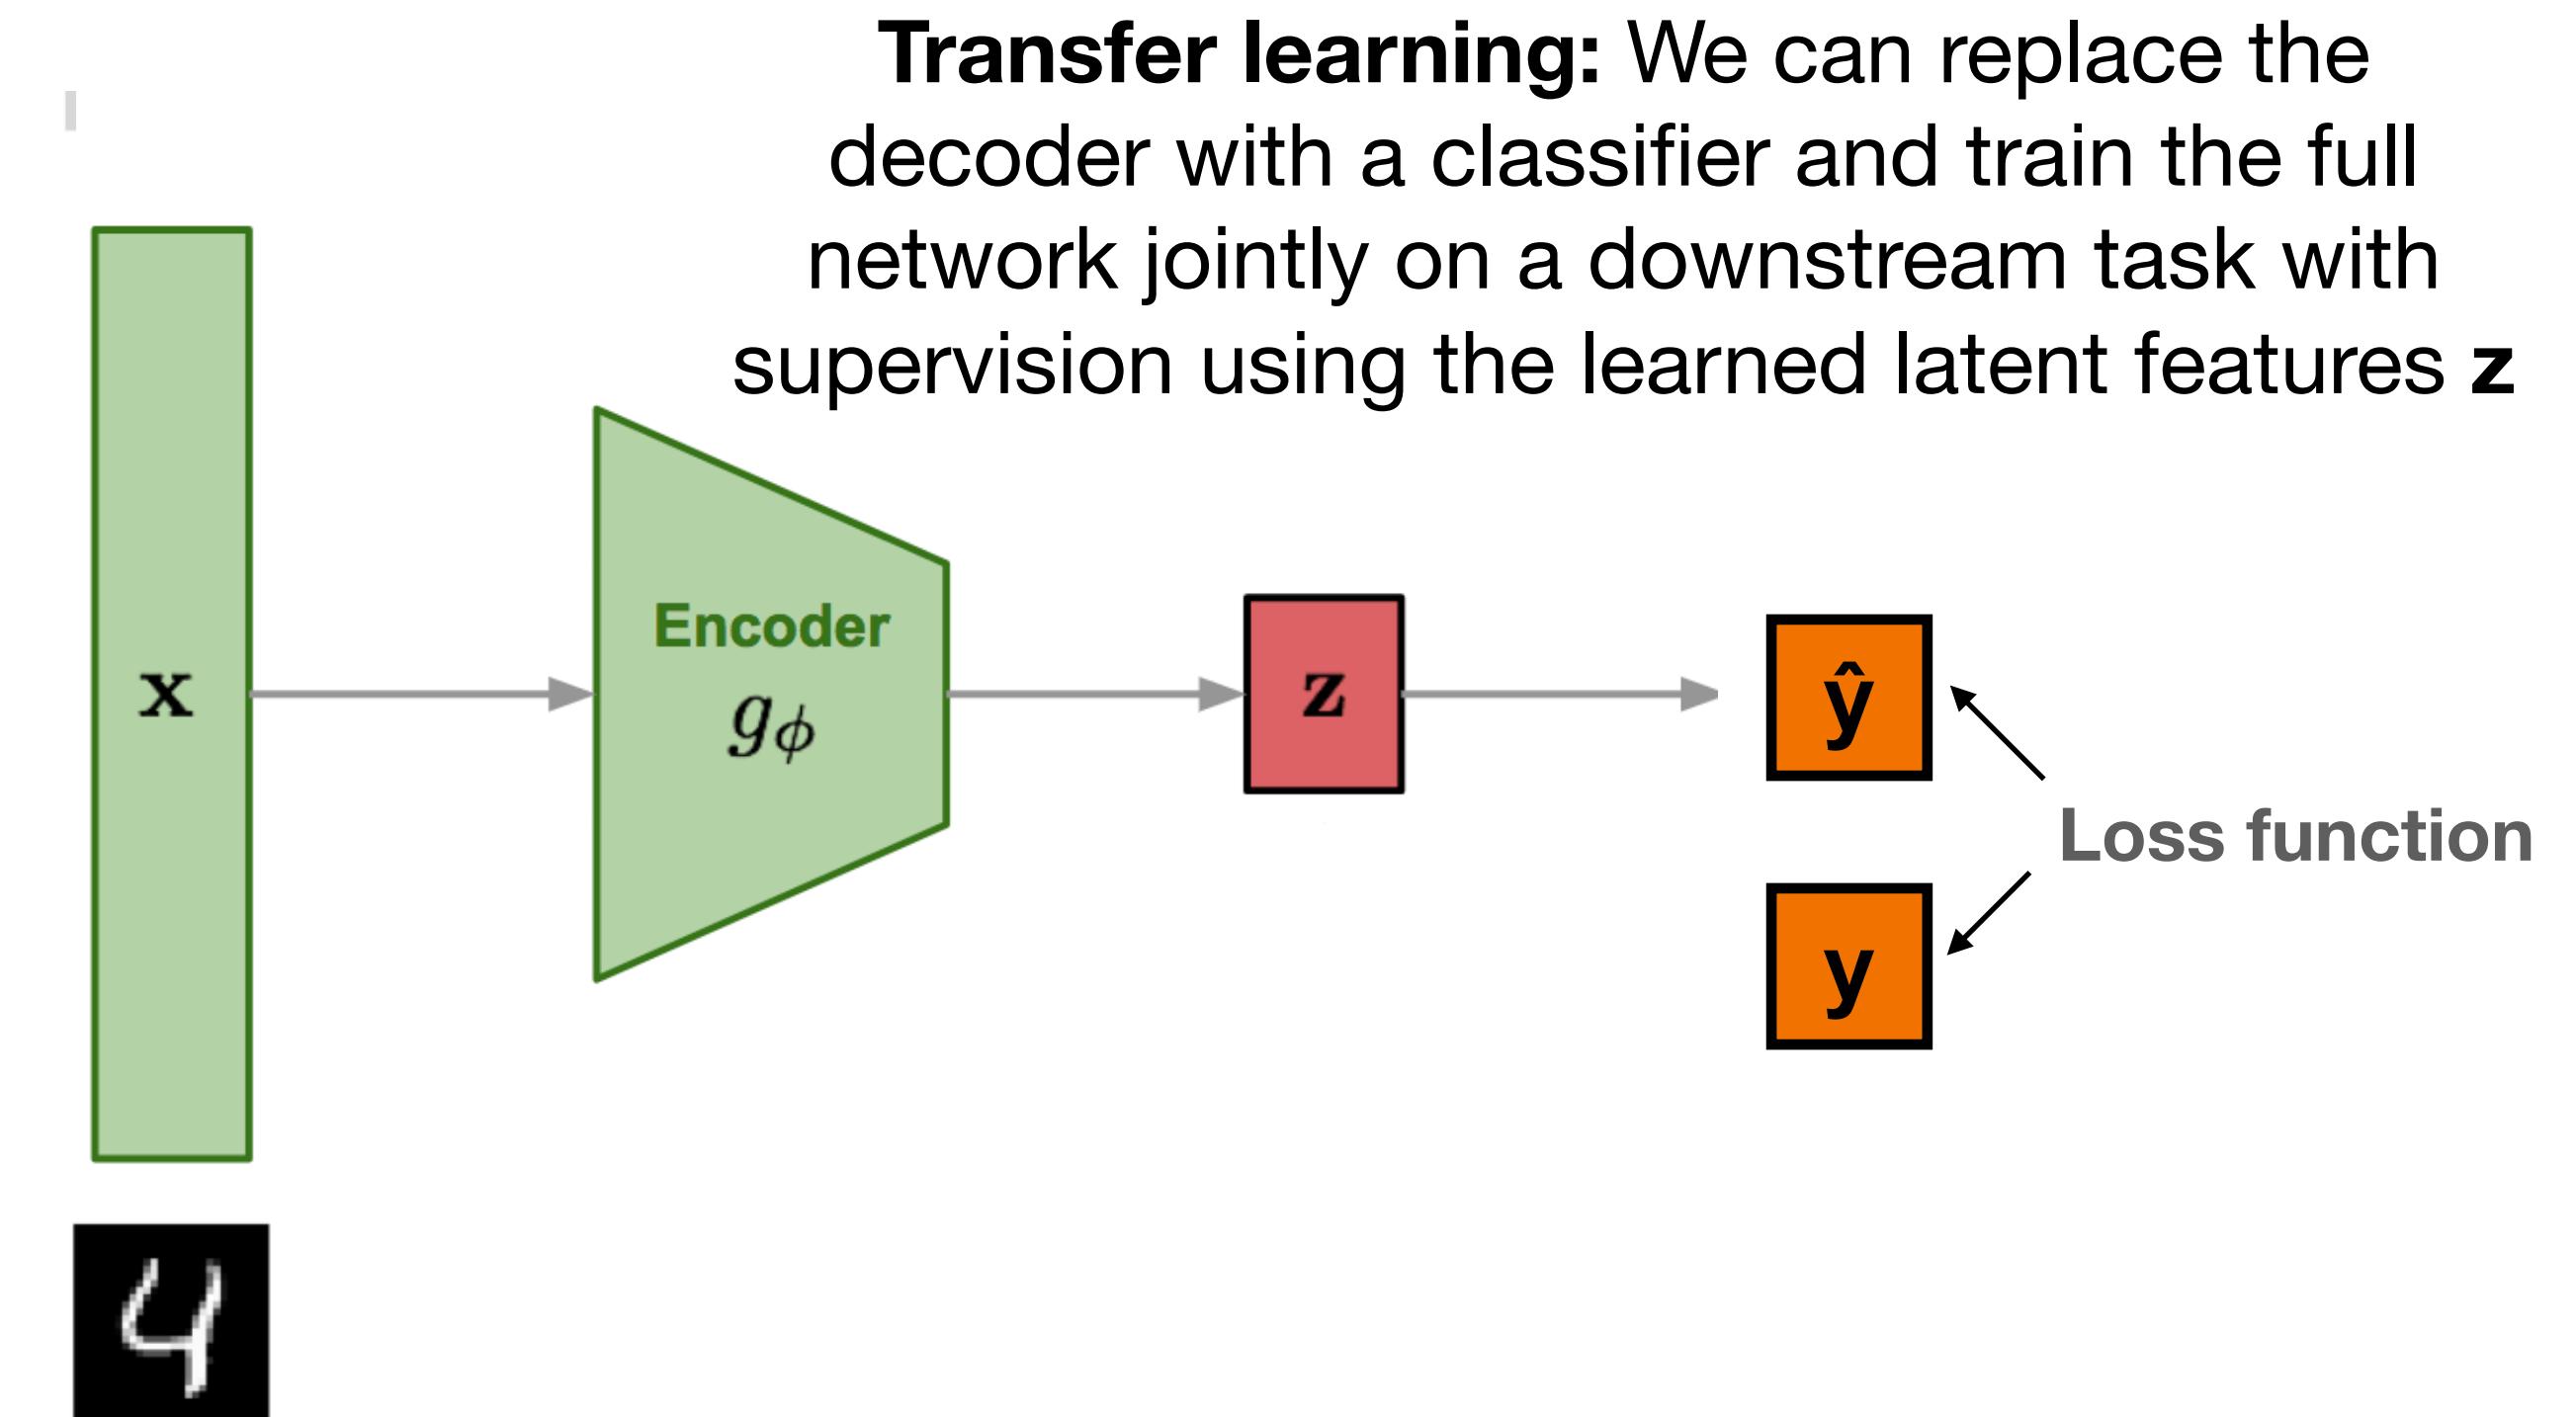


Figure modified from <https://lilianweng.github.io/posts/2018-08-12-vae/>

Variational Autoencoders

- We cannot sample from \mathbf{z} in the vanilla Autoencoder
- To achieve this, let's map our inputs \mathbf{x} to a parametrized distribution $p_\theta(\mathbf{z})$ rather than a *fixed* vector \mathbf{z}
- Relationship between \mathbf{x} and \mathbf{z} defined in Bayesian terms:

Prior $p_\theta(\mathbf{z})$

Likelihood $p_\theta(\mathbf{x}|\mathbf{z})$

Posterior $p_\theta(\mathbf{z}|\mathbf{x})$

Variational Autoencoders

Assuming we have found a good set of parameters θ , we can **generate an example $x^{(i)}$** as follows:

1. Sample $z^{(i)}$ from $p_\theta(z)$
2. Decode $x^{(i)}$ via conditional probability distribution

$$p_\theta(x | z = z^{(i)})$$

The optimal parameters θ^* maximize the probability of generating a real example $x^{(i)}$:

$$\theta^* = \arg \max_{\theta} \prod_{i=1}^n p_\theta(x^{(i)}) \rightarrow \theta^* = \arg \max_{\theta} \sum_{i=1}^n \log p_\theta(x^{(i)})$$

(Use log likelihood because sum is easier to work with than product)

Prior $p_\theta(z)$

Likelihood $p_\theta(x|z)$

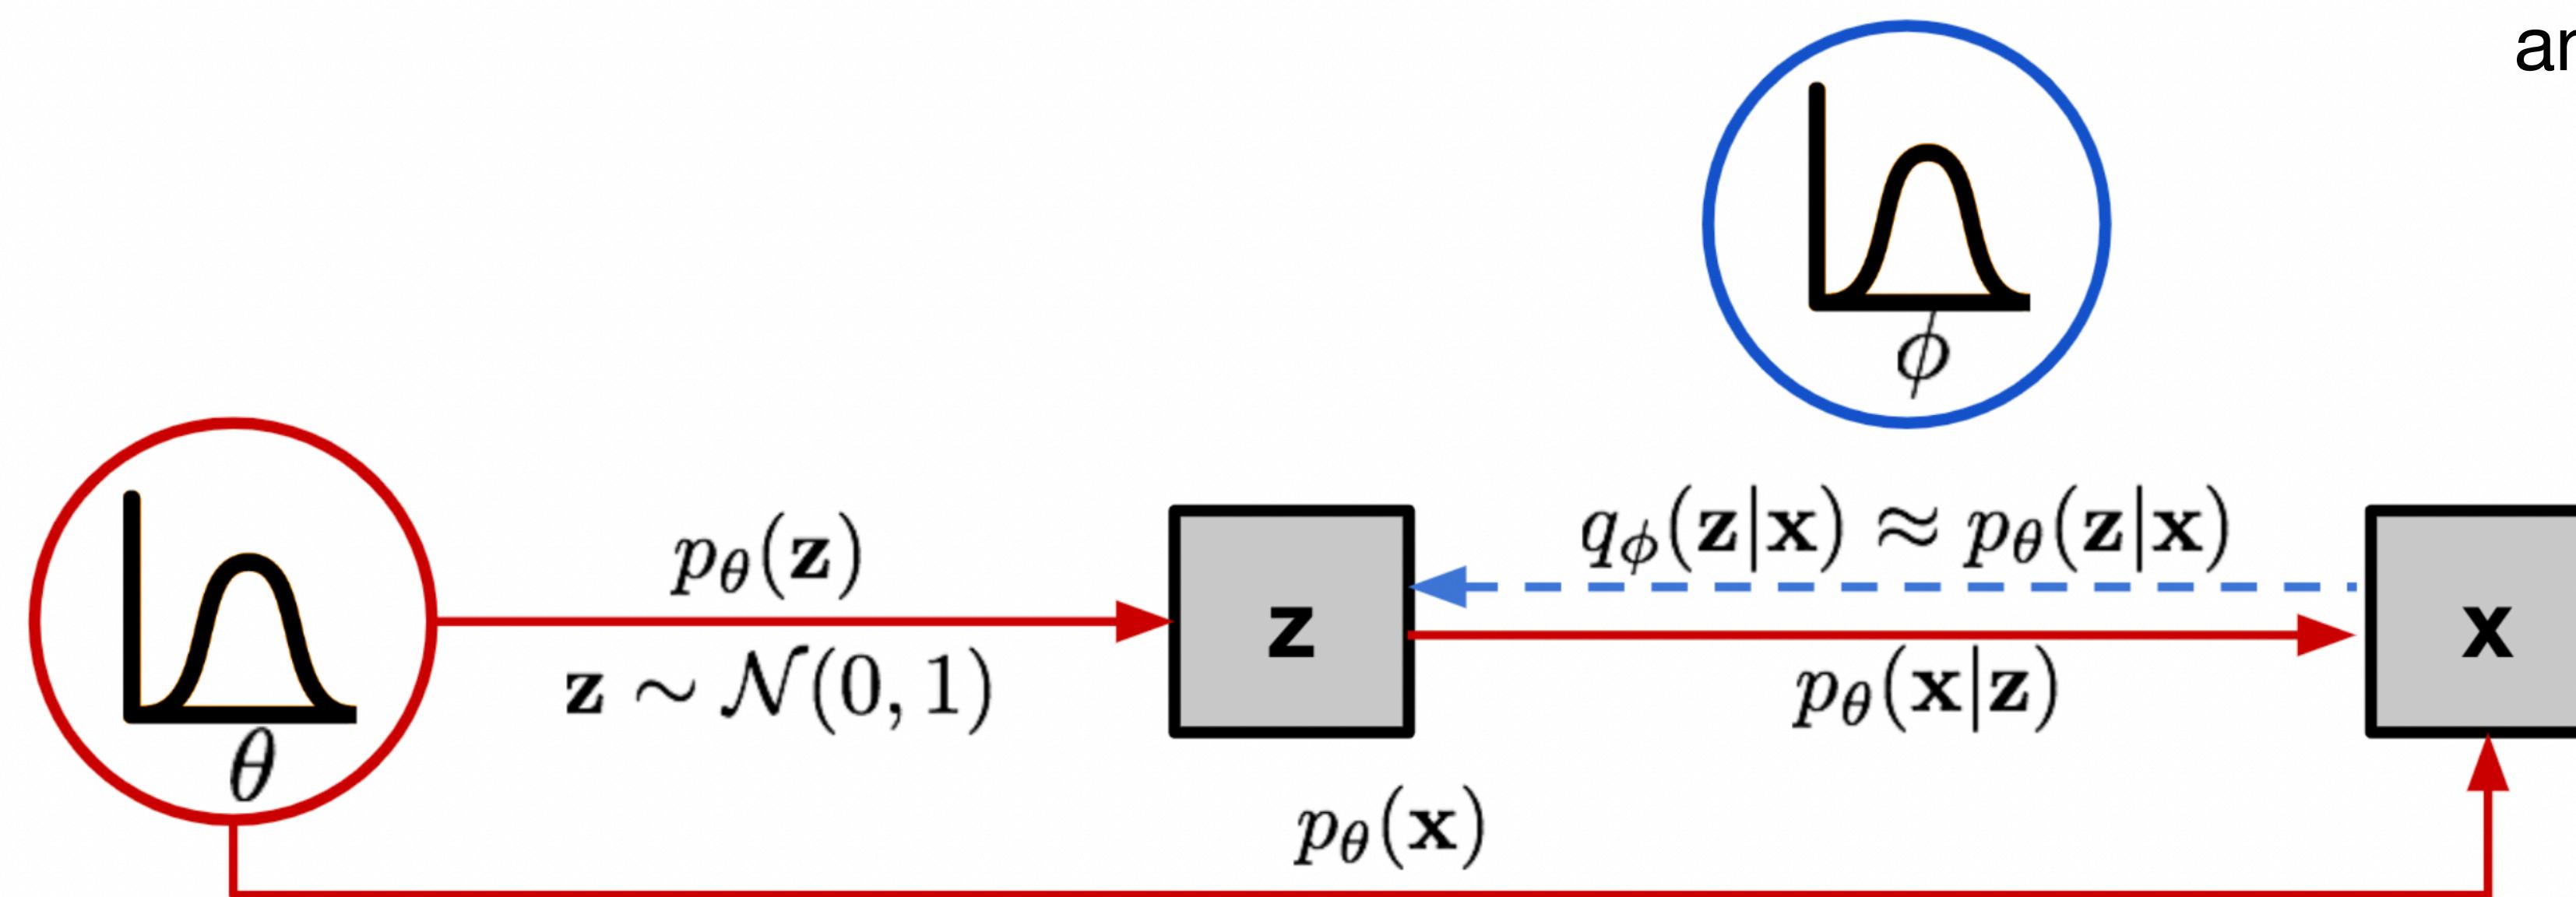
Posterior $p_\theta(z|x)$

$$p_\theta(x^{(i)}) = \int p_\theta(x^{(i)}|z)p_\theta(z)dz$$

This is cool but intractable
because we can't really sum over
all possible z

Variational Autoencoder (VAE)

Graphical Model Illustration



The **integral over z** is **intractable**,
and the **posterior $p_\theta(z|x)$** is also
intractable

Therefore, we approximate the posterior via $q_\phi(z|x)$

We call $q_\phi(z|x)$ the **probabilistic encoder**
and $p_\theta(x|z)$ the **probabilistic decoder**

Fig. 6. The graphical model involved in Variational Autoencoder. Solid lines denote the generative distribution $p_\theta(\cdot)$ and dashed lines denote the distribution $q_\phi(z|x)$ to approximate the intractable posterior $p_\theta(z|x)$.

Variational Autoencoder (VAE)

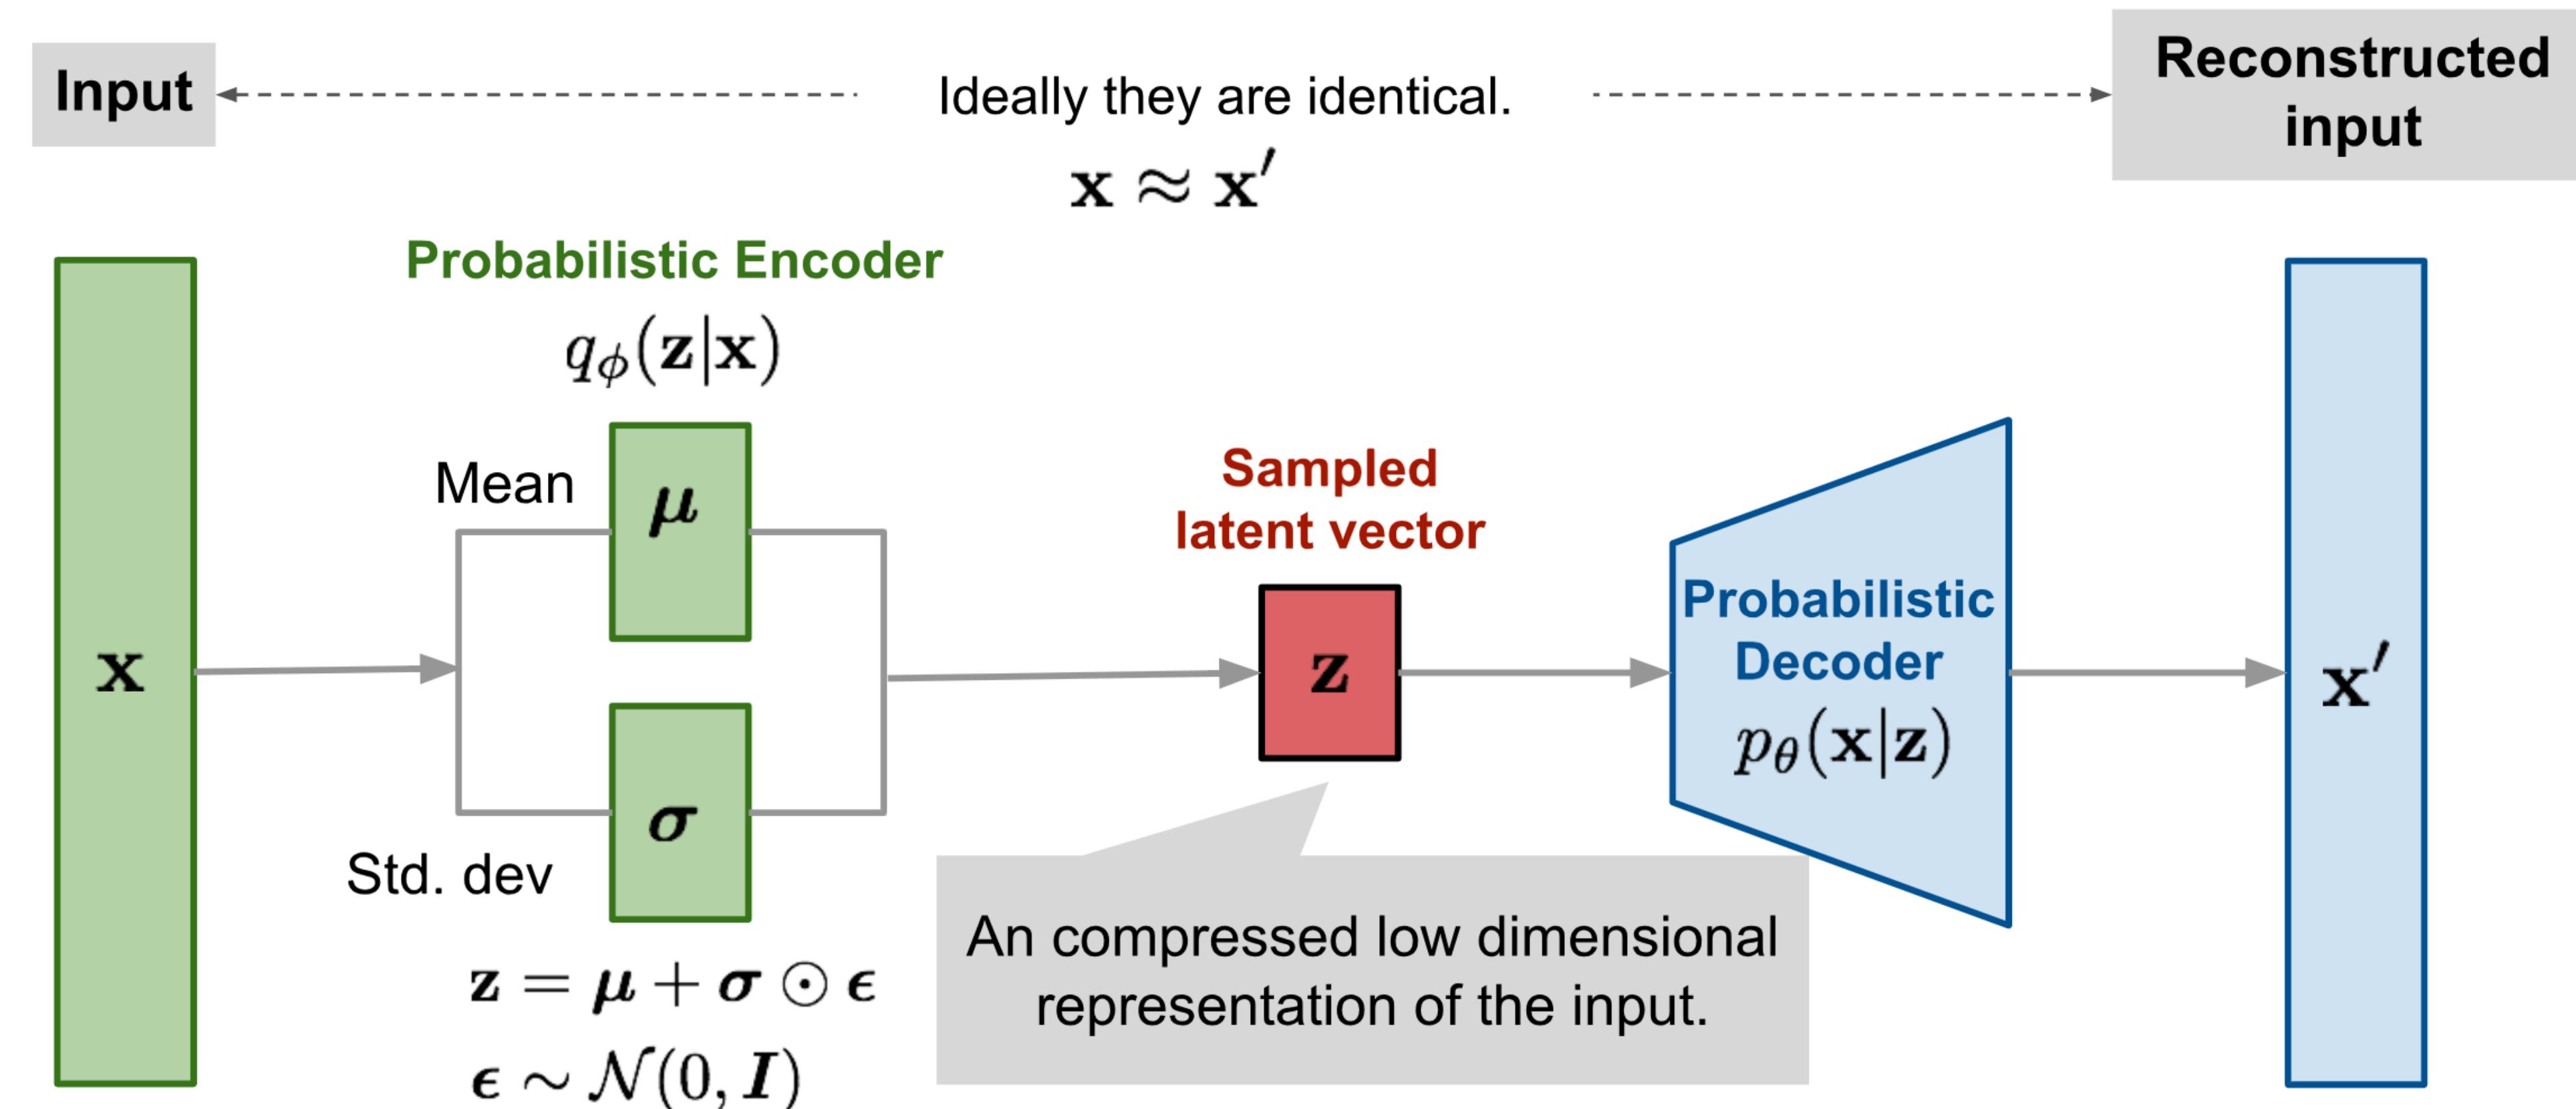


Figure taken from <https://lilianweng.github.io/posts/2018-08-12-vae/>

Variational Autoencoder (VAE)

Loss Function – Evidence Lower Bound (ELBO)

We need our probabilistic decoder distribution $q_\phi(z | x)$ to be close to the posterior distribution $p_\theta(z | x)$

(Otherwise our latent space may be something completely different)

We can make the distributions similar by **minimizing** the Kullback-Leibler (KL) Divergence

$$D_{\text{KL}}(q_\phi(z|x) || p_\theta(z|x))$$

*KL-Divergence quantifies how much information is lost when representing the left-hand distribution by the right-hand distribution (it is not symmetrical!)

Variational Autoencoder (VAE)

Loss Function – Evidence Lower Bound (ELBO)

$$\begin{aligned} & D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p_{\theta}(\mathbf{z}|\mathbf{x})) \\ &= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}|\mathbf{x})} d\mathbf{z} \\ &= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})p_{\theta}(\mathbf{x})}{p_{\theta}(\mathbf{z}, \mathbf{x})} d\mathbf{z} && ; \text{ Because } p(z|x) = p(z,x)/p(x) \\ &= \int q_{\phi}(\mathbf{z}|\mathbf{x}) \left(\log p_{\theta}(\mathbf{x}) + \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}, \mathbf{x})} \right) d\mathbf{z} \\ &= \log p_{\theta}(\mathbf{x}) + \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z}, \mathbf{x})} d\mathbf{z} && ; \text{ Because } \int q(z|x) dz = 1 \\ &= \log p_{\theta}(\mathbf{x}) + \int q_{\phi}(\mathbf{z}|\mathbf{x}) \log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{x}|\mathbf{z})p_{\theta}(\mathbf{z})} d\mathbf{z} && ; \text{ Because } p(z,x) = p(x|z)p(z) \\ &= \log p_{\theta}(\mathbf{x}) + \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} [\log \frac{q_{\phi}(\mathbf{z}|\mathbf{x})}{p_{\theta}(\mathbf{z})} - \log p_{\theta}(\mathbf{x}|\mathbf{z})] \\ &= \log p_{\theta}(\mathbf{x}) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x}) \| p_{\theta}(\mathbf{z})) - \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) \end{aligned}$$

Variational Autoencoder (VAE)

Loss Function – Evidence Lower Bound (ELBO)

We end up with this equation:

$$D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}|\mathbf{x})) = \log p_{\theta}(\mathbf{x}) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z})) - \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z})$$

Which we can arrange to the following one:

$$\boxed{\log p_{\theta}(\mathbf{x}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}|\mathbf{x}))} = \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}))$$

This is exactly what we want to maximize:

The log likelihood of our data \mathbf{x}
minus

the difference between the probabilistic decoder
and the true posterior distributions



We can maximize the left-hand side by maximizing the right-hand side, or in other words, we can take the negative of the right-hand side as our loss function and minimize it

Variational Autoencoder (VAE)

Loss Function – Evidence Lower Bound (ELBO)

We end up with this equation:

$$D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}|\mathbf{x})) = \log p_{\theta}(\mathbf{x}) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z})) - \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z})$$

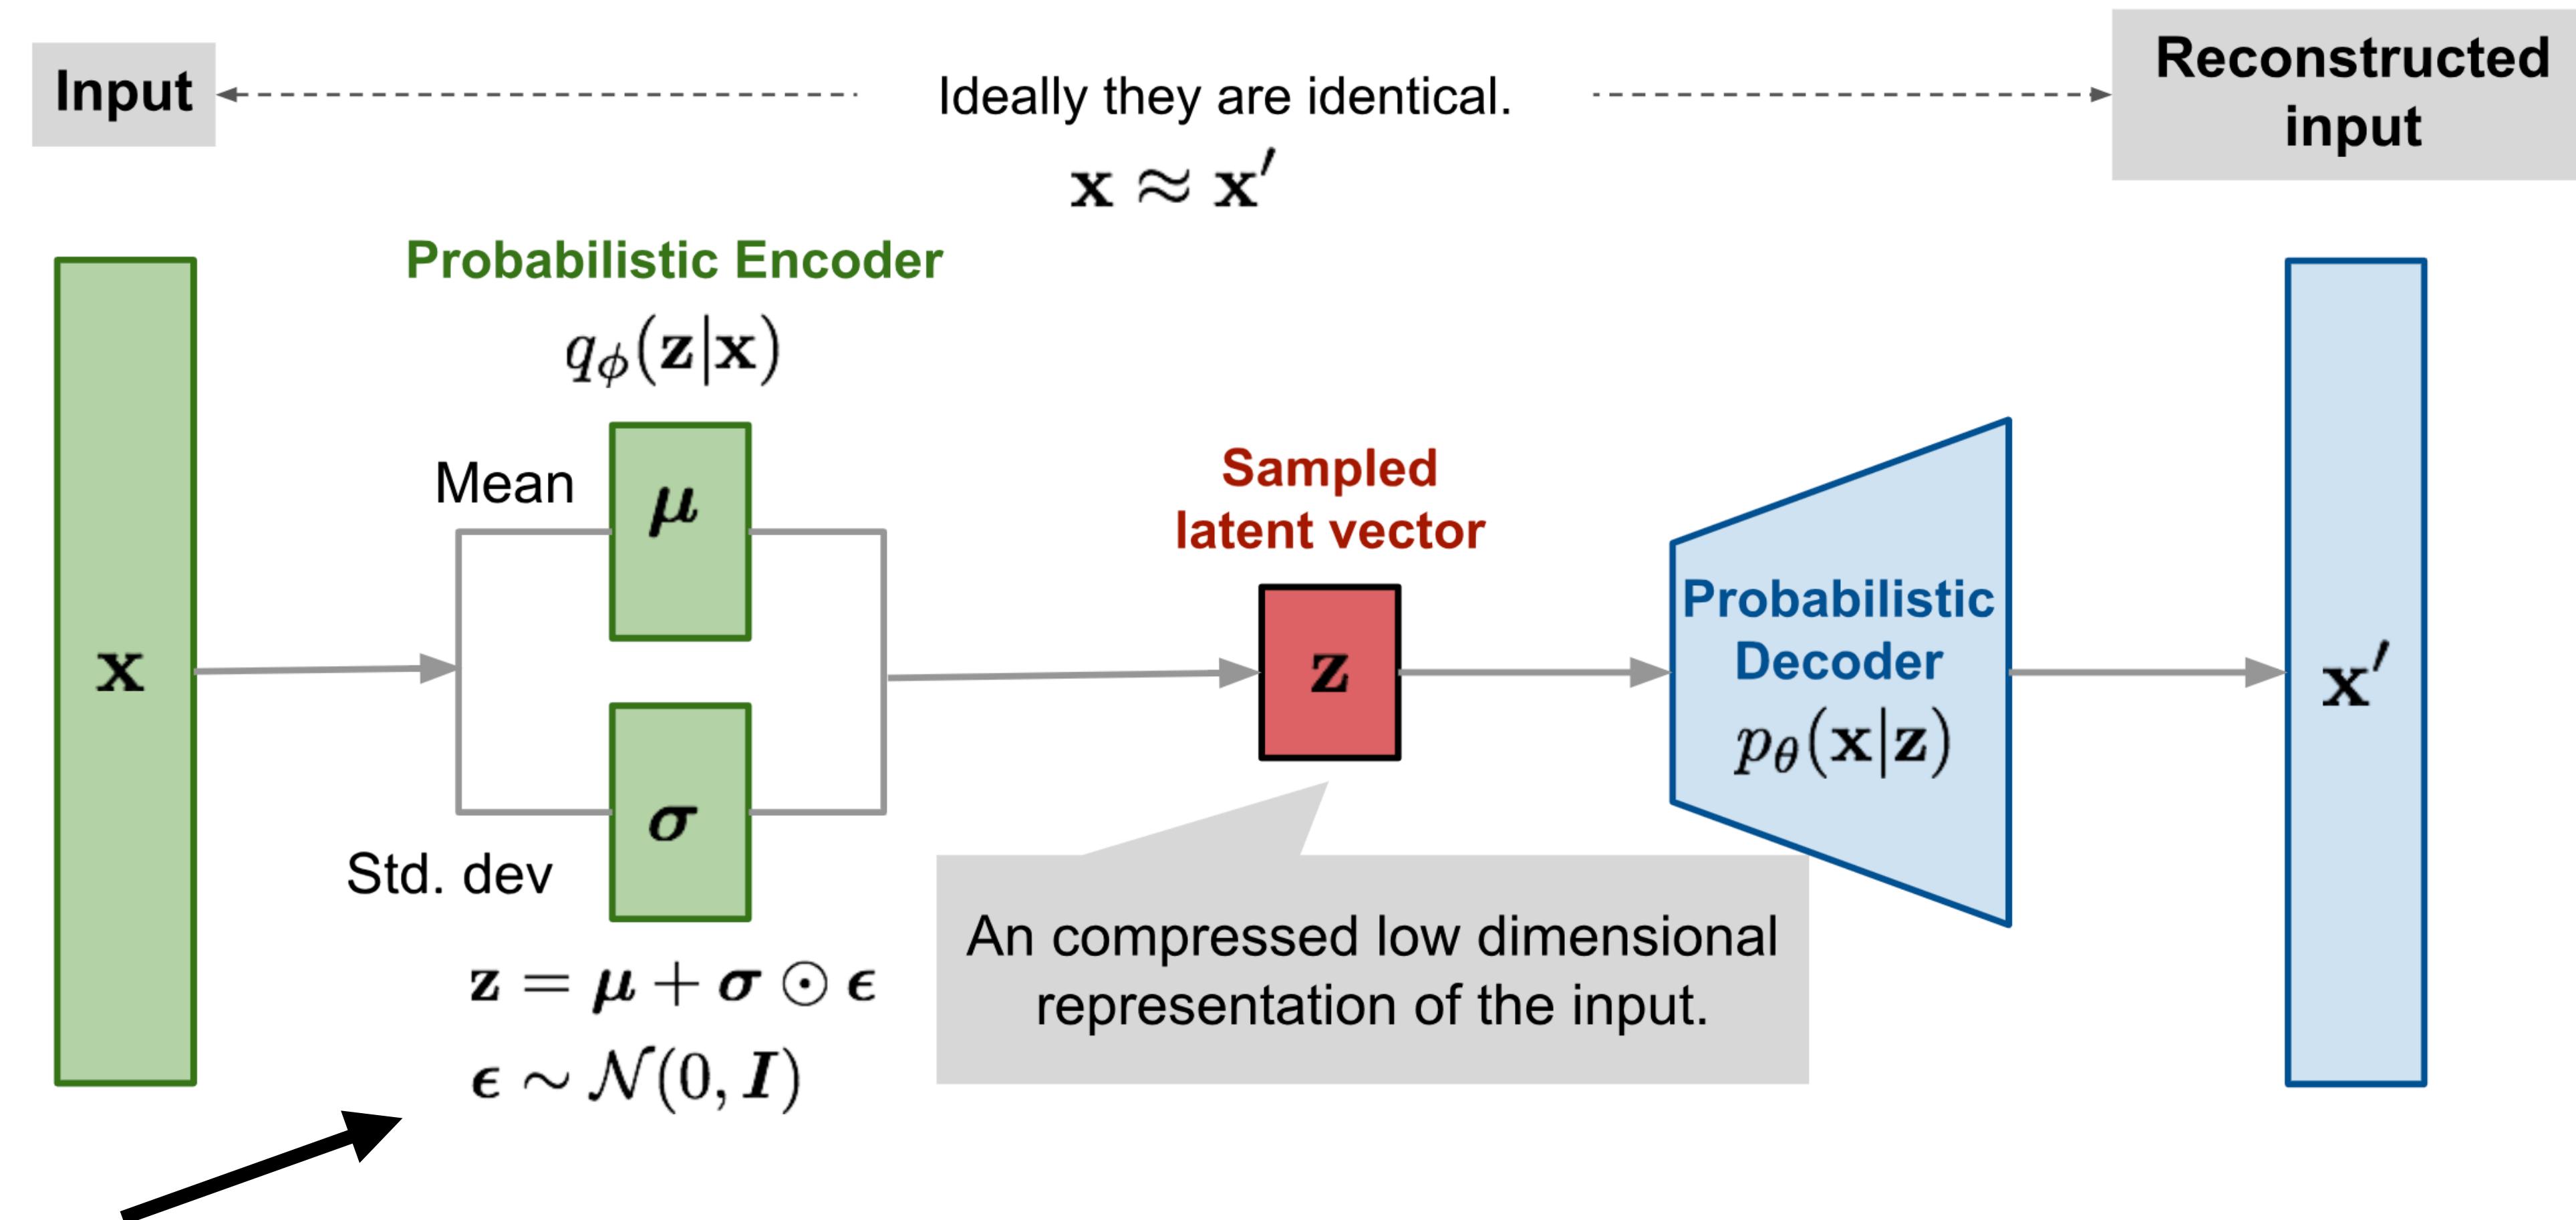
Which we can arrange to the following one:

$$\log p_{\theta}(\mathbf{x}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}|\mathbf{x})) = \mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) - D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}))$$



$$\begin{aligned} L_{\text{VAE}}(\theta, \phi) &= -\log p_{\theta}(\mathbf{x}) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z}|\mathbf{x})) \\ &= -\mathbb{E}_{\mathbf{z} \sim q_{\phi}(\mathbf{z}|\mathbf{x})} \log p_{\theta}(\mathbf{x}|\mathbf{z}) + D_{\text{KL}}(q_{\phi}(\mathbf{z}|\mathbf{x})\|p_{\theta}(\mathbf{z})) \\ \theta^*, \phi^* &= \arg \min_{\theta, \phi} L_{\text{VAE}} \end{aligned}$$

Variational Autoencoder (VAE)

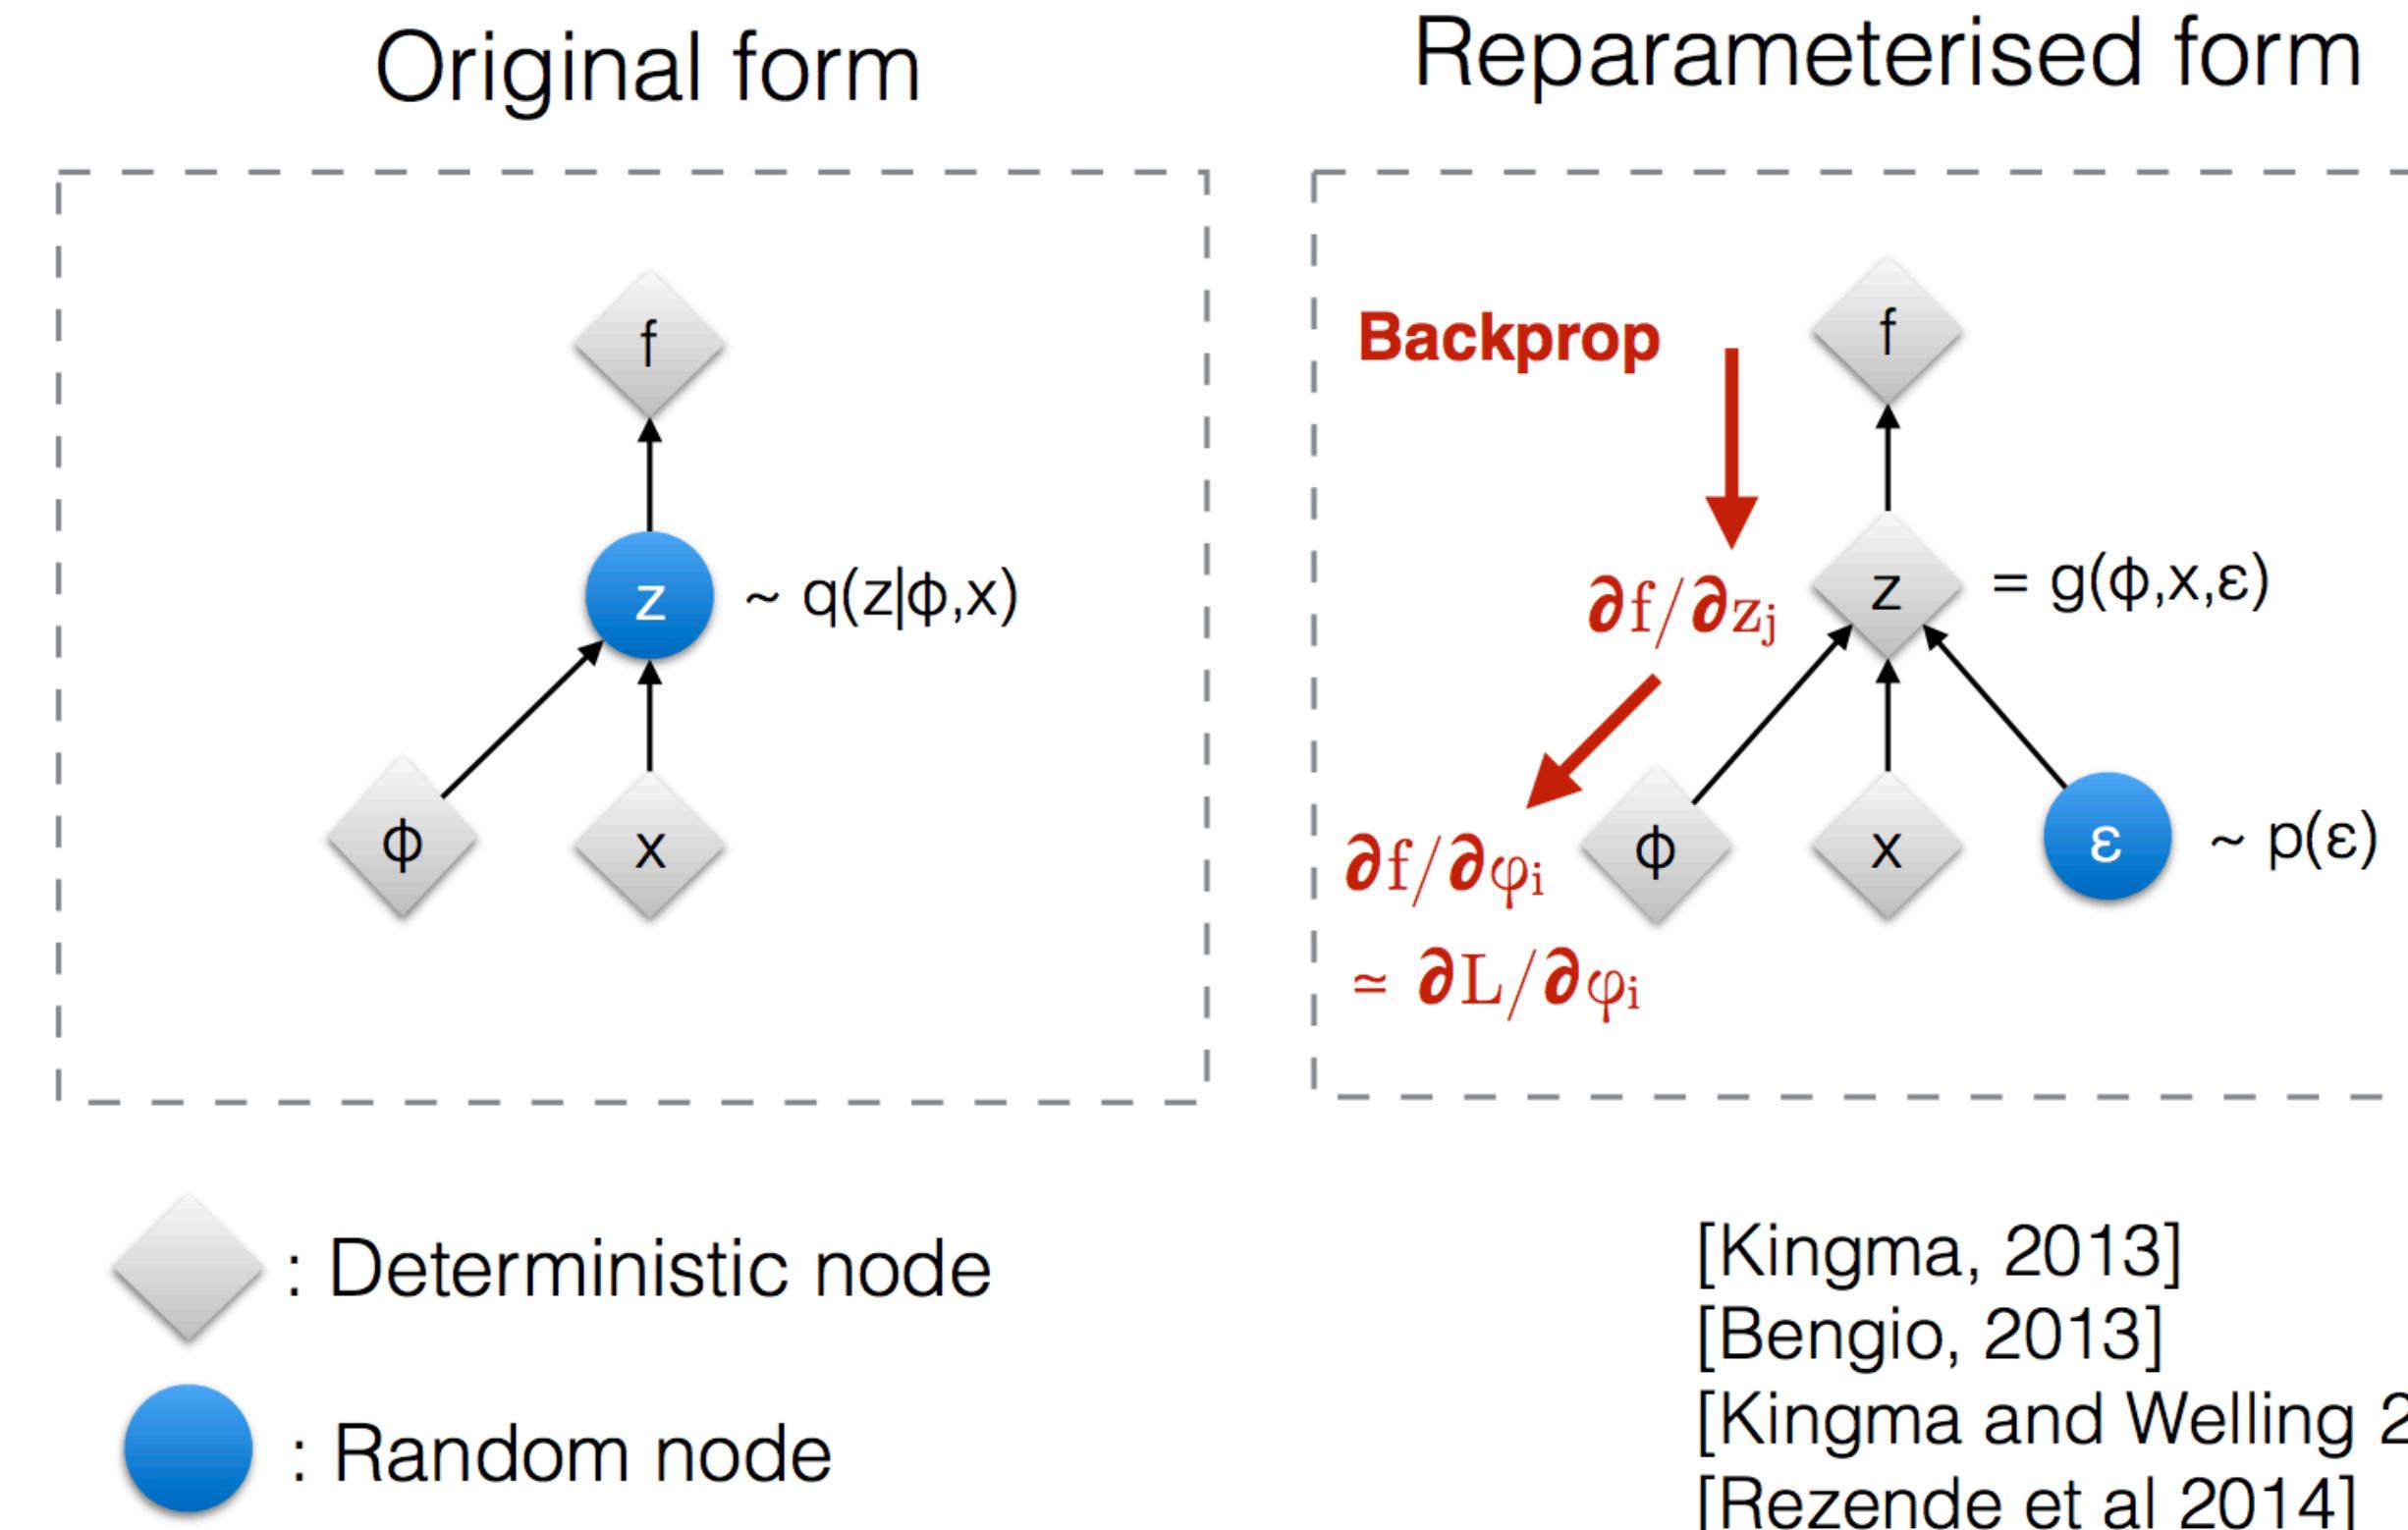


Note: **Sampling z is not a differentiable operation**, so we apply the **reparametrization trick**

Figure taken from <https://lilianweng.github.io/posts/2018-08-12-vae/>

Variational Autoencoder (VAE)

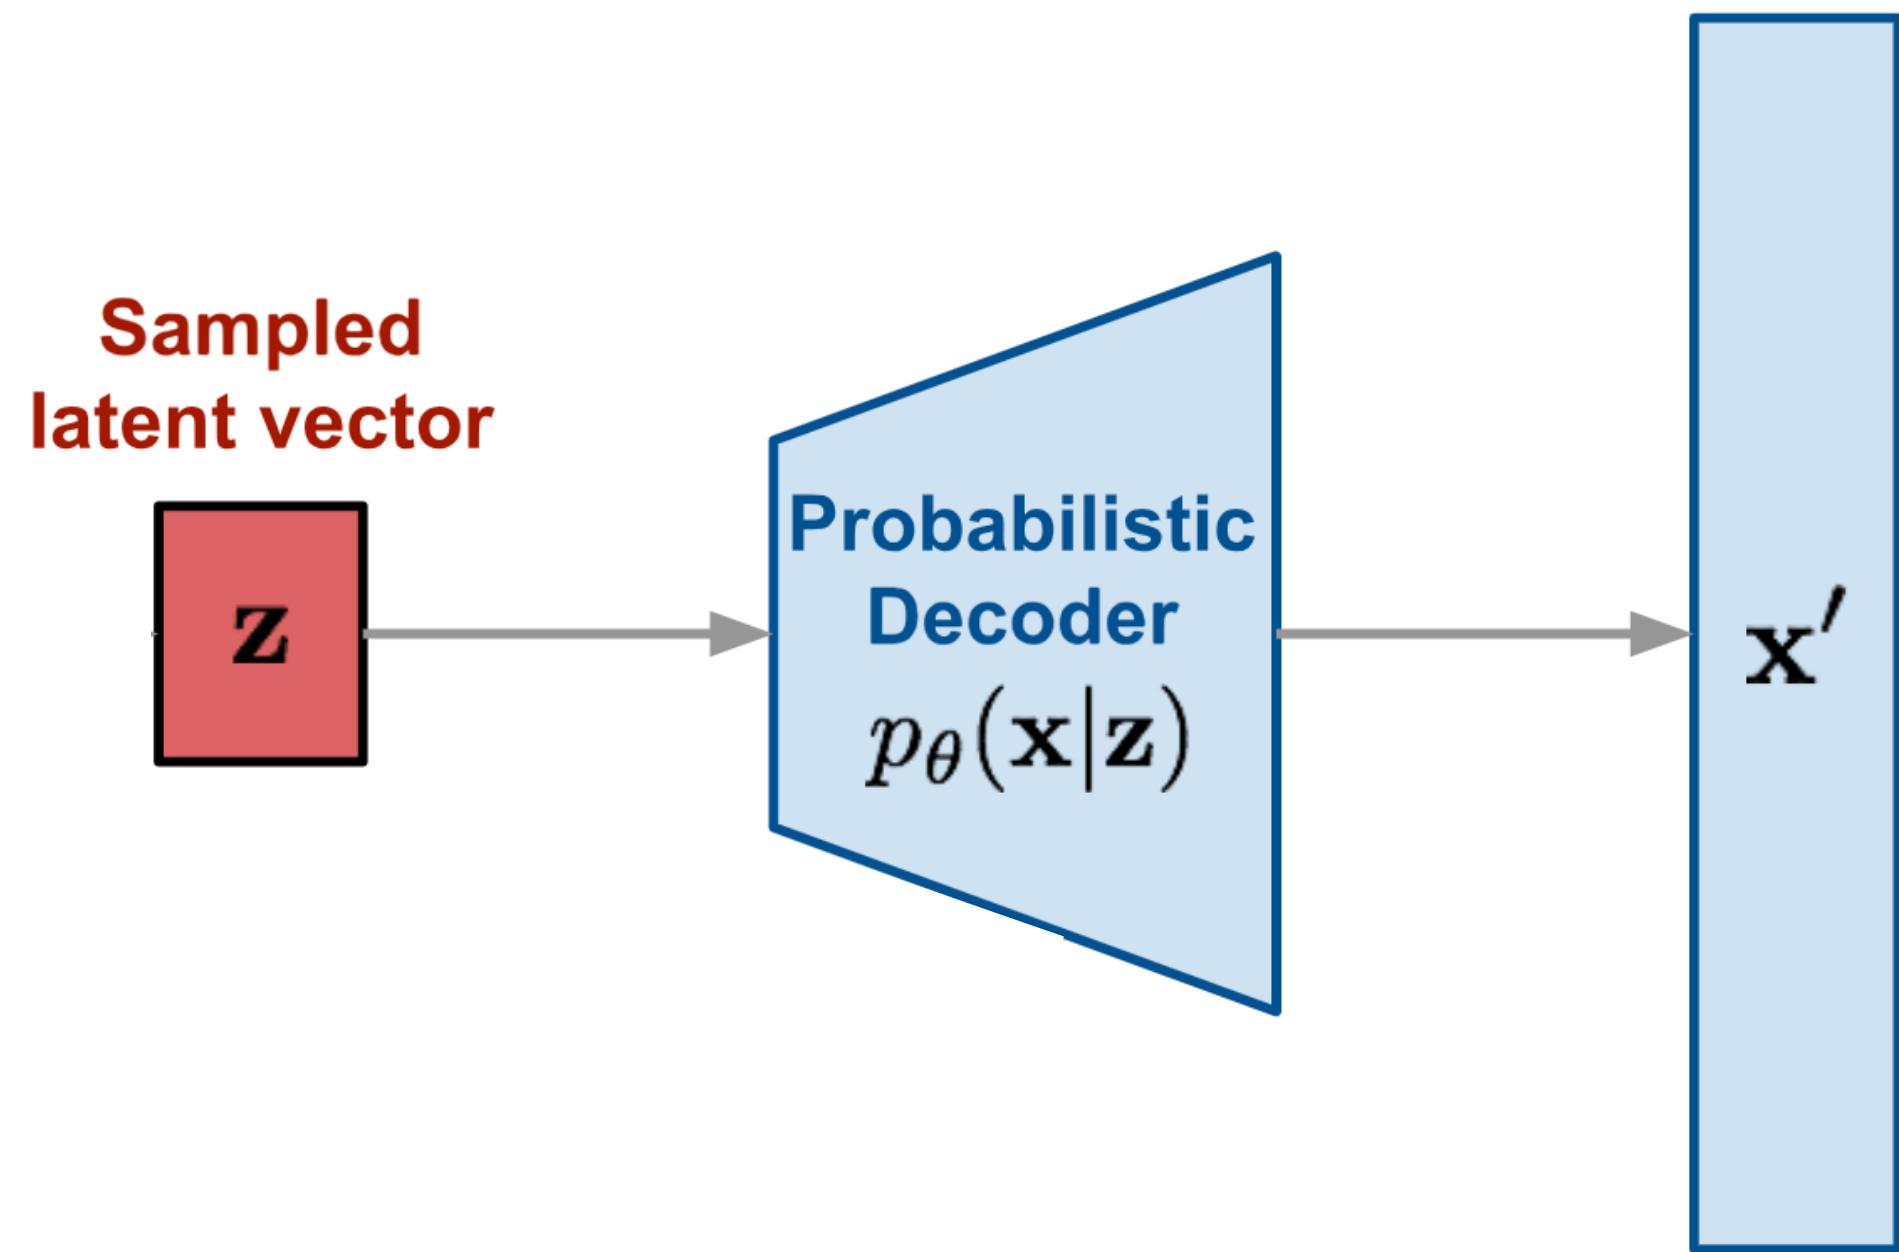
Reparametrization Trick



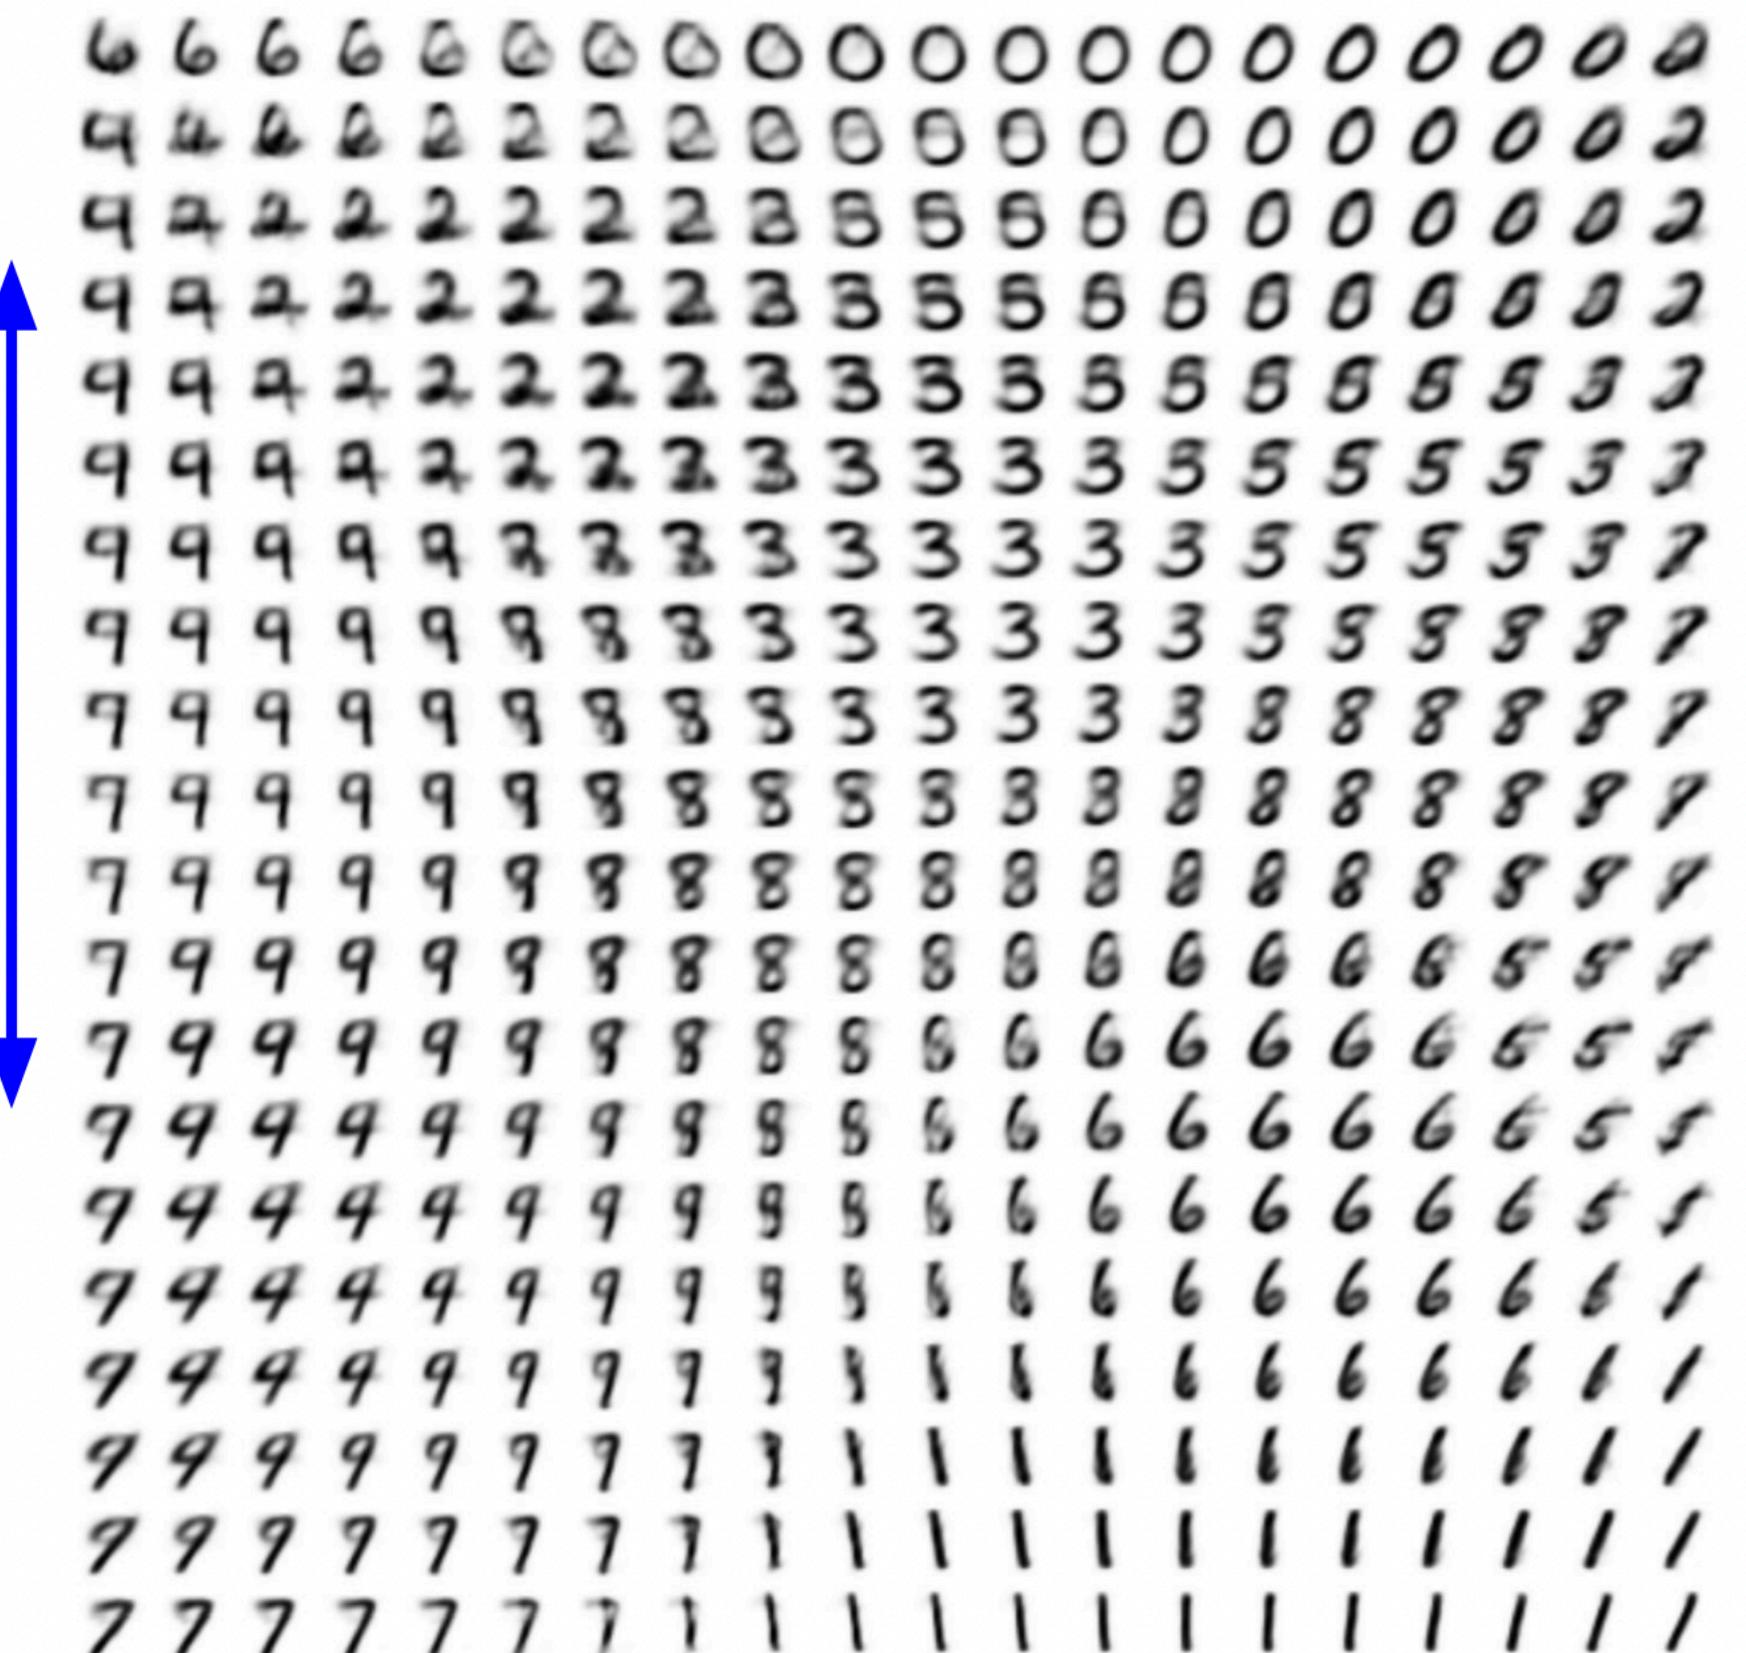
Variational Autoencoder (VAE)

Generating Data

Just sample $z \sim N(0, 1)$ and decode



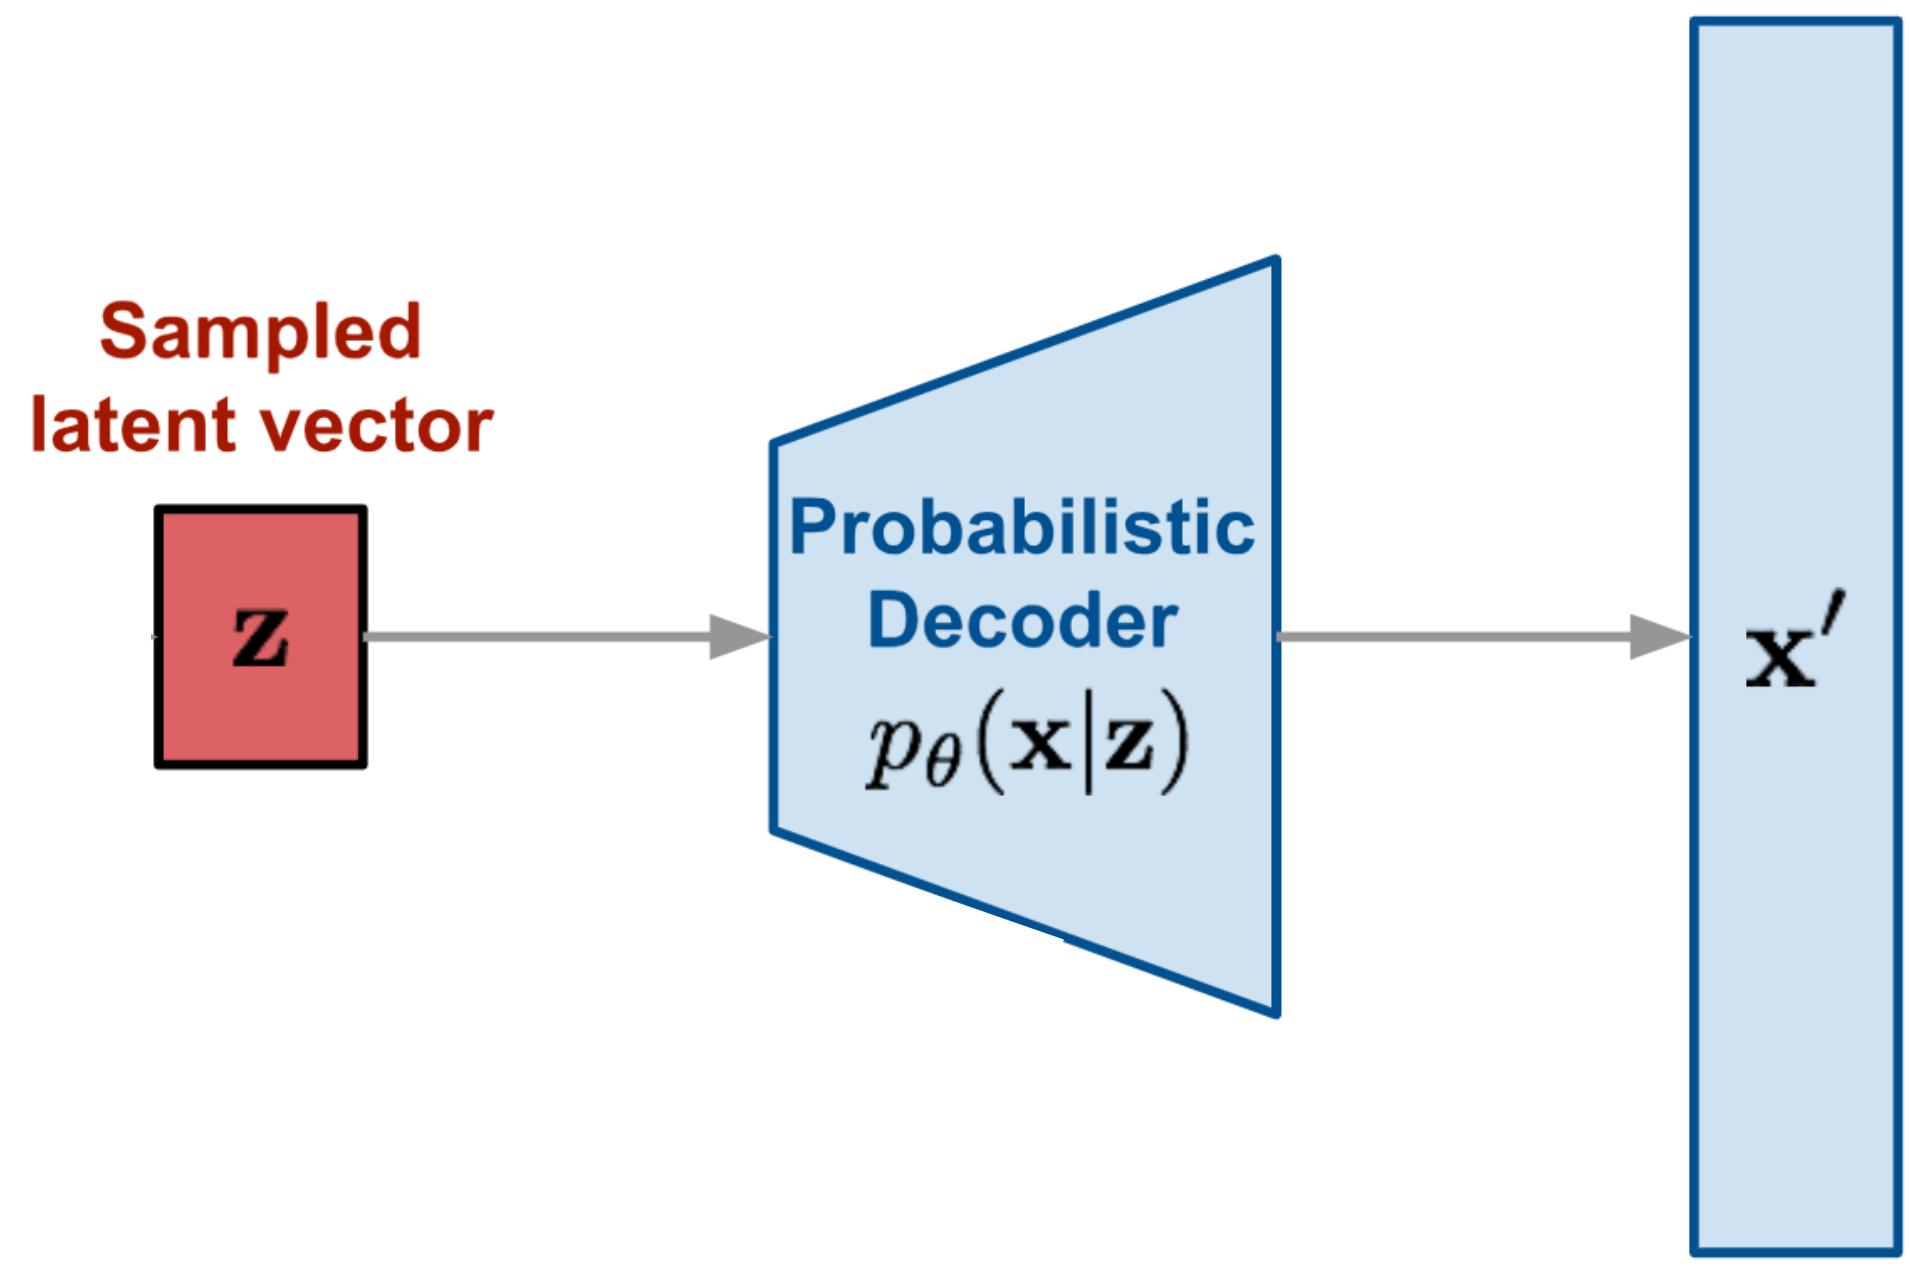
Data manifold for 2-d z



Variational Autoencoder (VAE)

Generating Data

Just sample $z \sim N(0, 1)$ and decode



Degree of smile

Vary z_1



Vary z_2 Head pose

Variational Autoencoder

Pros and Cons

- **Pros:**
 - Interpretable latent space enabling e.g. interpolation
 - Probabilistic encoder can yield feature representations for transfer to supervised downstream tasks
- **Cons:**
 - We only optimize a lower bound on the data likelihood
 - Sample quality is worse than that of GANs (which we will now talk about)

Today's Lecture

1. Supervised vs Unsupervised Learning
2. Generative Modeling
3. Applications of generative models
4. Fully-visible belief networks (FVBNs) & Autoregressive Decoders
5. (Variational) Autoencoders
- 6. Generative Adversarial Networks (GANs)**

Generative models – Taxonomy

Maximum likelihood estimation:

Find the set of parameters θ^* that maximizes the likelihood of our data x

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} \log p_{\text{model}}(x | \theta)$$

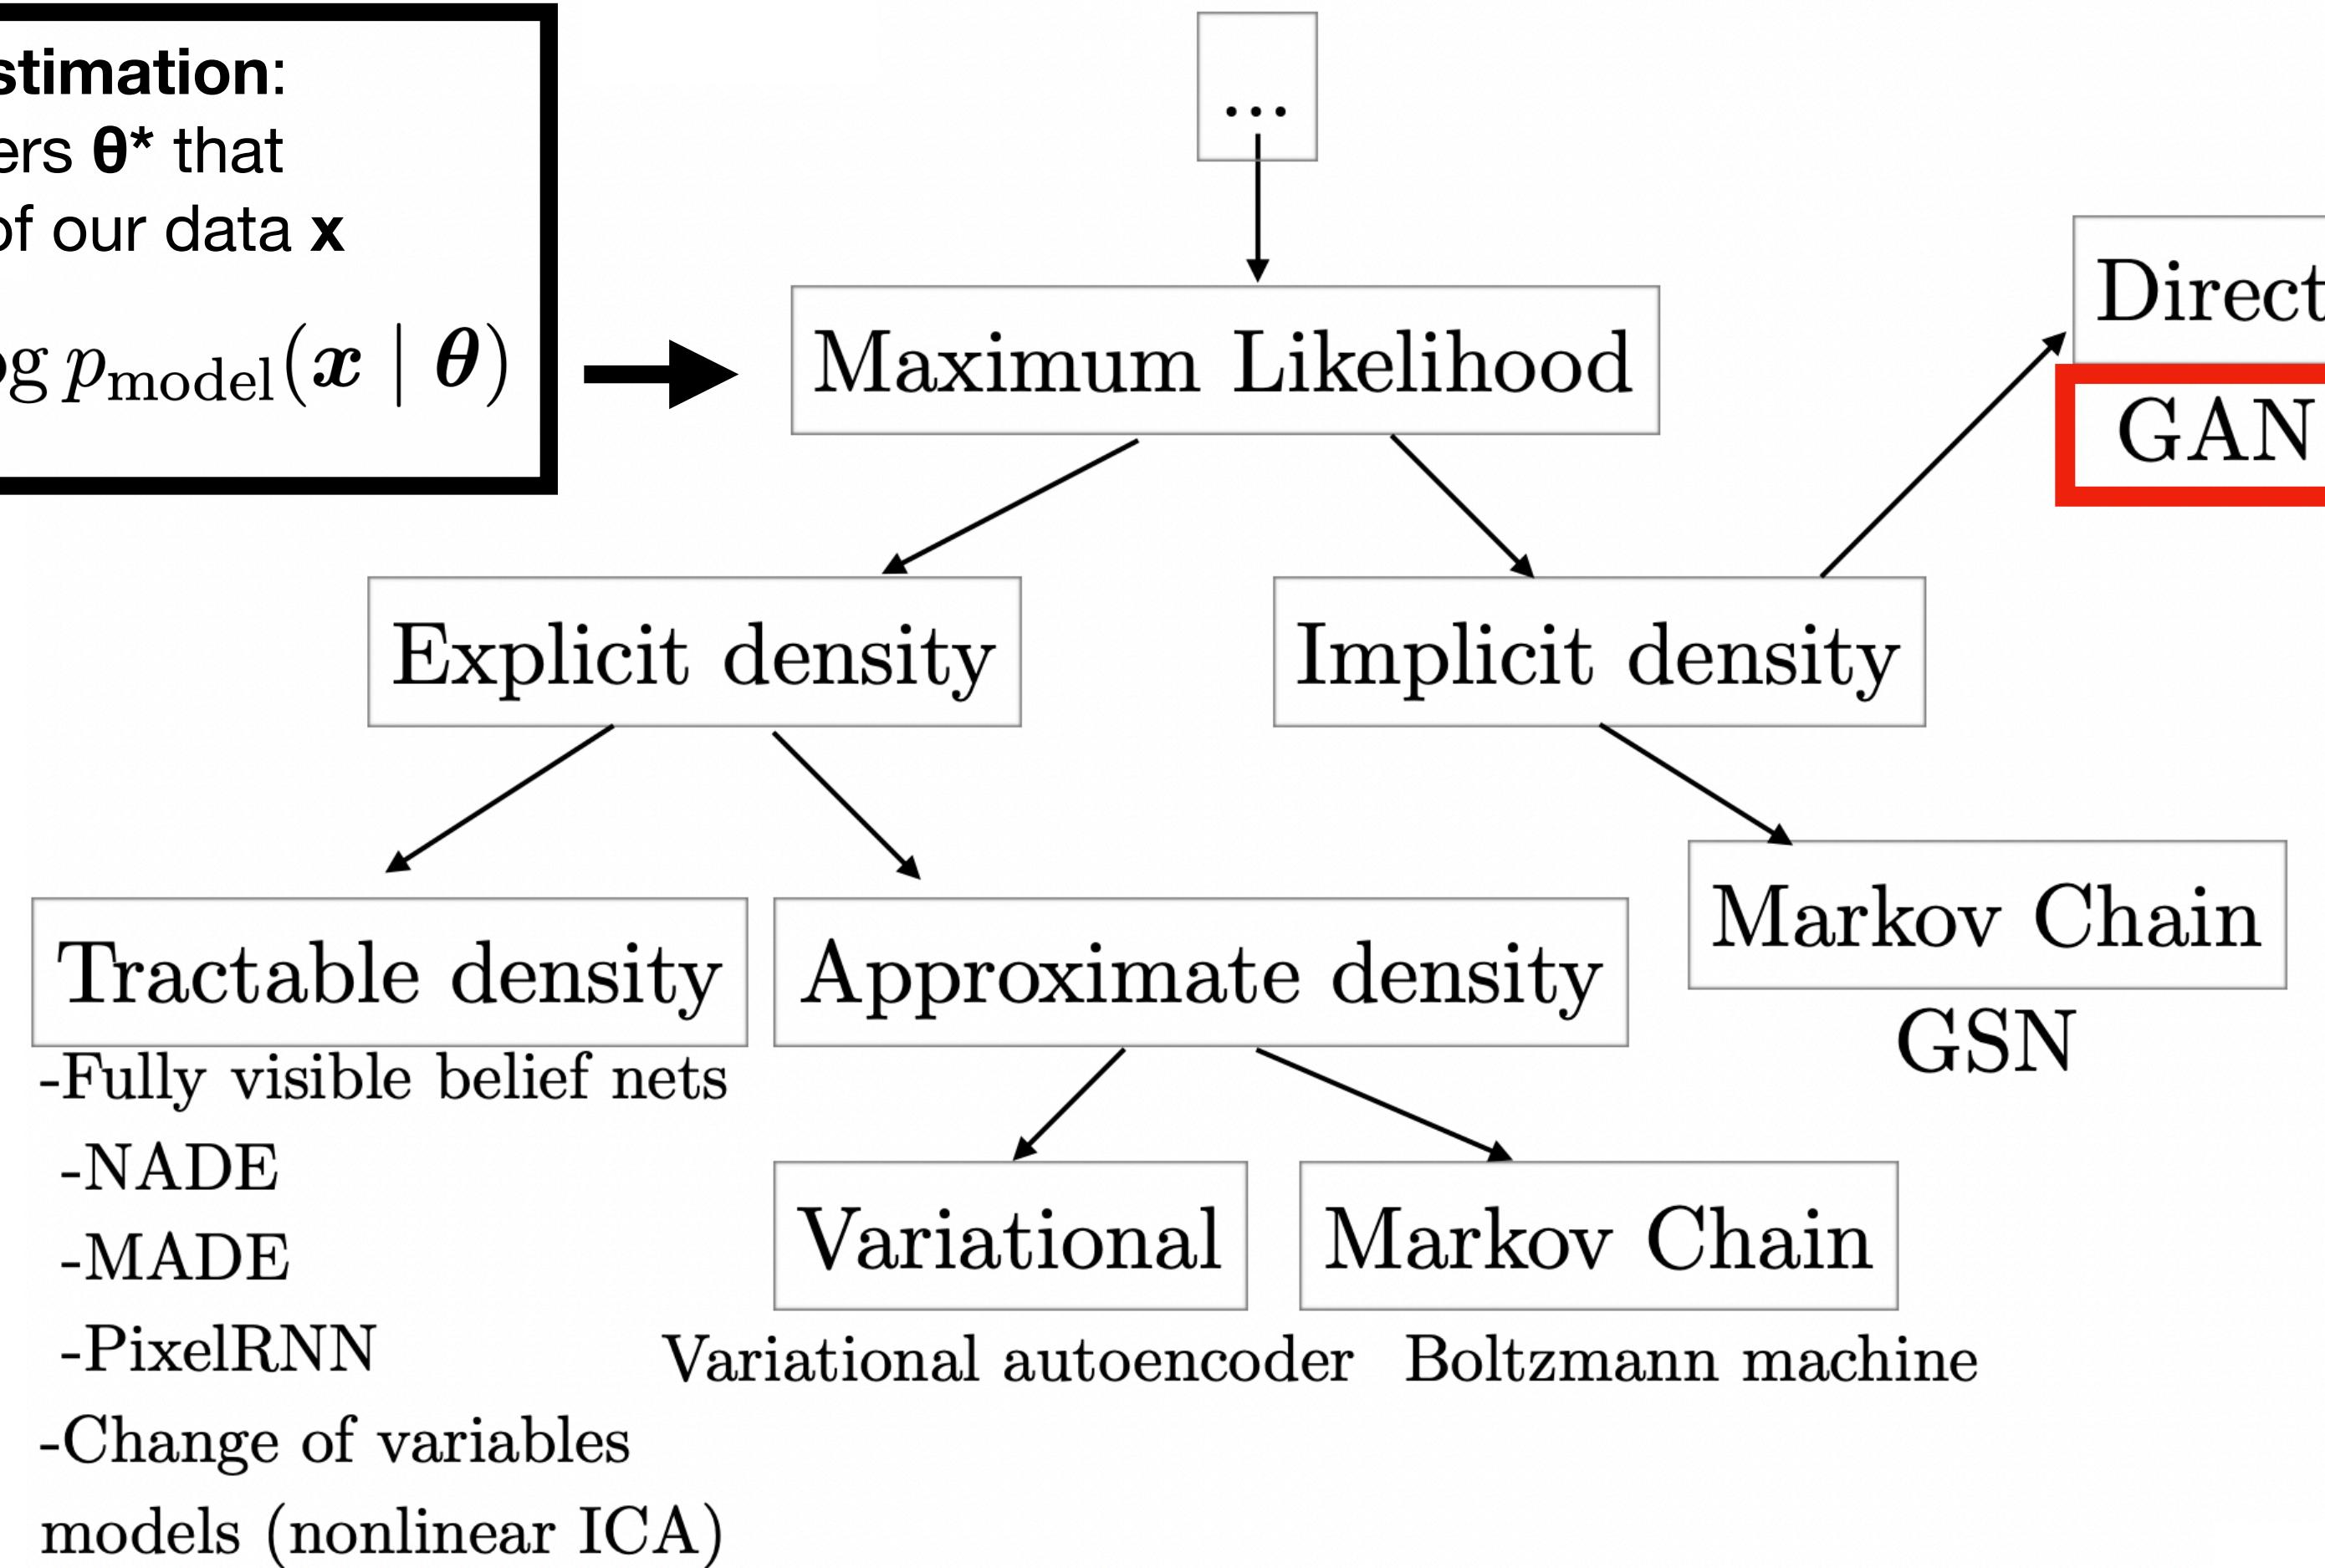


Figure taken from <https://www.iangoodfellow.com/slides/2016-12-04-NIPS.pdf>

Generative Adversarial Networks (GANs)

- **Implicit** density estimation, i.e., we are **only** interested in **learning to sample** from the data distribution **without actually knowing it**
- **Game-theoretic approach** that involves a **generator** and a **discriminator** network with **competing training objectives**

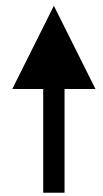


Generative Adversarial Networks (GANs)



Objective: Generate images that look **real**

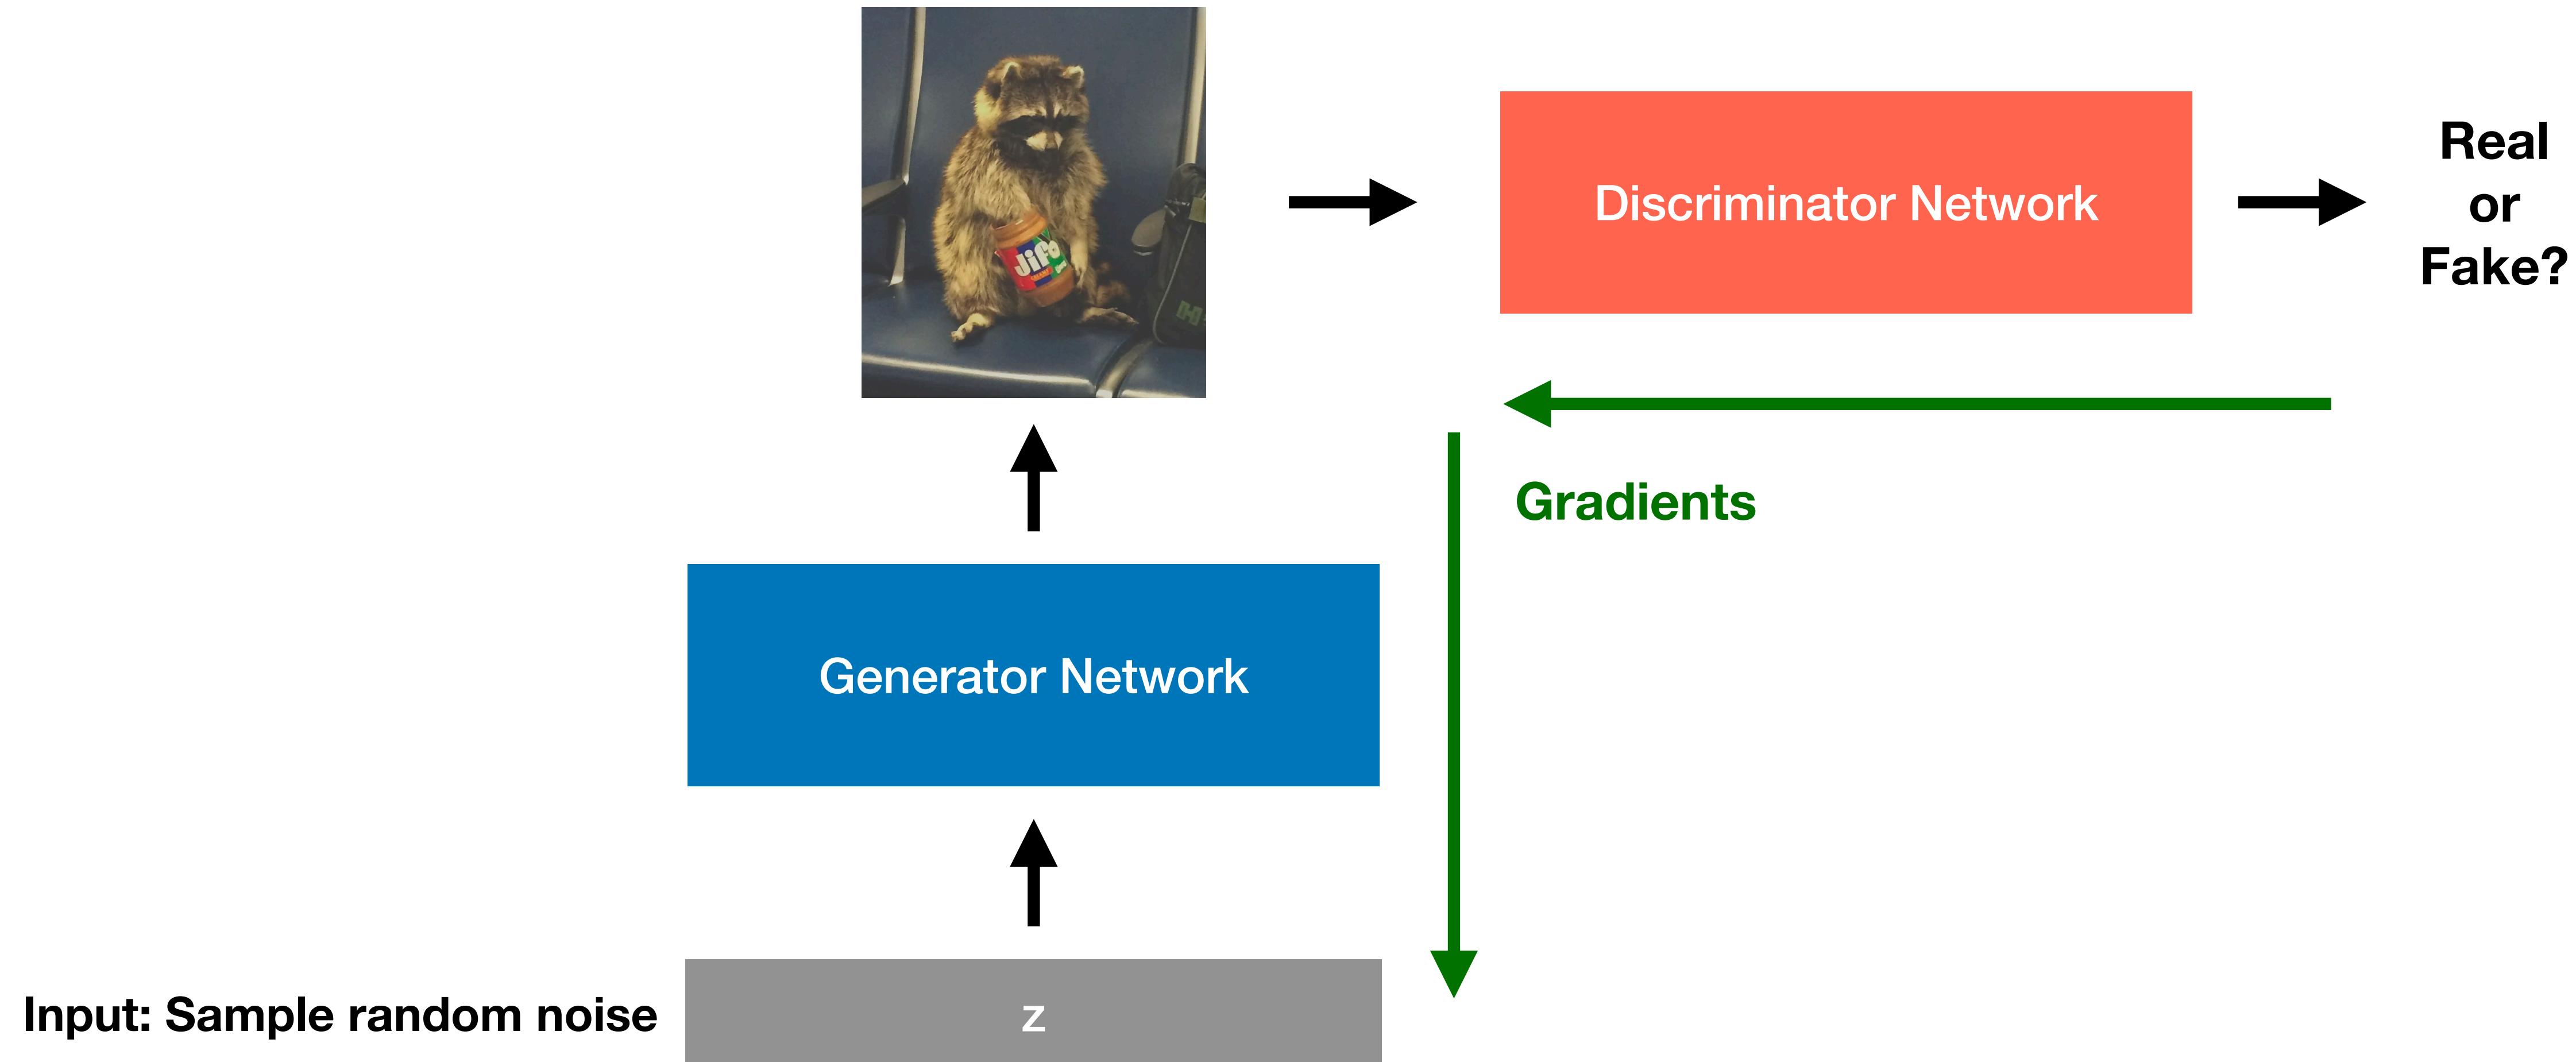
Generator Network



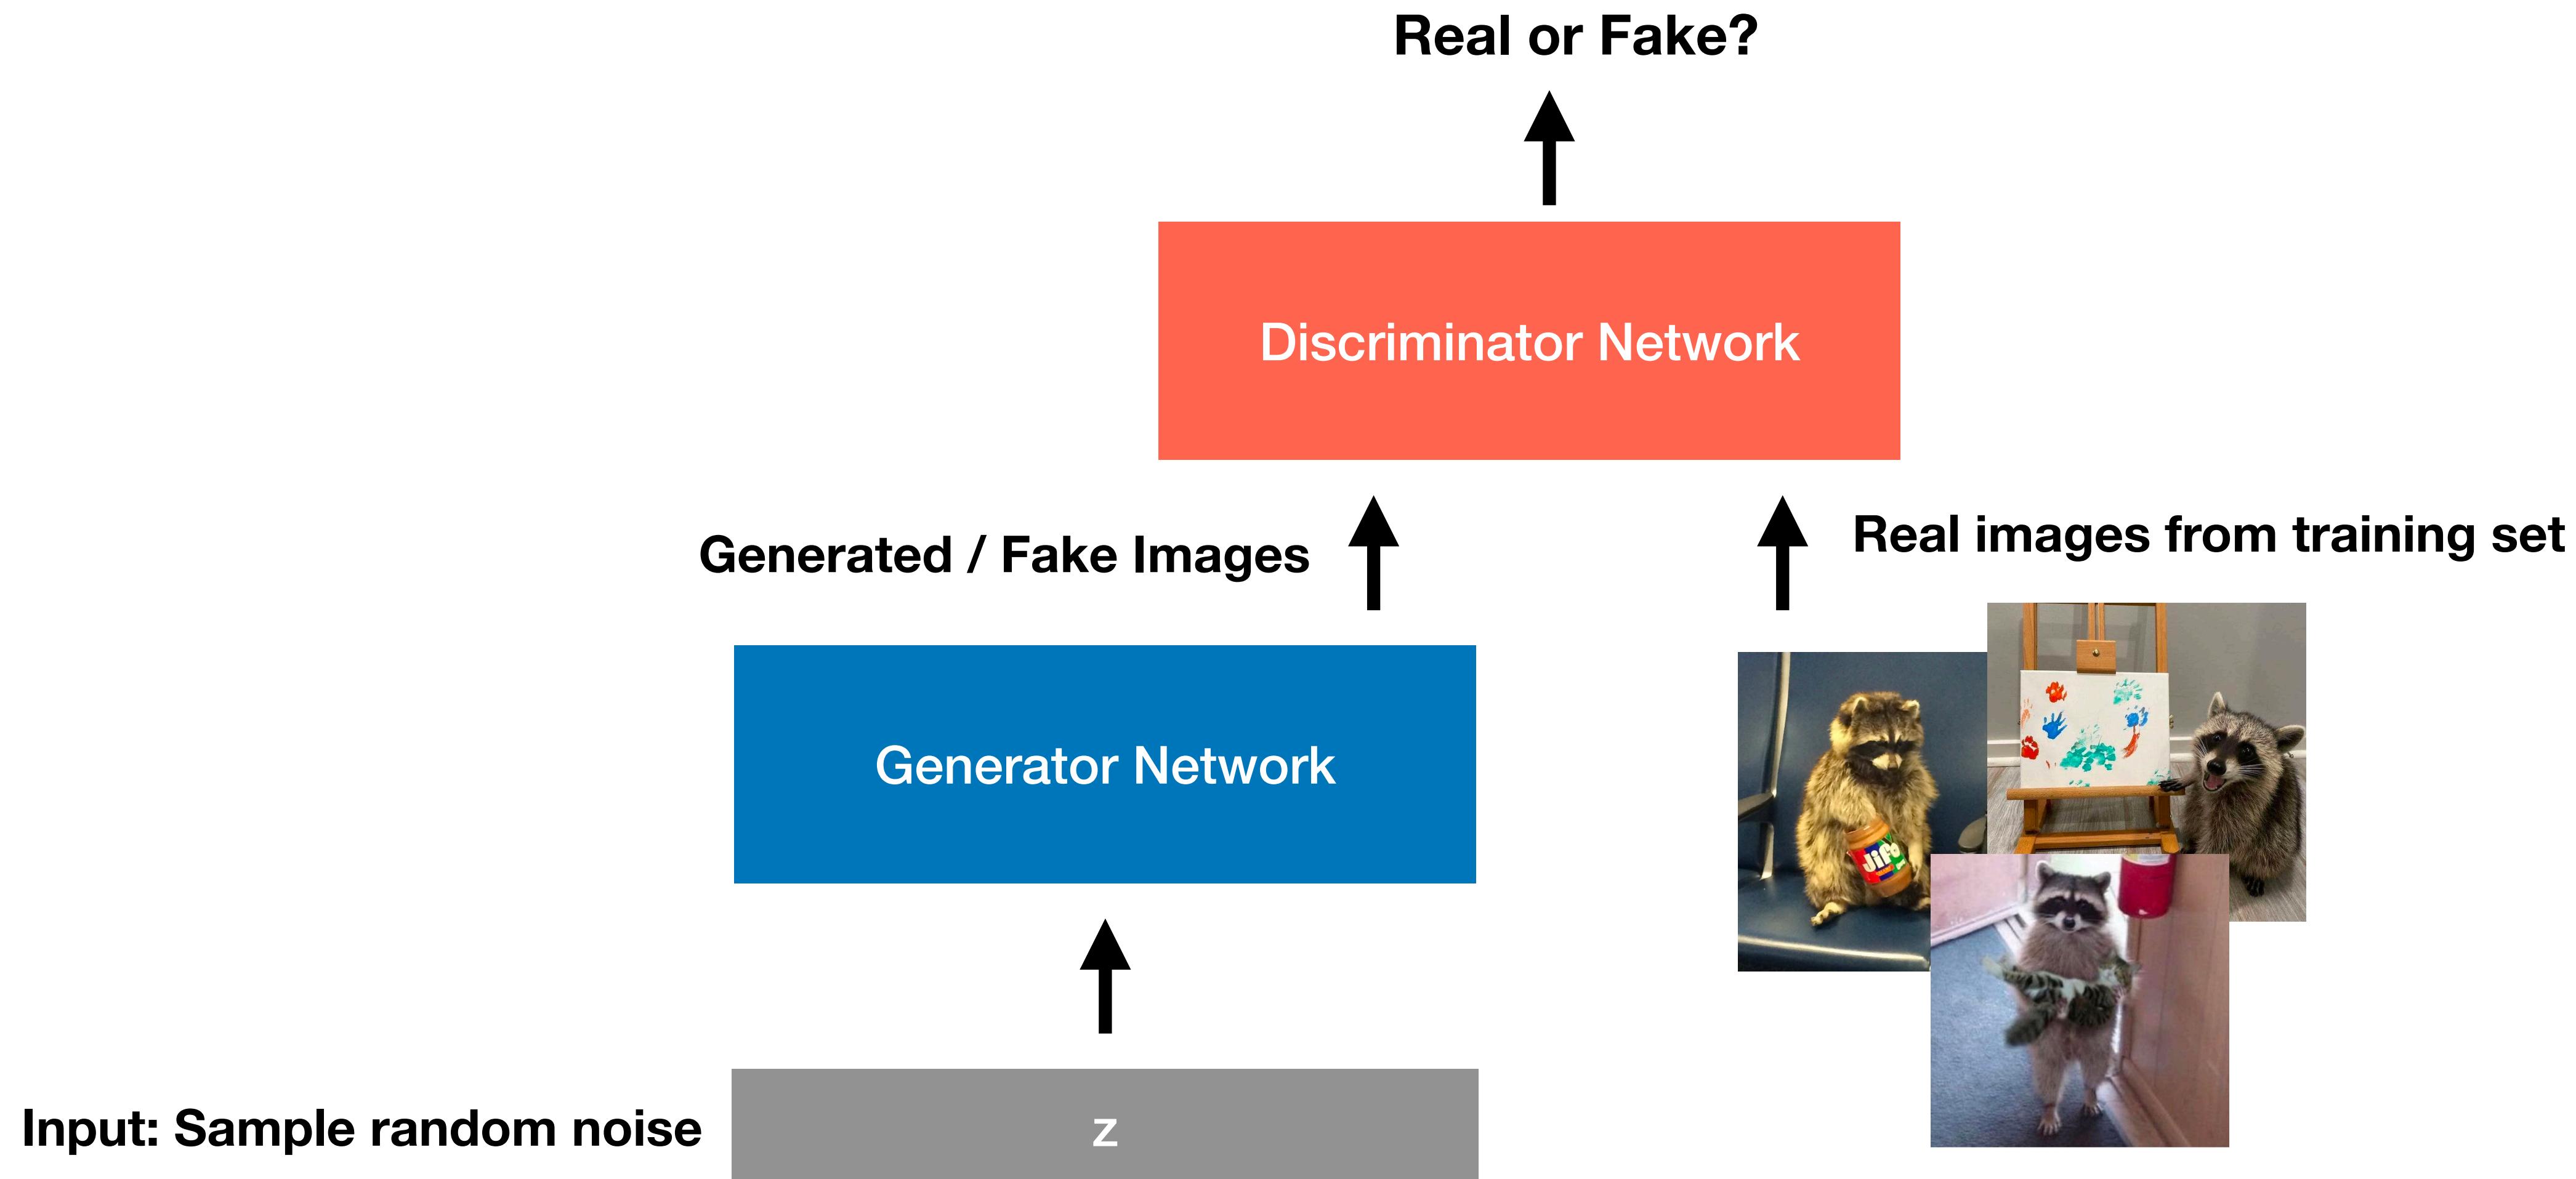
Input: Sample random noise

z

Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)

Two-Player Minimax Game

- **Generator Objective:** Fool discriminator by generating images that look real
- **Discriminator Objective:** Learn to distinguish between “fake” and “real” images

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Discriminator output for real image x

Discriminator output for fake image created by generator from sampled noise

Discriminator objective

Generator objective

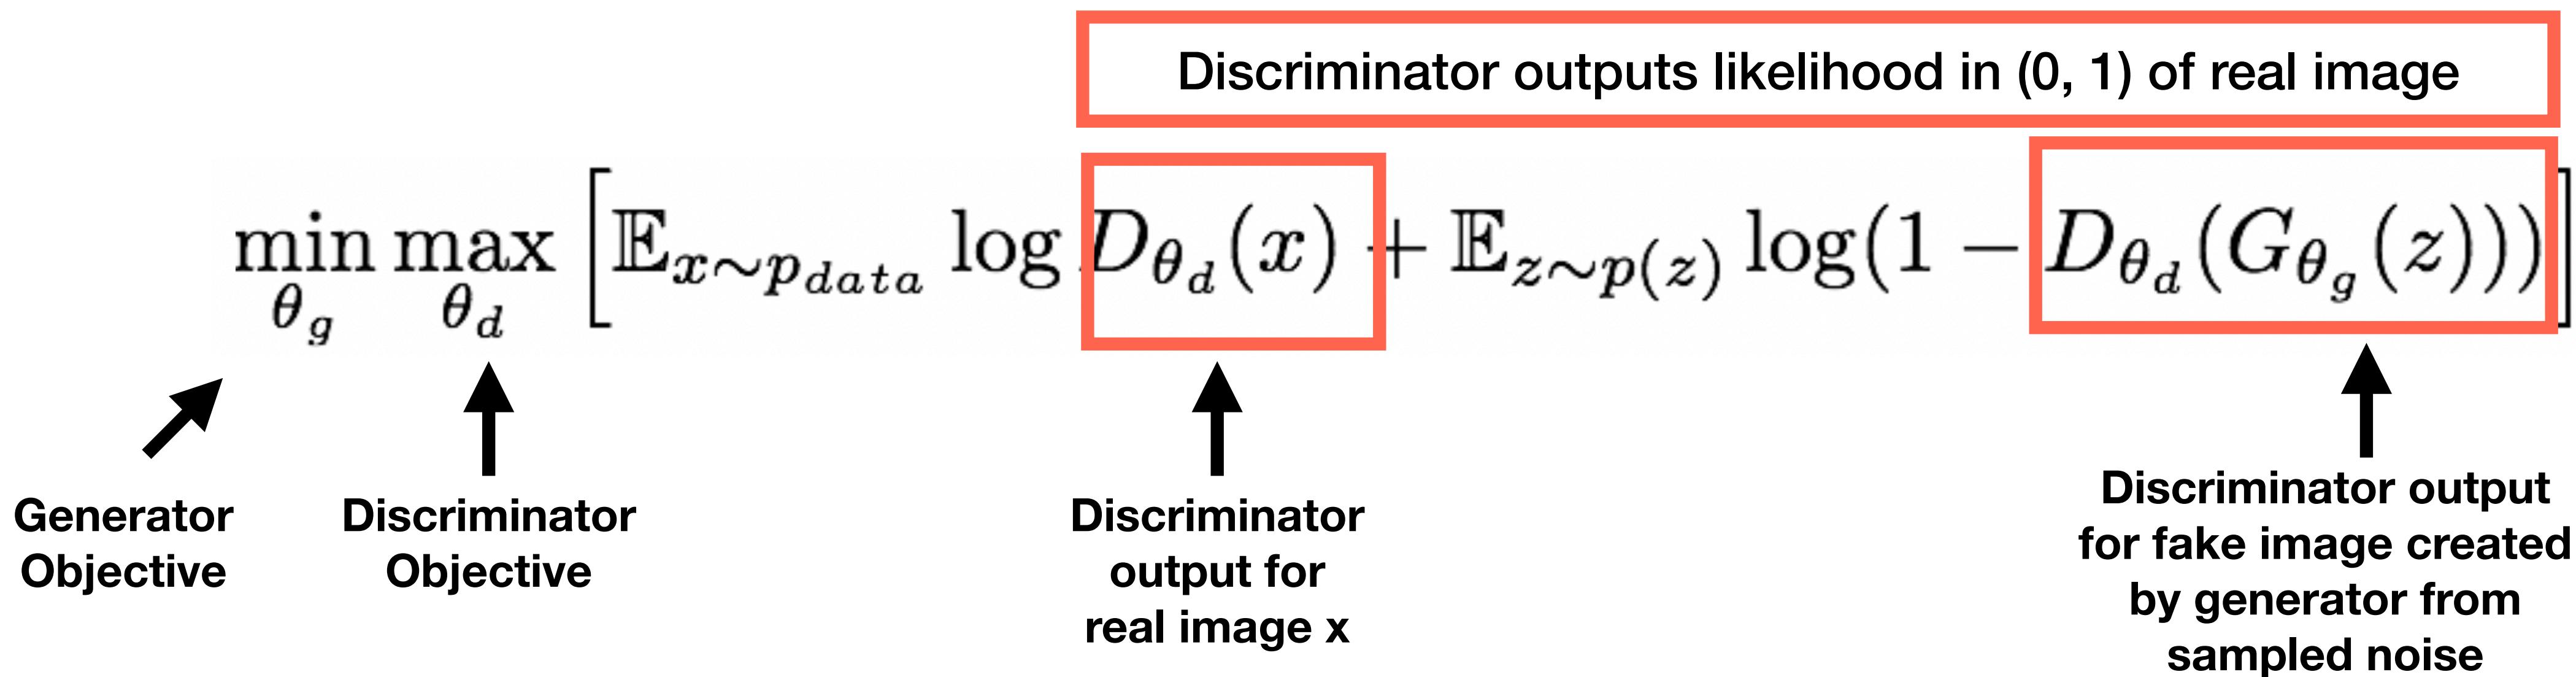
Discriminator outputs likelihood in (0, 1) of real image

The diagram illustrates the GAN loss function with annotations explaining each term. Arrows point from the text labels to the specific components in the equation:

- An arrow points from "Generator Objective" to the term $G_{\theta_g}(z)$.
- An arrow points from "Discriminator Objective" to the term $D_{\theta_d}(x)$.
- An arrow points from "Discriminator output for real image x" to the term $D_{\theta_d}(x)$.
- An arrow points from "Discriminator output for fake image created by generator from sampled noise" to the term $D_{\theta_d}(G_{\theta_g}(z))$.
- An annotation above the equation states "Discriminator outputs likelihood in (0, 1) of real image".

Generative Adversarial Networks (GANs)

Two-Player Minimax Game



- **Discriminator objective**: $D(x)$ close to 1, $D(G(z))$ close to 0
- **Generator objective**: $D(G(z))$ close to 1

Generative Adversarial Networks (GANs)

Two-Player Minimax Game

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between two optimization problems:

1. Gradient ascent on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. Gradient ascent on generator

$$\max_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(D_{\theta_d}(G_{\theta_g}(z)))$$

Note: **Maximizing** the probability of the discriminator being wrong is more stable than **minimizing E[log(1-D(G(z)))]** (the probability of the discriminator being right) because of gradient flow for bad generated samples in the beginning

Generative Adversarial Networks (GANs)

Two-Player Minimax Game

After enough training, the generator and discriminator reach the **Nash-Equilibrium** at which **neither can improve**

Now we can throw away the discriminator and happily generate samples with the generator



Generative Adversarial Networks (GANs)

Samples

- Try it out yourself: <https://thispersondoesnotexist.com/>

- Interpolation:



Suggested and Further Resources

- **The Illustrated GPT-2:** <https://jalammar.github.io/illustrated-gpt2/>
- **Autoencoder Blog Post** (Includes **Vector-Quantized VAE**, a state-of-the-art VAE used in 2022): <https://lilianweng.github.io/posts/2018-08-12-vae/>
- **NIPS 2016 Tutorial: Generative Adversarial Networks**, Ian Goodfellow
- **StyleGAN Paper** (Karras et al., 2019)
- **FAIR PyTorch GAN Zoo:** https://github.com/facebookresearch/pytorch_GAN_zoo